



Azure

Azure Messaging

Standards Matter!

Clemens Vasters

Principal Architect – Azure Messaging, Microsoft

@clemensv

OASIS AMQP Technical Committee Chair

OASIS MQTT TC Member

CNCF CloudEvents Architect

OPC UA PubSub Architect

Service Bus

Event Hubs

Event Grid

Relay

IoT

Agenda

What are the Azure Messaging services?

Messaging Patterns and Protocol Characteristics

Push/Pull, Queues, Ingestors/Streams, Routers/Distributors

Messaging Protocol Standards

MQTT, AMQP, HTTP, Kafka, gRPC

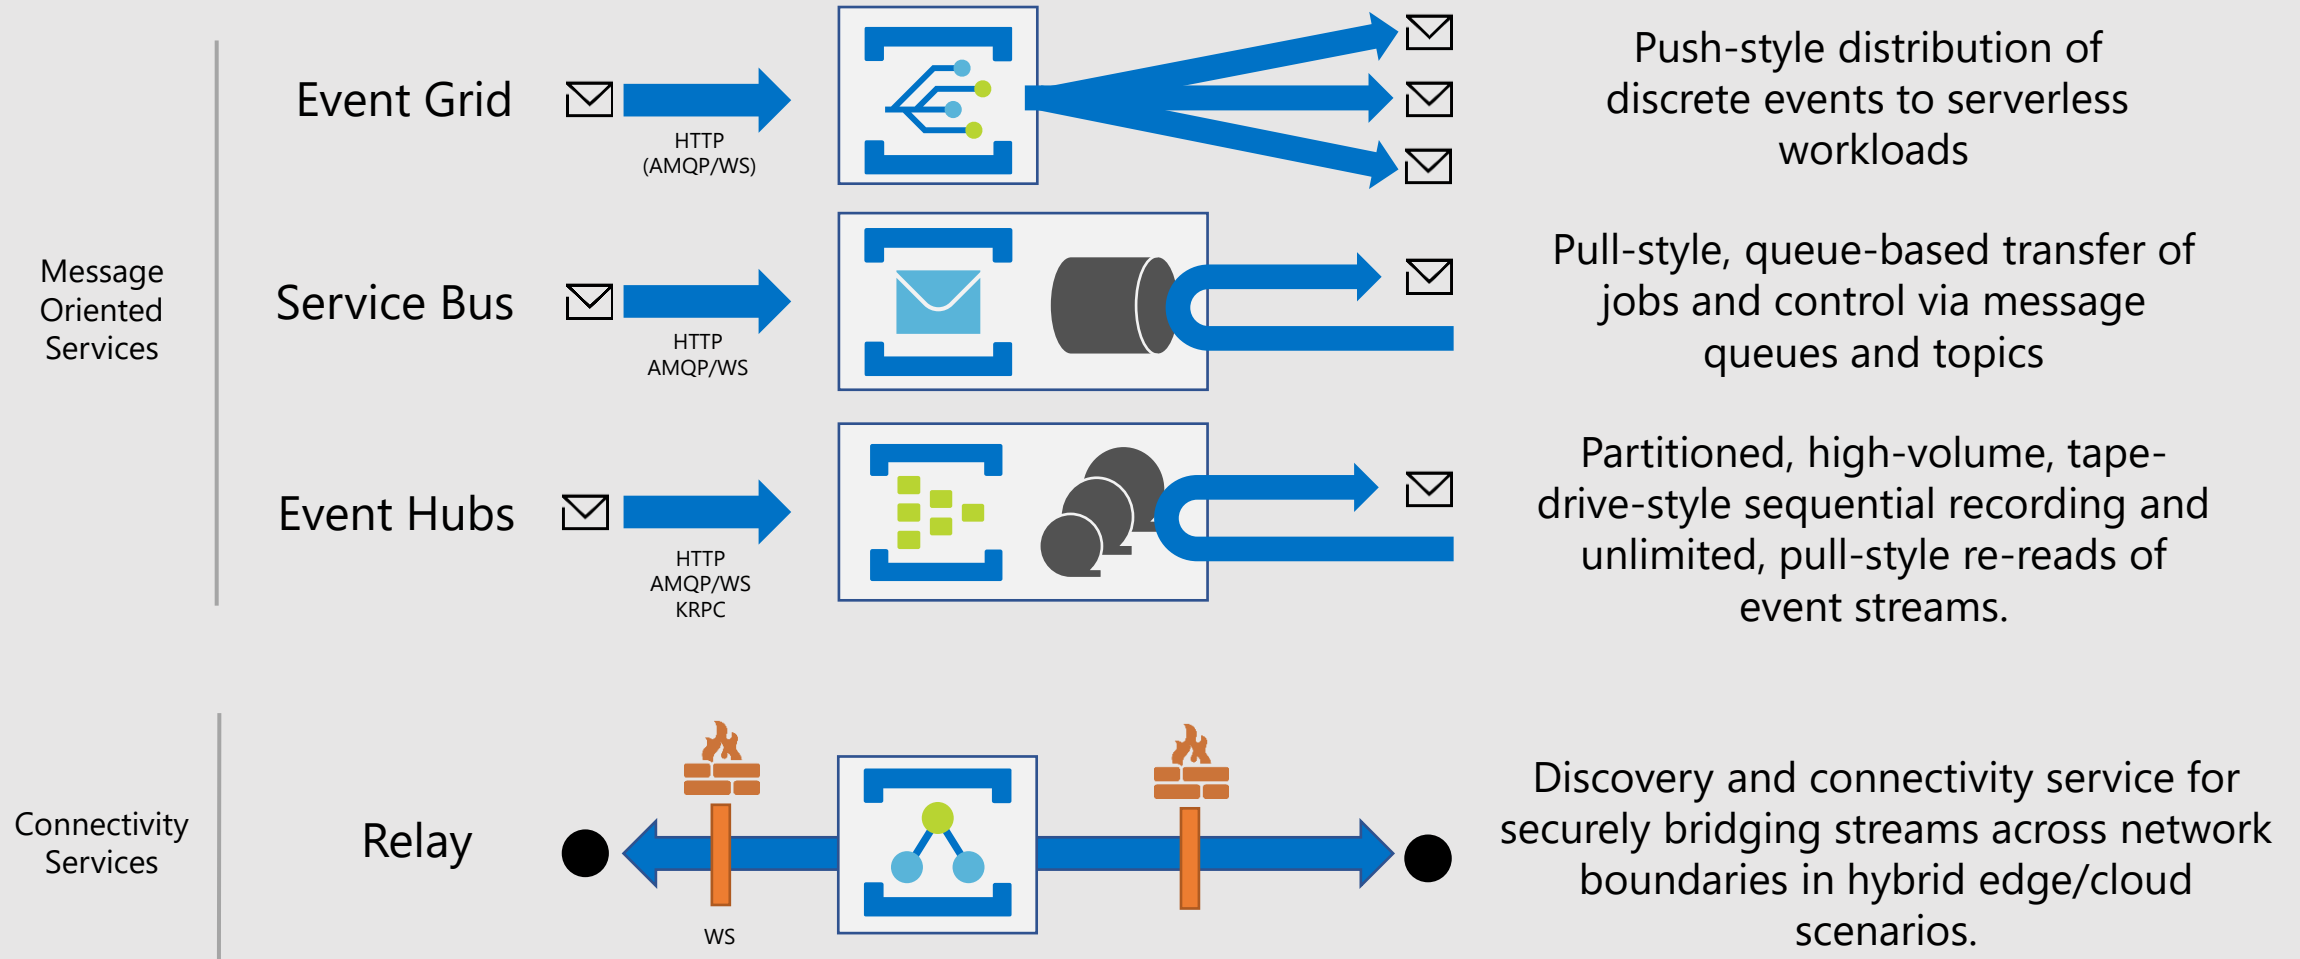
Encoding Standards

XML, CSV, JSON, CBOR, Avro, Thrift, Protobuf

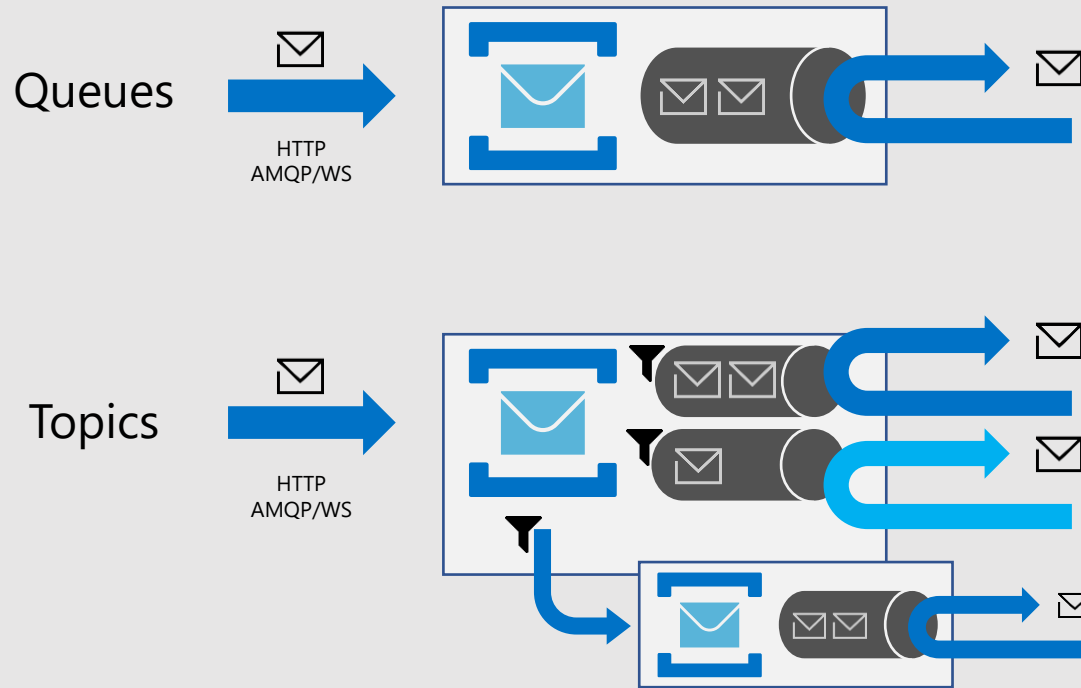
Abstraction Standards

OPC UA, CNCF CloudEvents, JMS 2.0

Azure Eventing and Messaging Core Services



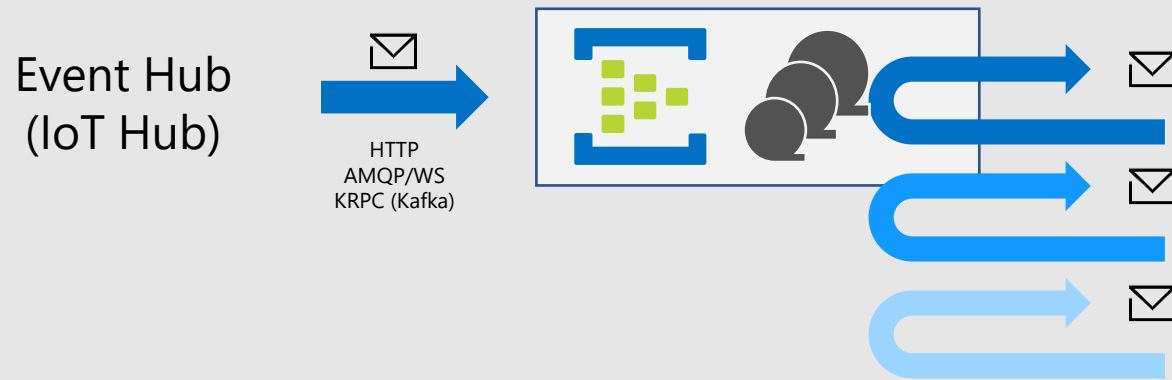
Service Bus Architectural Patterns



- Assignment of work with load-aware balancing
 - Load-leveling for "spiky" workload traffic shapes
 - Transactional, once-and-only-once processing
 - Multiplex handling of in-order message sequences
 - Deduplication, deferral, and "poison" handling
-
- All of the above, plus:
 - Copies to 100s of concurrent subscribers
 - Filter rules and message markup
 - Message routing

Service Bus is a "swiss army knife" for messaging-driven workloads.


Event Hubs Architectural Patterns



- Ingestion and storage of large event streams
 - > 2 Gigabyte per second if required
- Separation of event streams into partitions
- Client-chosen offsets into event stream allow arbitrary reads and re-reads during retention
- Retention of raw event data from 1 up to 90 days
- Automated archival into Avro containers for subsequent batch-style processing
- Publisher policies for data origin attestation and access control

Event Hubs is a high-scale, high-availability, multi-protocol event stream engine used for collecting and consolidating events for real-time analytics and other high-throughput computations

Event Grid Architectural Patterns



- Ingestion and push-style distribution of discrete events (events not correlated into streams) to interested subscribers.
- Per-subscriber application of simple and complex filters to select particular events of interest
- Abuse protection for event publishers
- Event schema mapping and support for CNCF CloudEvents 1.0 standard and bindings
- Multitenancy support for SaaS applications.
- Simple integration with a catalog of available event sources and sinks.

Event Grid is the Azure-wide eventing backplane for distributing and handling discrete events raised at the platform level, by custom applications, and by partner platforms.

Messaging Patterns



Jobs and Commands

Intents


Expectations
Conversations
Contracts
Control Transfer
Value Transfer

Events

Facts


History
Context
Order

Jobs and Commands



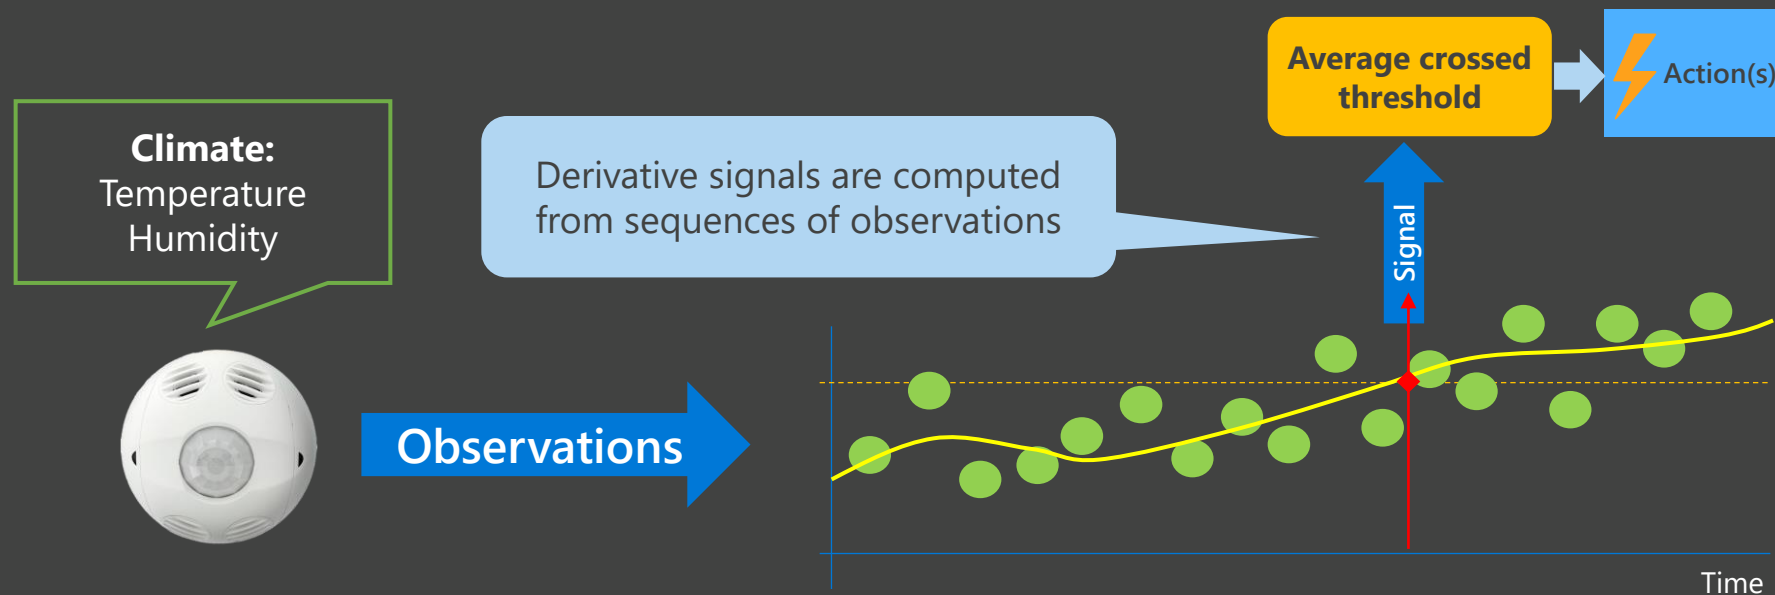
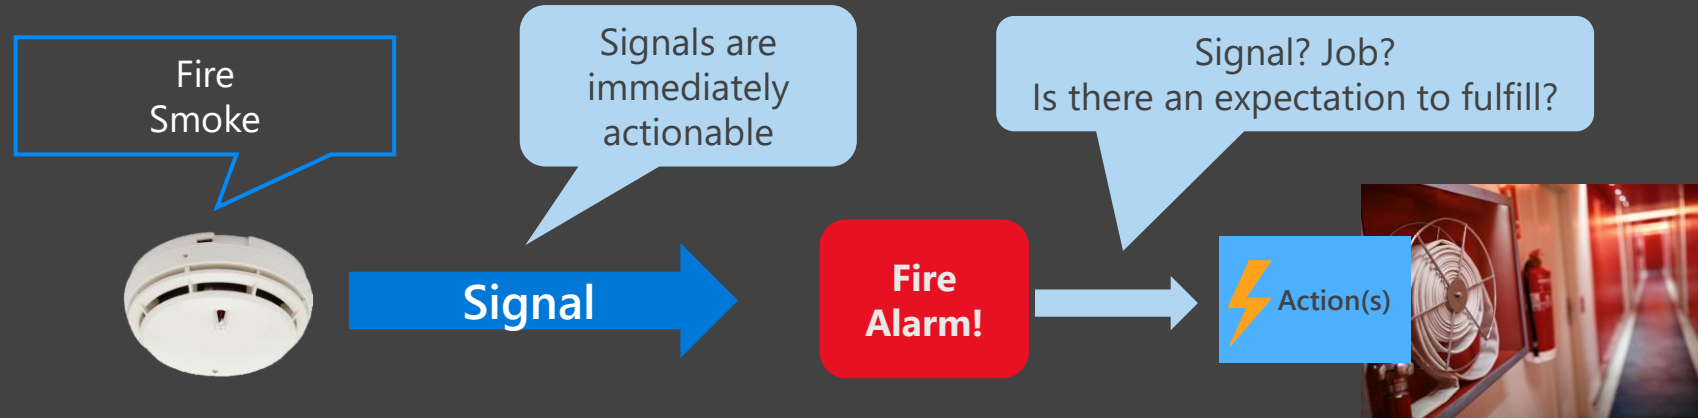
One message, one receiver

Events



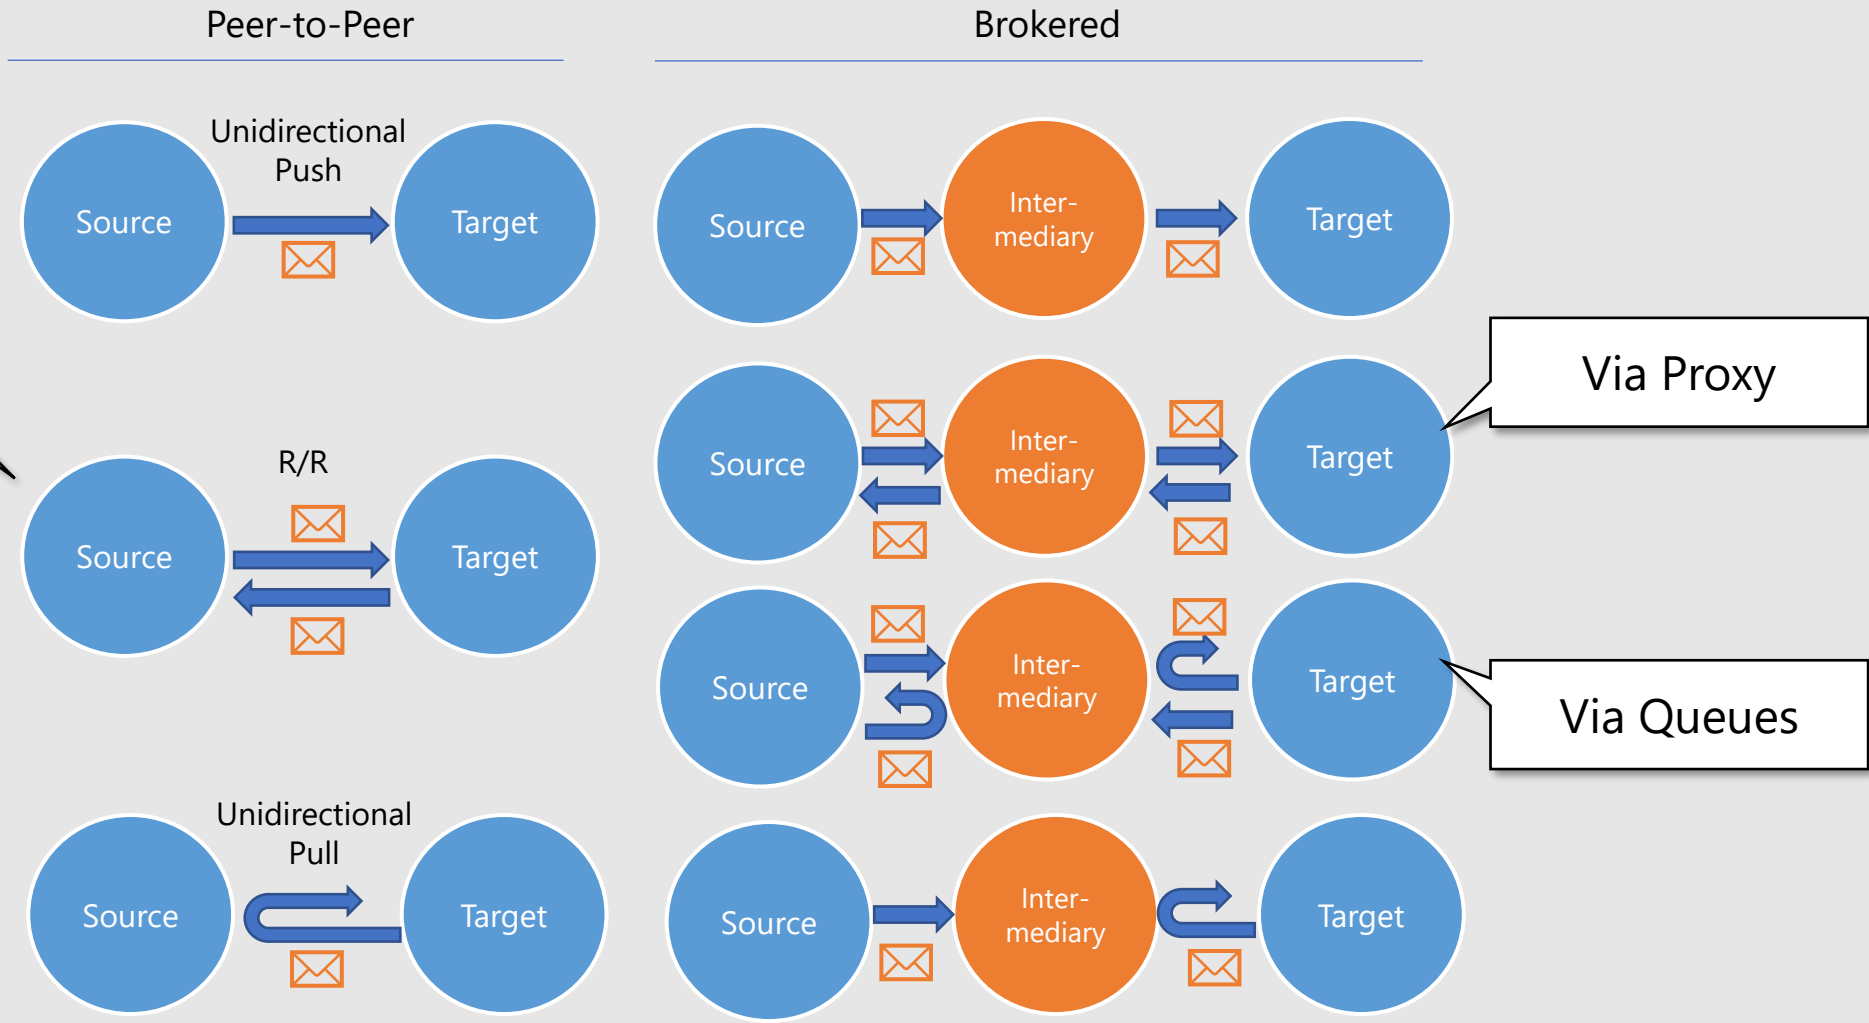
One event, 0 to many receivers

Events: Observations, Signals, Jobs

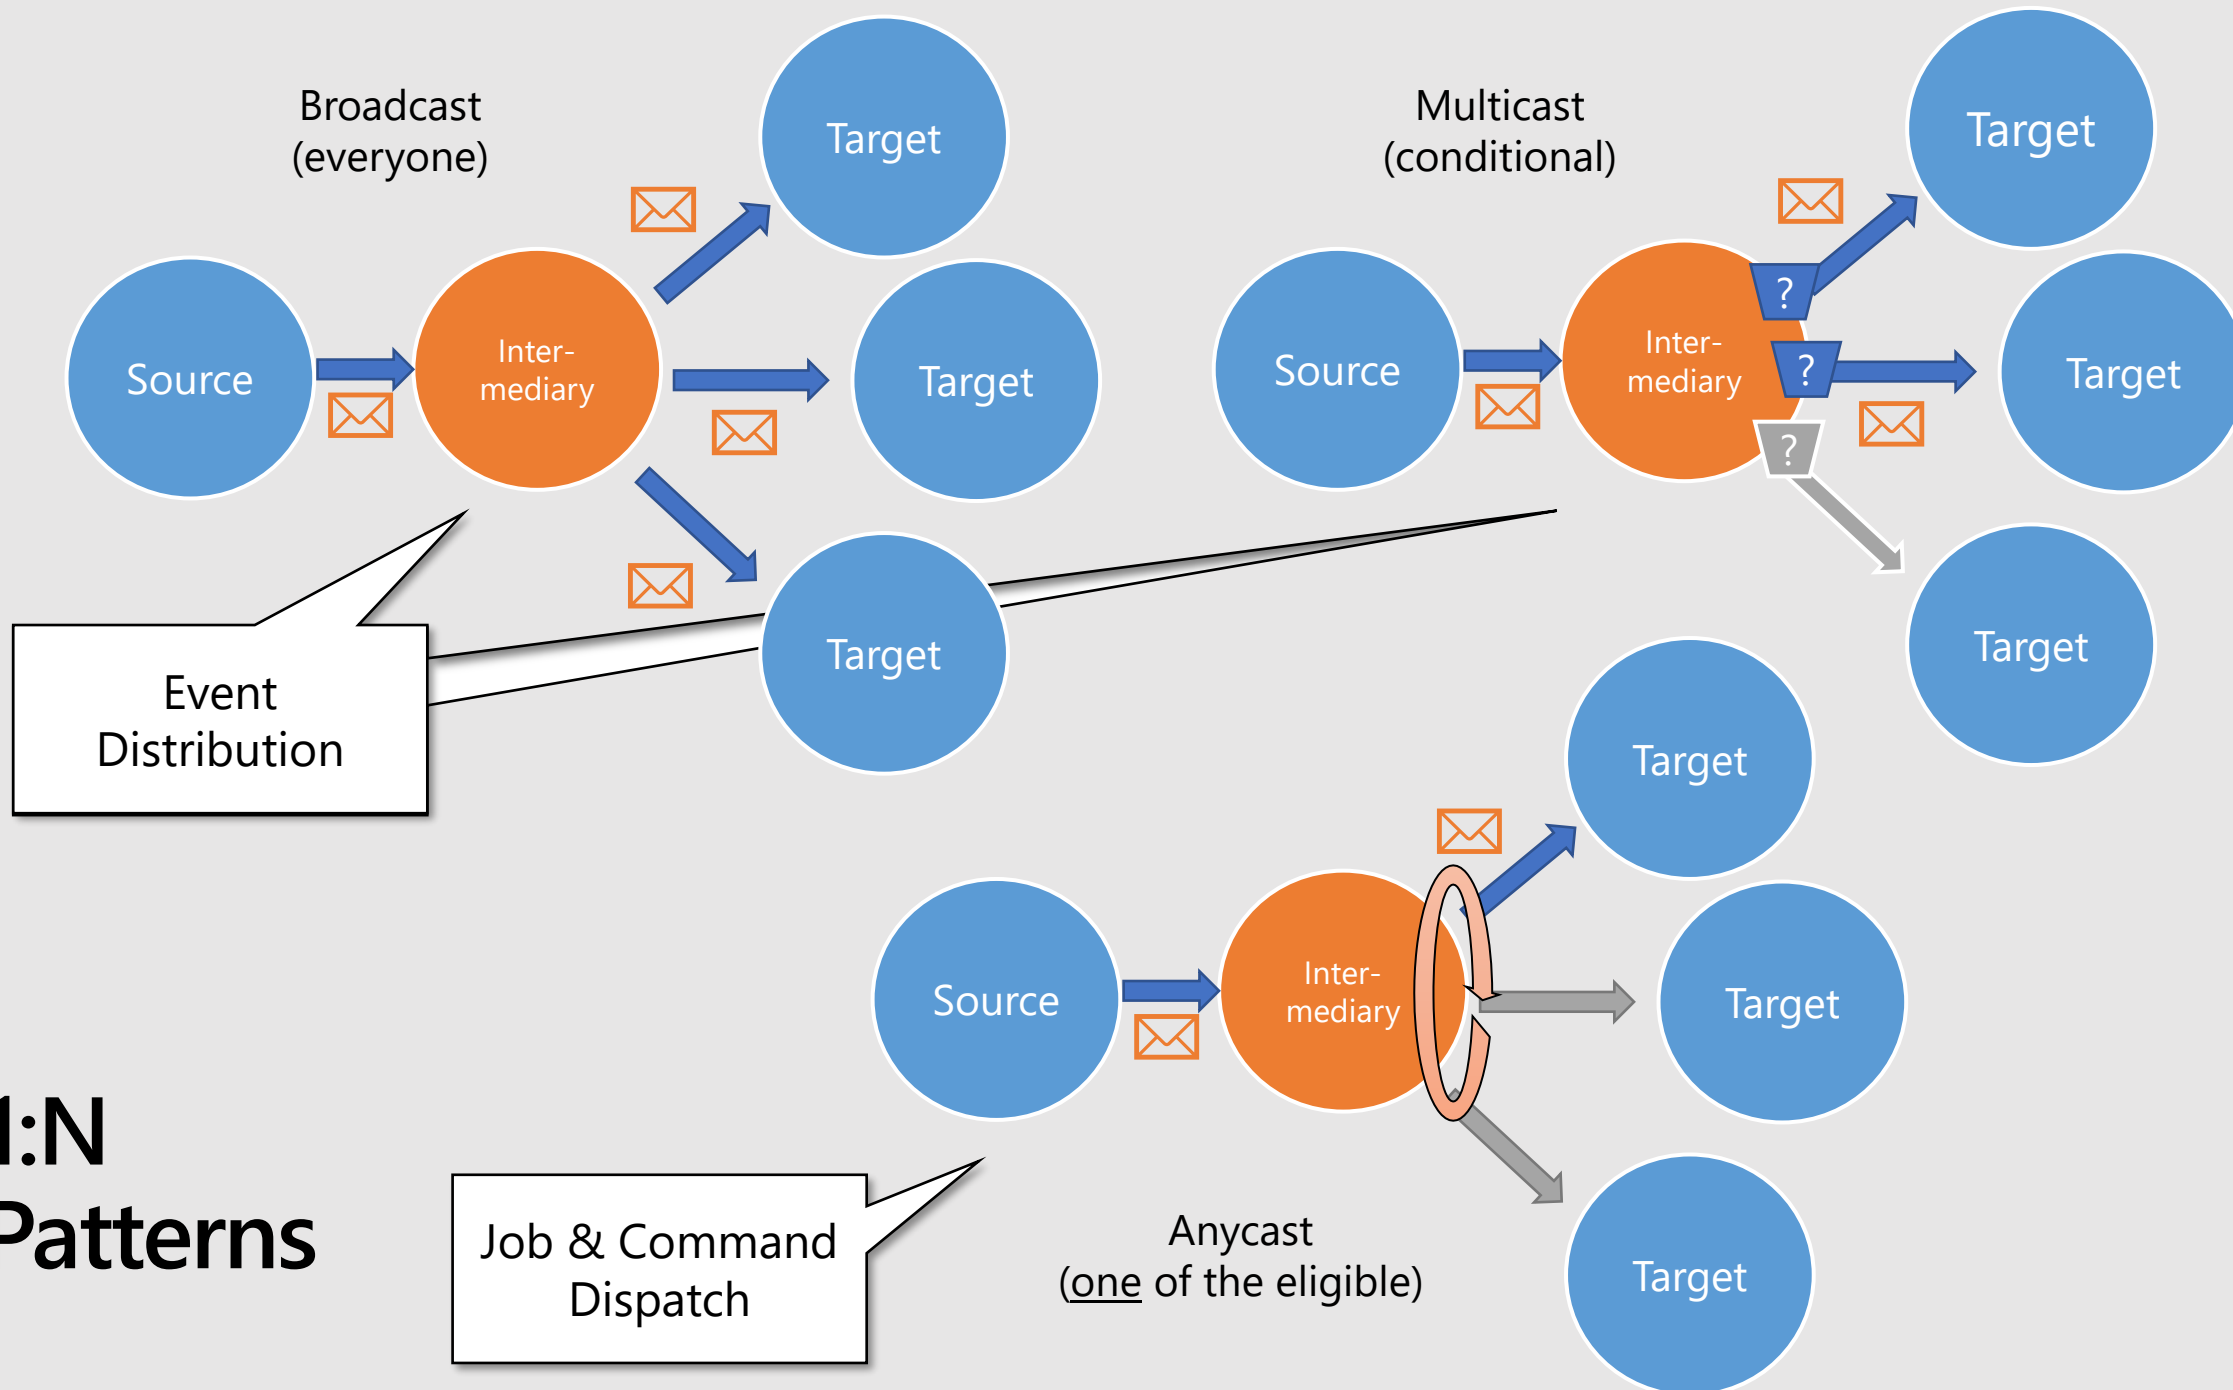


HTTP & RPC

1:1 Patterns




1:N Patterns




Job & Command Dispatch

Beyond a single hop: Event Journeys


Device Status Change or Alarm



Web site visitor statistics



Package delivery

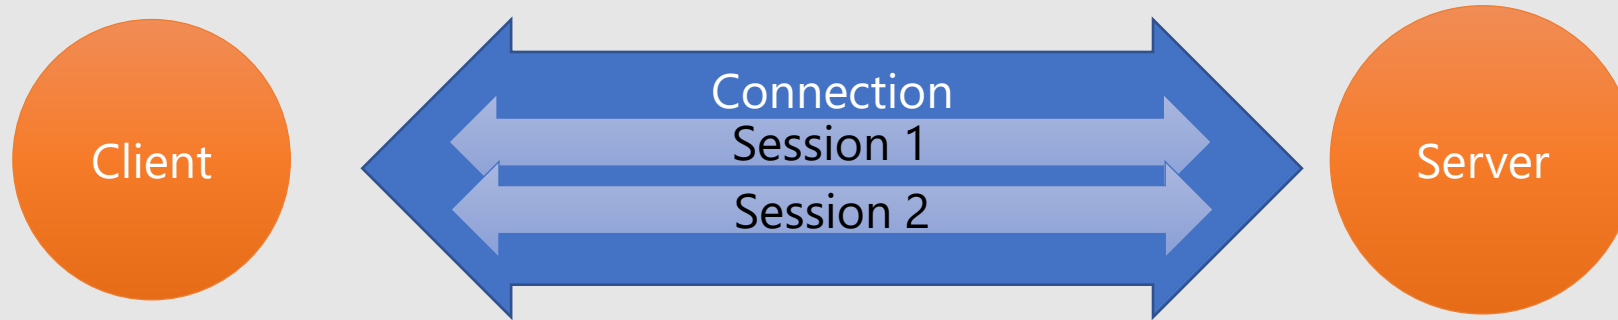


Symmetry



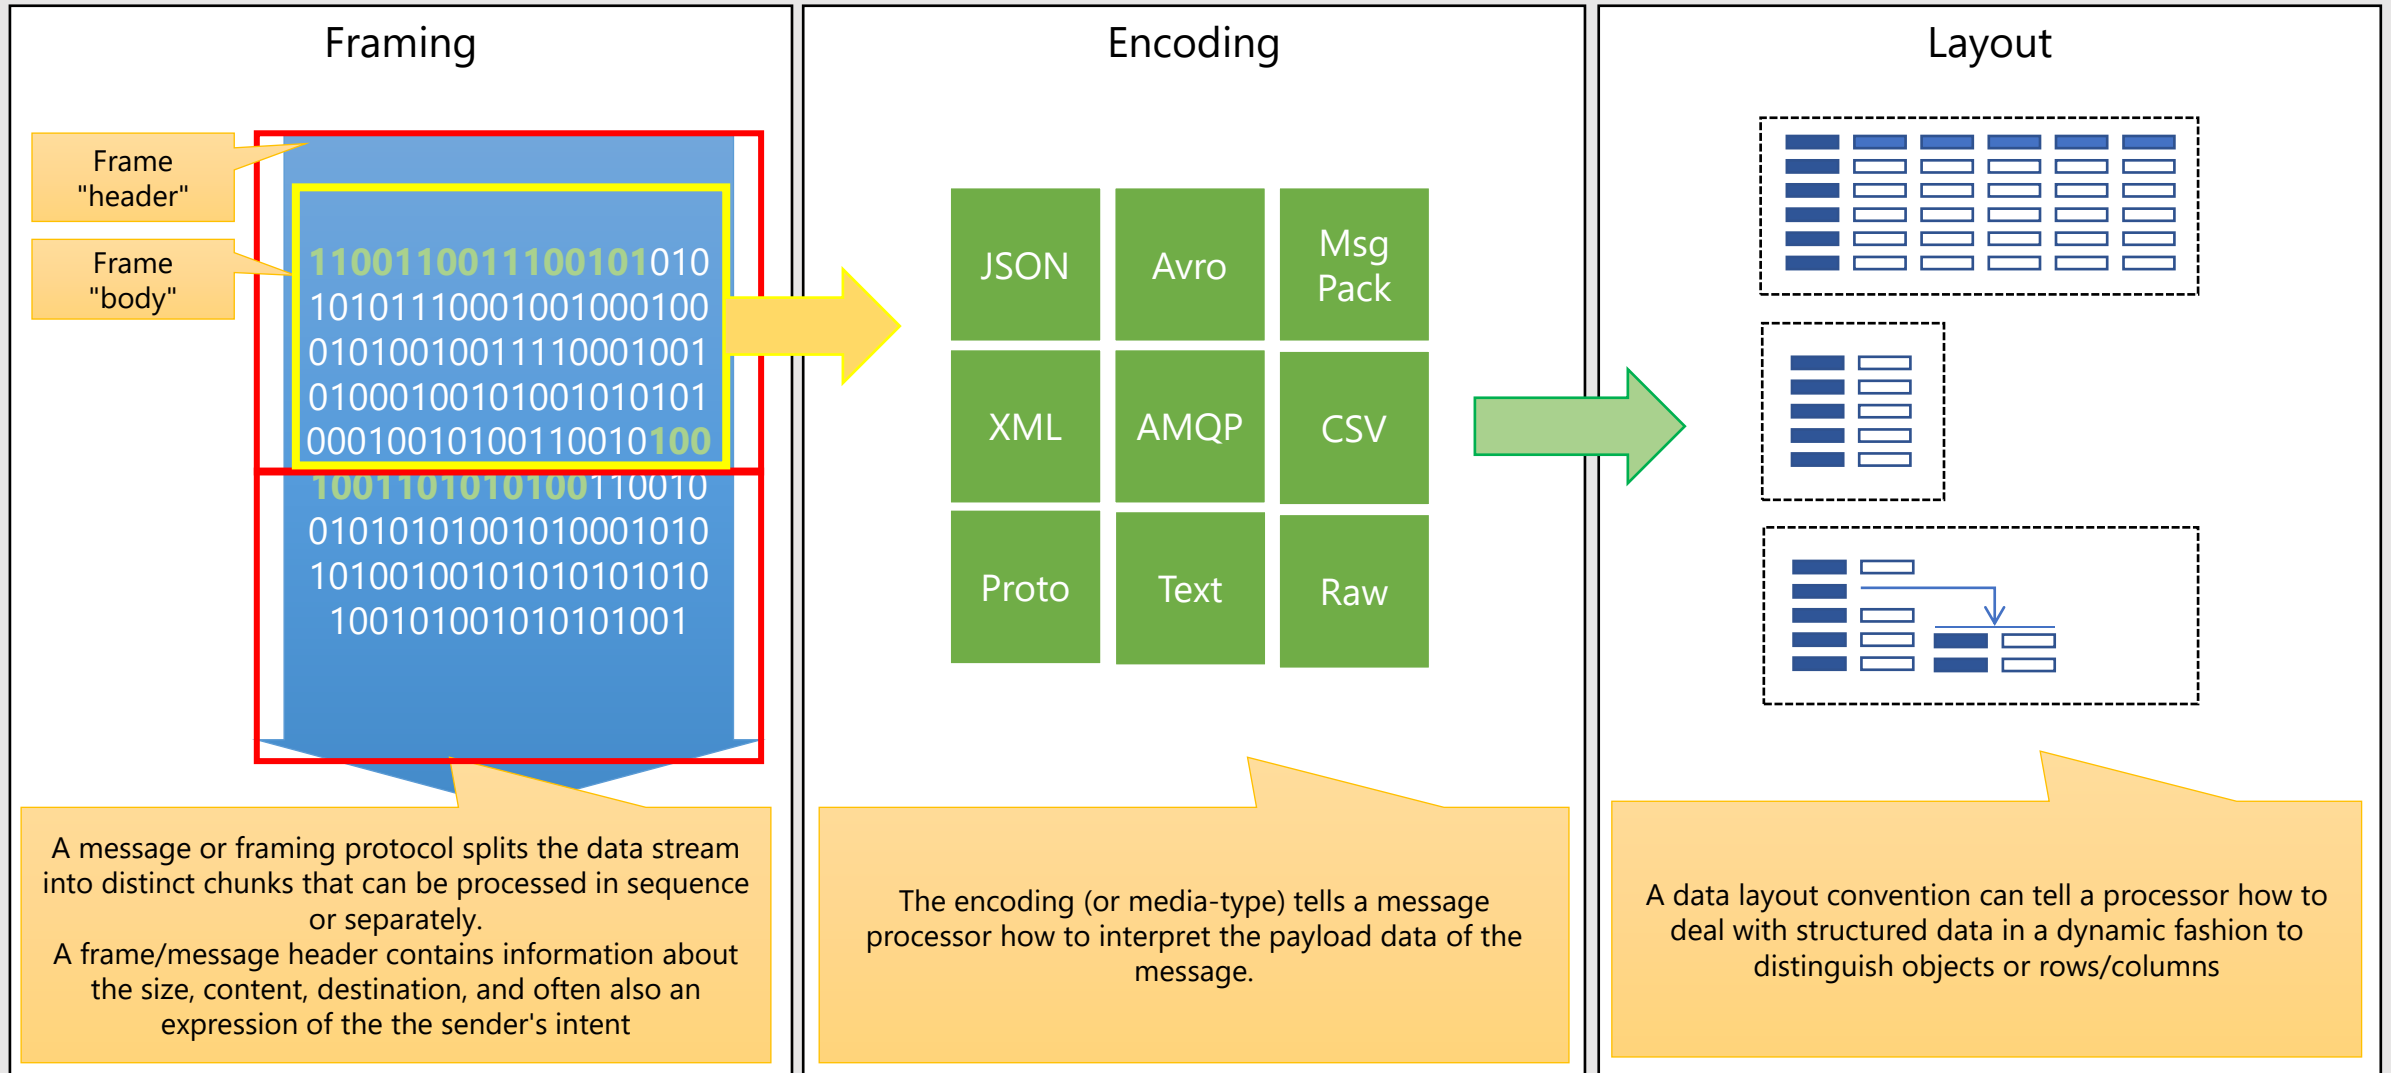
- A protocol is **symmetric** when it allows all of its supported gestures (except for connection establishment) independent of who initiated the connection.

Multiplexing

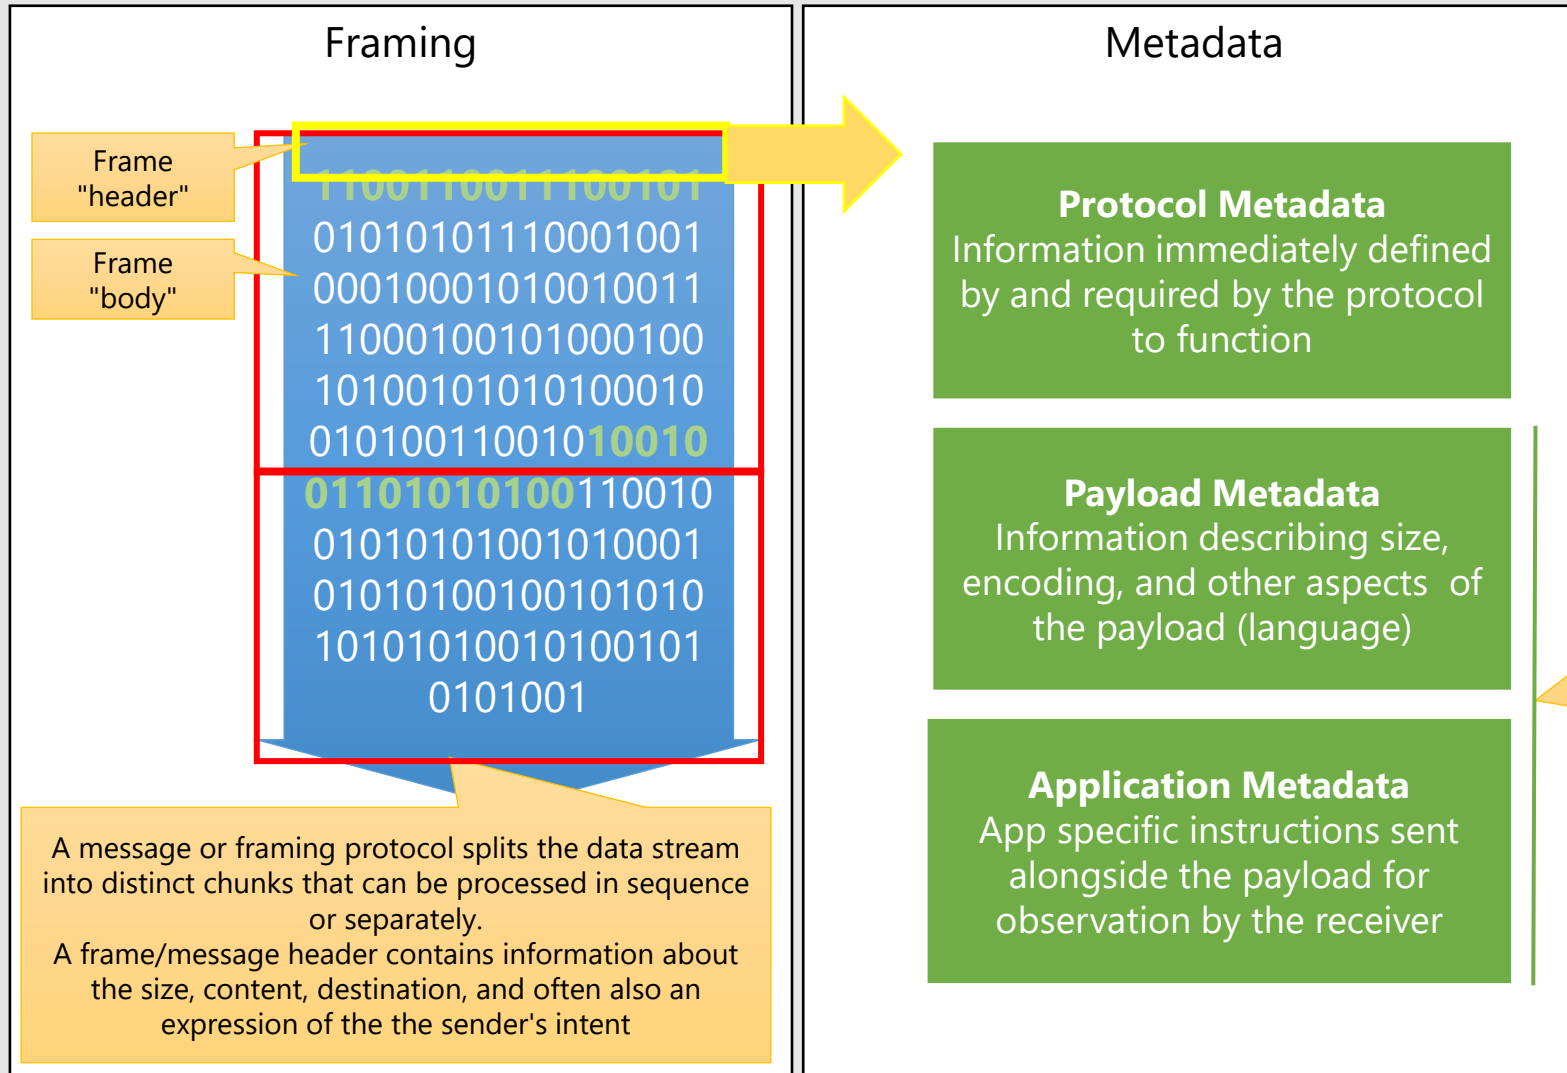


- **Multiplexing** allows a singular network connection to be used for multiple concurrent communication sessions (or links)
- Establishing connections can be enormously costly, multiplexing saves the effort for further connections between parties

Framing, Encoding, Data Layout



Metadata



- Not all protocols allow for payload and application metadata, requiring externally agreed conventions establishing mutual understanding of message content

Transfer Assurances



- Reliable protocols allow transfer of frames more reliably than underlying protocol layers
 - Compensating for data loss, preventing duplication, ensuring order
- Various strategies to compensate for data loss
 - Resend on negative acknowledgment („data didn't get here")
 - Resend on absence of acknowledgment
 - Send duplicates of frames
- **Common Transfer Assurances**
 - "Best Effort" or "At Most Once" – no resend, not reliable
 - "At Least Once" – frame is resent until it is understood that it has been delivered at least once
 - "Exactly Once" – frame is delivered exactly once [see next]

Standards

Application Protocols and APIs



Our Standardization Engagement Principles

Open door development whenever possible (ex. W3C, CNCF)

- Develop in public

- Share with public

- Work product royalty free

Consortium standards orgs case-by-case (ex. OASIS, OPC UA)

- Develop in private

- Share with public when finalized

- Work product royalty free or RAND at minimum

Standards are not useful unless they're broadly accessible to anyone interested in implementing or otherwise working with them.

Why should YOU care?

Betting on standards.

Using a common, standardized message model increases the interoperability between parts of your own applications.

The CloudEvents project already did the work of figuring out what the essential attributes of an event are and how those best map to common transports and protocols

Using standard APIs makes it easier to switch between broker products, e.g. migrate from on-premises to the cloud

Using standard protocols allows tapping into a broad ecosystem of client software for most runtimes and accommodates many frameworks.

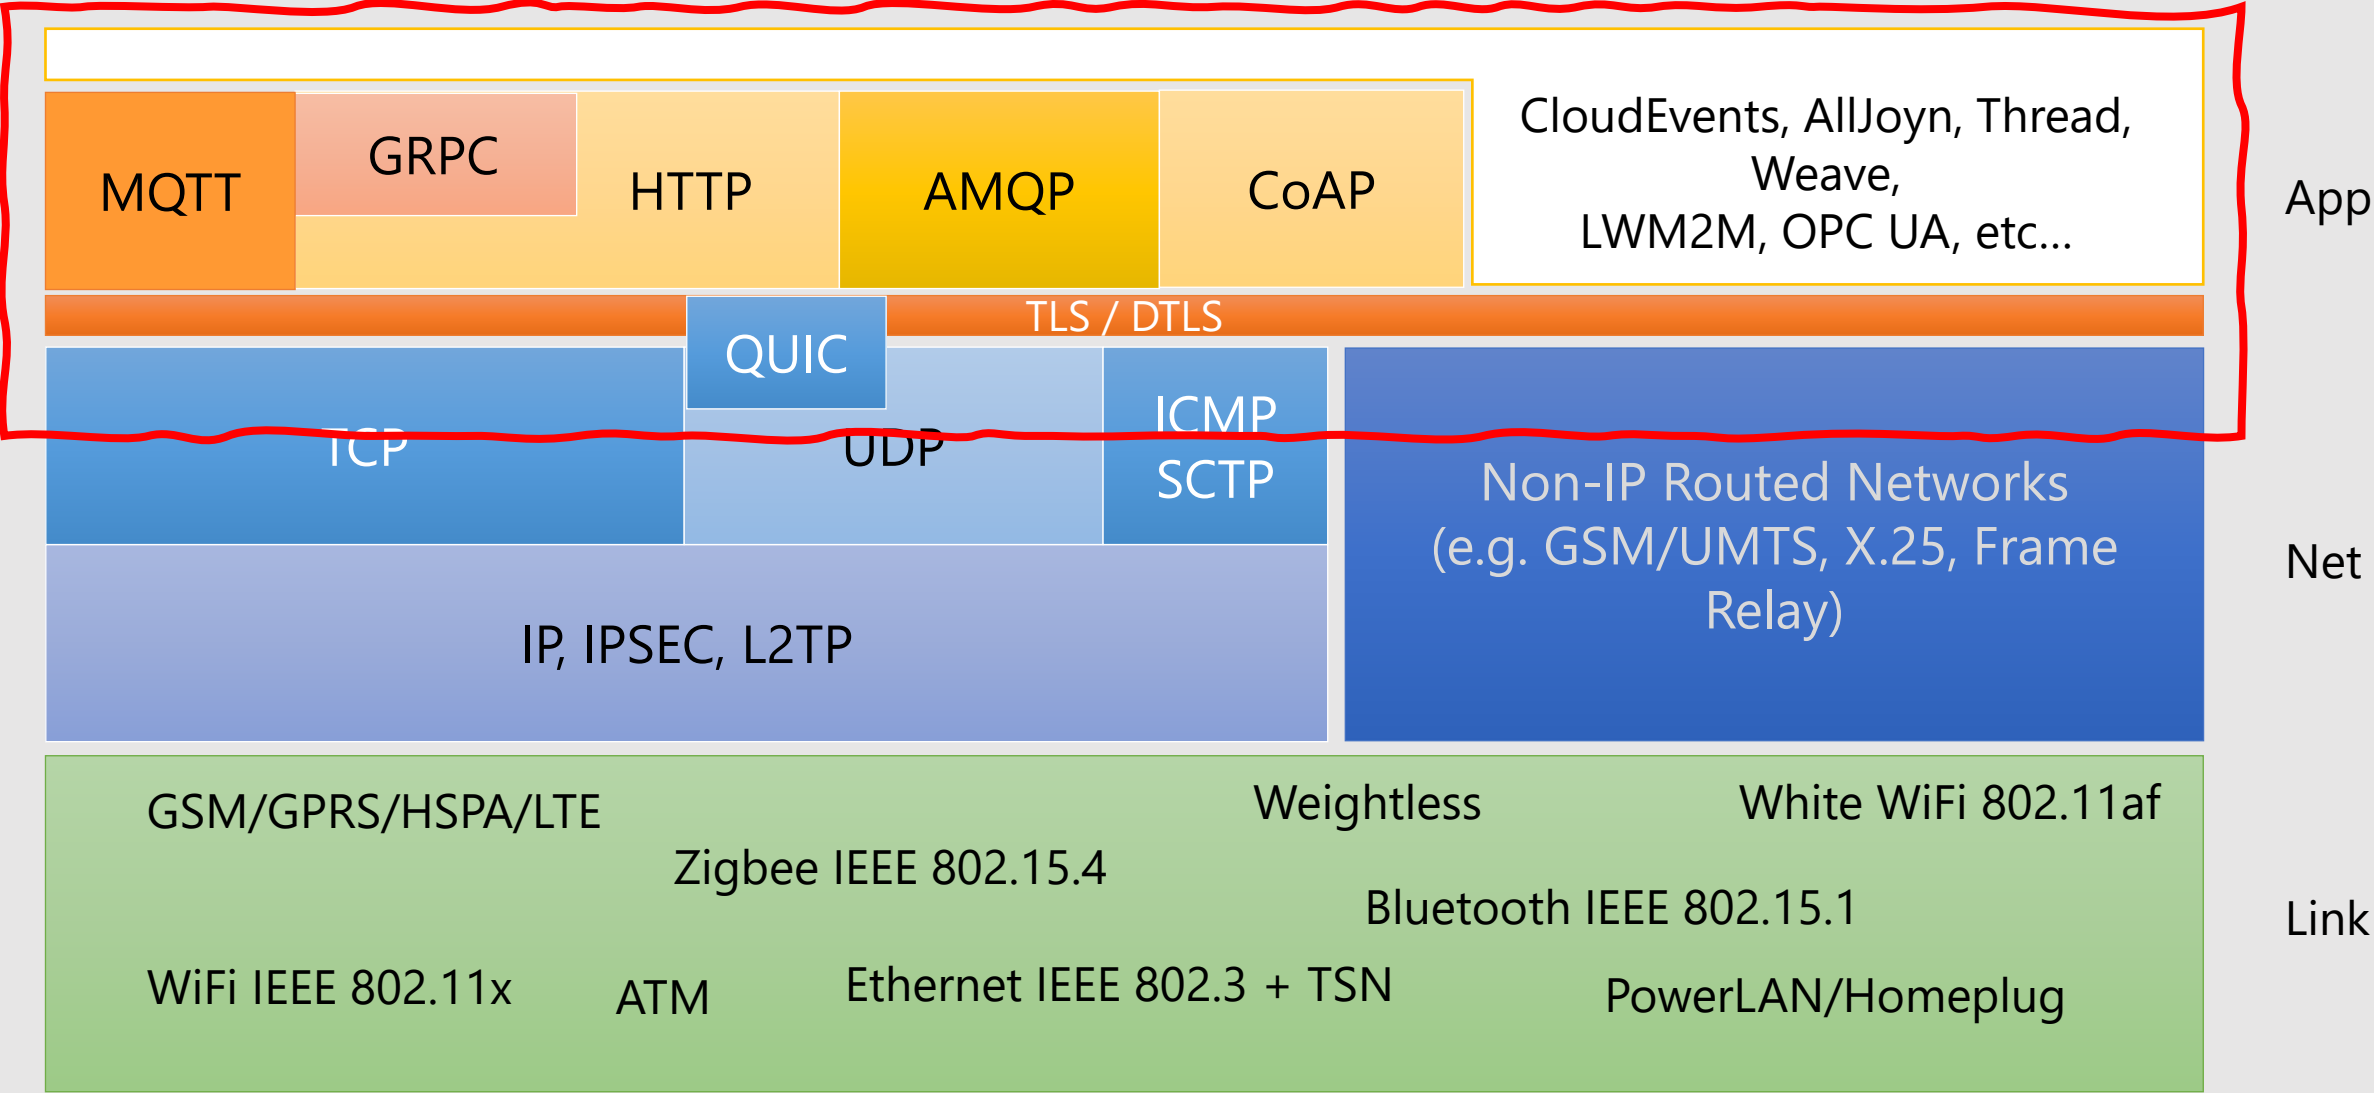
What is a Protocol?

A (network-) protocol defines rules and conventions to allow flow of information from some party to another party (and back, sometimes)

Making information flow between parties can get very complicated, so the job of making that happen is split up into layers that each focus on specific aspects, such as representing bits as radio waves.

The layers usually compose such that higher, more abstract concepts provide a choice of which lower layer to leverage, like Internet protocol routing over wired connections or radio.

Application Protocols

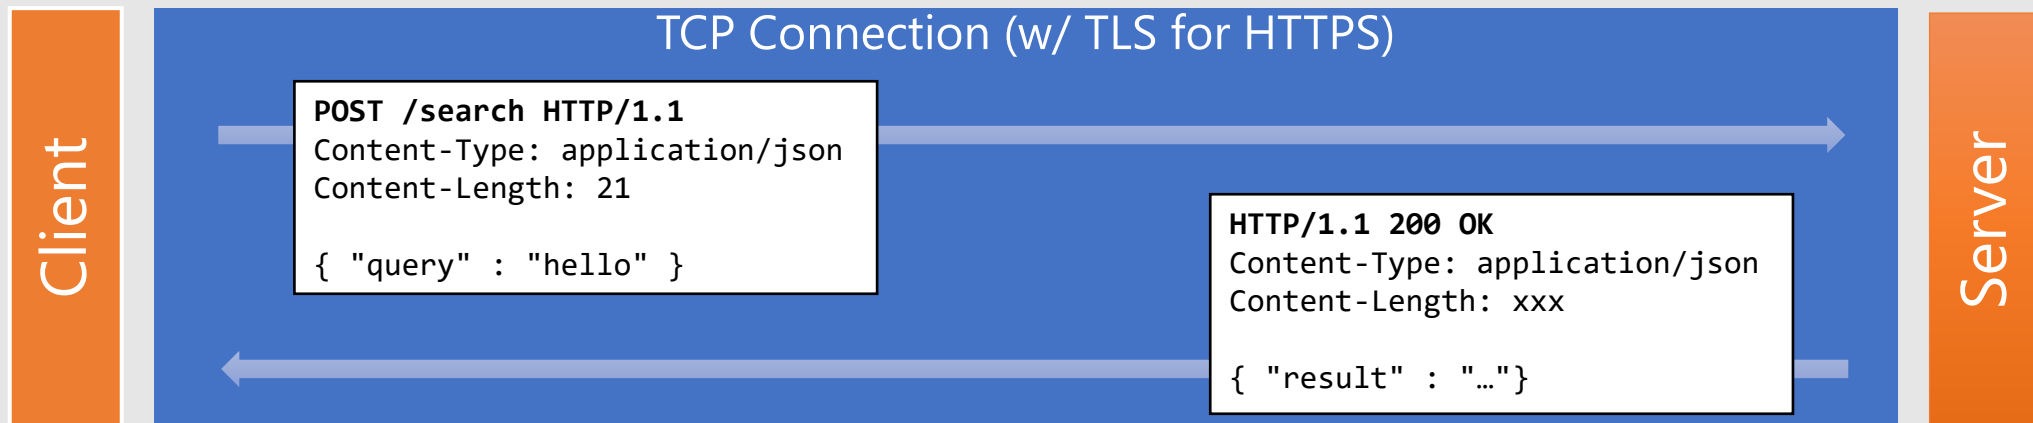


HTTP 1.1

HTTP 1.1 is the Application Protocol for the web

- Simple structure, text based, ubiquitous
- Client-initiated (asymmetric) request/response flow
- No multiplexing
- HTTP embodies the principles of "Representational State Transfer".
- REST is not a protocol, it's the architectural foundation for the WWW.

Patterns	ReqResp
Symmetric	No
Multiplexing	No
Encodings	Variable
Metadata	Yes
Assurances	-

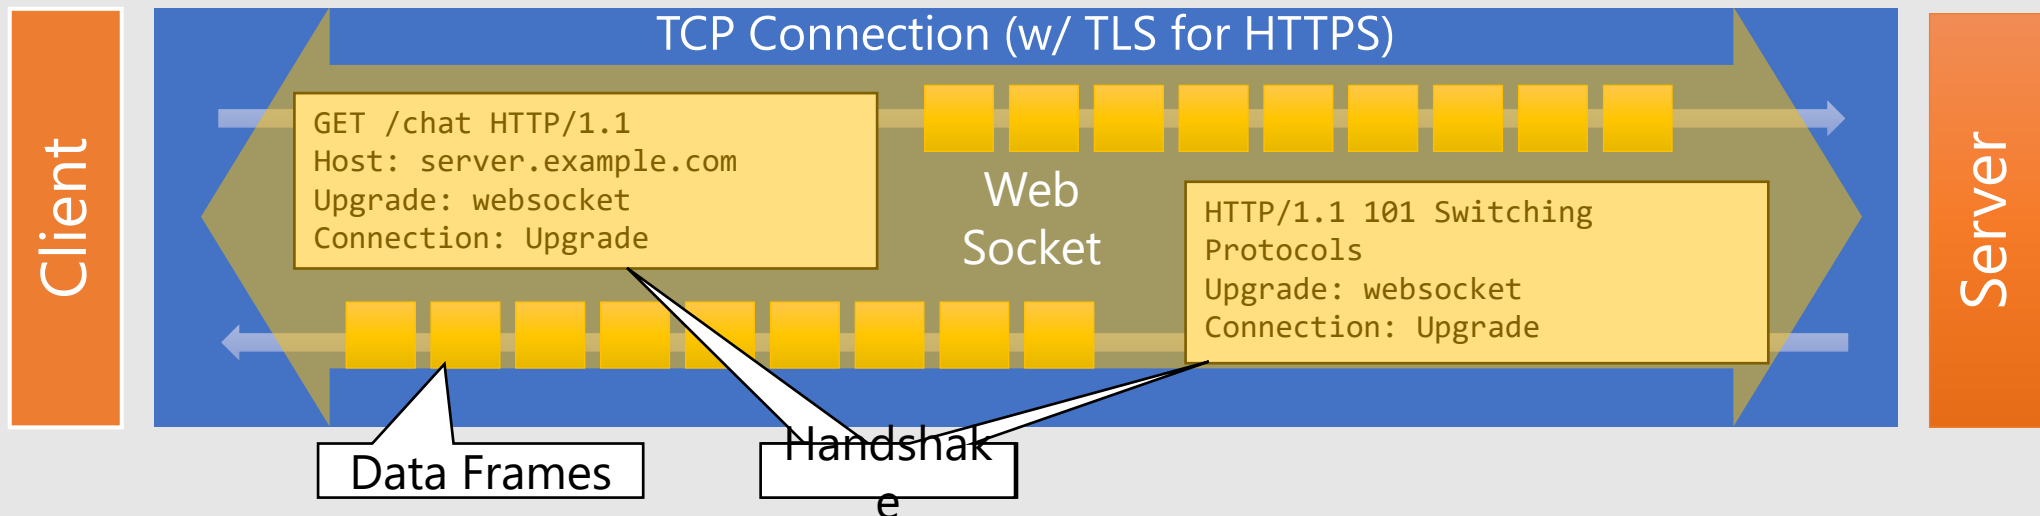


Web Sockets

Web Sockets is a Stream Tunneling Protocol

- Allows using the HTTP 1.1 port (practically only HTTPS) for bi-directional, non-HTTP stream transfer
- Web Sockets by itself is neither a Messaging or an Application Protocol, as it defines no encoding or semantics for the stream.
- Web Sockets can tunnel AMQP, MQTT, CoAP/TCP, etc.

Patterns	Duplex
Symmetric	No
Multiplexing	No
Encodings	Fixed
Metadata	No
Assurances	-

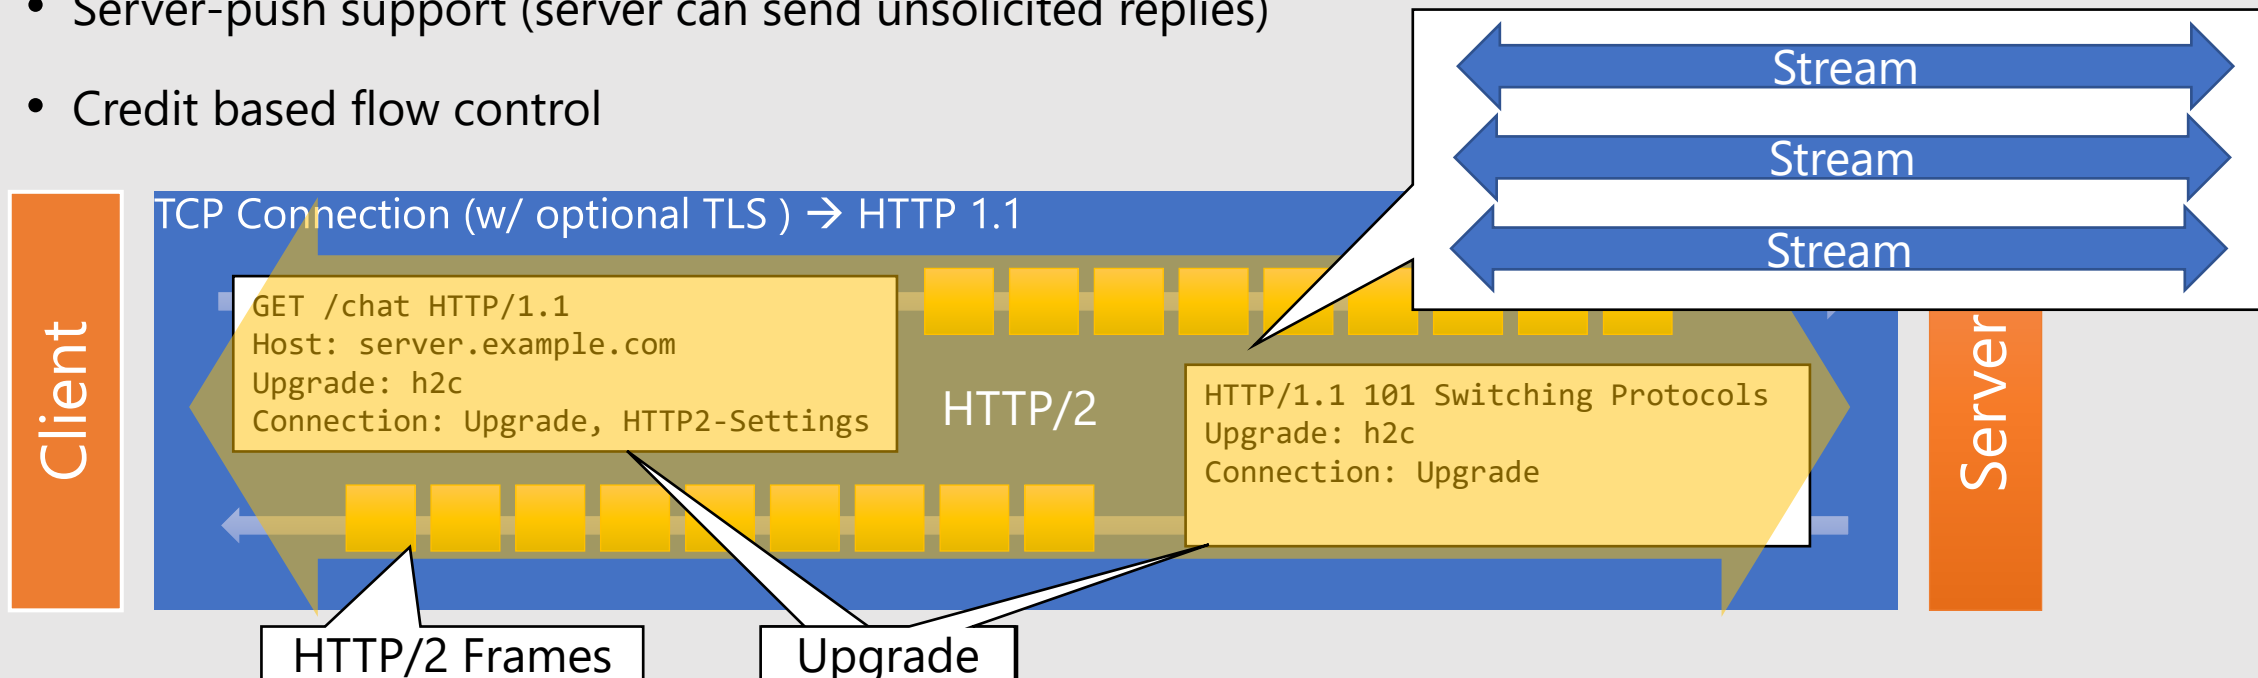


HTTP/2 (+ GRPC)

HTTP/2 is an **Application Protocol**; successor of HTTP 1.1

- Same semantics and message model, different implementation
- Multiplexing support, binary standard headers, header compression.
- Uses Web Socket like upgrade for backward compatible integration with HTTP 1.1, no WS support
- Server-push support (server can send unsolicited replies)
- Credit based flow control

Patterns	RR, OW/SC
Symmetric	No
Multiplexing	Yes
Encodings	Variable
Metadata	Yes
Assurances	-




HTTP/3

HTTP/3 is an **Application Protocol**; coexists with HTTP/2

- Same semantics and message model, different implementation
- Largely a redo of HTTP/2, moving to QUIC
- Multiplexing support via QUIC (UDP Streams), binary standard headers, header compression.
- QUIC runs over UDP port 443. TLS built in.

Patterns	RR, OW/SC
Symmetric	No
Multiplexing	Yes
Encodings	Variable
Metadata	Yes
Assurances	-




OASIS MQTT

MQTT is a lightweight **Publish and Subscribe Protocol**

- Easy to implement for publishers and subscribers, assumes broker
- Rigorously optimized for minimizing wire overhead
- Publish/Subscribe gestures are explicit elements of the protocol; “subscribe” = “receive”
- Often used for submitting and subscribing to telemetry and sharing state changes amongst peers, can model request/reply routes on top

MQTT 3.1.1 most used

MQTT 5.0 new with many improvements



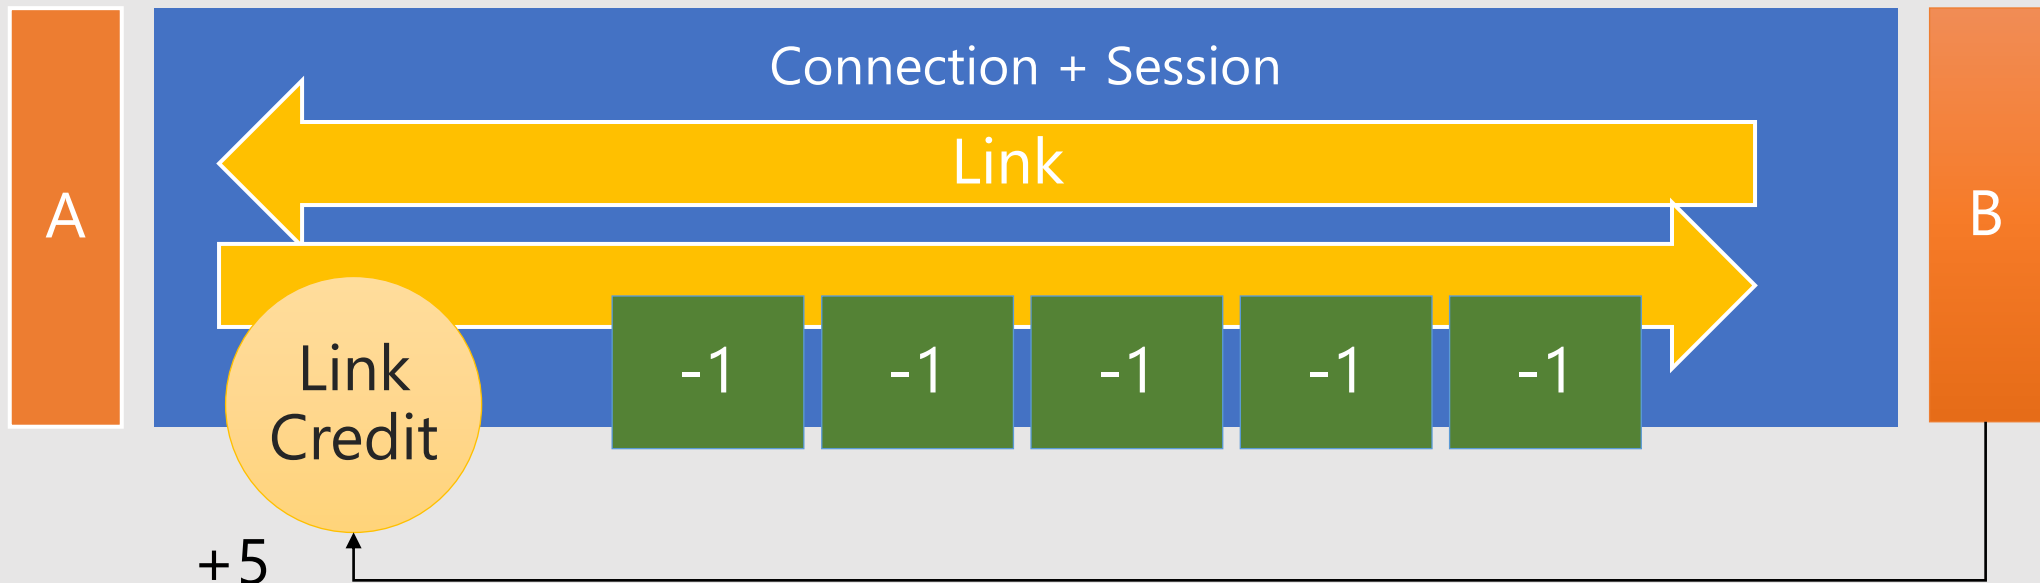
OASIS AMQP

Advanced Message Queuing Protocol

AMQP 1.0 is a symmetric, reliable **Message Transfer Protocol** with support for multiplexing and flow control

- Supports queuing, pub/sub, filters, one-way, request-response, streams
- No topology assumptions, multi-hop routing facilities

AMQP 0.9 (what RabbitMQ uses) is an expired draft with a fixed topology model



Azure Compatible, Generic AMQP 1.0 Clients

Client Stack	URL
Apache Qpid Proton-C (Go, Python, Ruby, C+ +)	https://qpid.apache.org/proton/
Apache Qpid Messaging API	https://qpid.apache.org/components/messaging-api
Apache Qpid Proton-J	https://qpid.apache.org/proton/
Apache NMS AMQP	https://activemq.apache.org/components/nms/providers/amqp/
AMQP .NET Lite (all variants of .NET, incl. Nano & Micro)	https://github.com/Azure/amqpnetlite
Azure AMQP (our own server stack)	https://github.com/Azure/azure-amqp
Azure uAMQP C (Python, PHP)	https://github.com/Azure/azure-uamqp-c https://github.com/Azure/azure-uamqp-python
Rhea (NodeJS)	https://github.com/amqp/rhea
Go AMQP	https://github.com/Azure/go-amqp
Vert.X AMQP Client	https://vertx.io/docs/vertx-amqp-client/java/

Apache Kafka

The Apache Kafka protocol is a project-specific request/response protocol for the Apache Kafka broker

- SASL preamble for authentication
- API keys identify operations, specific message types per key
- Effectively an RPC protocol tailored to Kafka's features
- Not a standard and fully controlled by the Kafka project, but quite popular

Patterns	Kafka
Symmetric	No
Multiplexing	No
Encodings	Variable
Metadata	Yes
Assurances	ALO

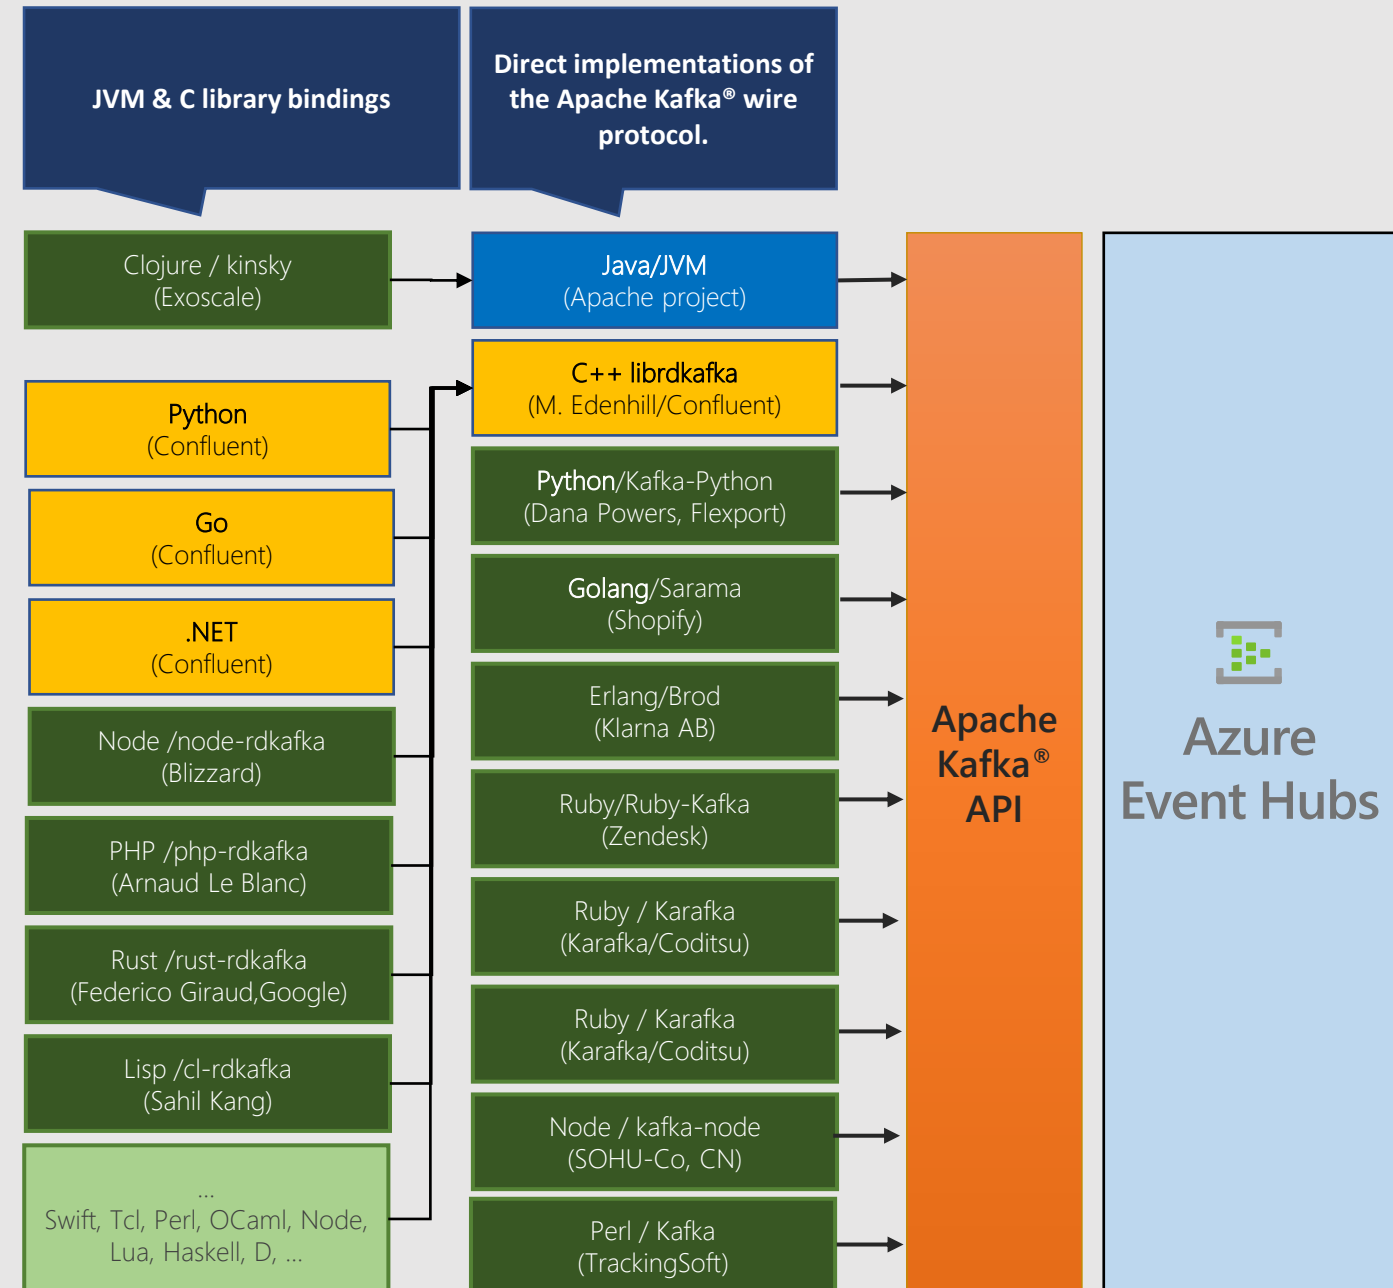


Apache Kafka® Clients

<https://cwiki.apache.org/confluence/pages/viewpage.action?pageId=30736606>

Azure Event Hubs is compatible with the producer and consumer APIs of all current Apache Kafka clients.

Only the Java/JVM client is part of the Apache Kafka project. All other clients are under various non-foundation ownerships and licenses. Azure does not provide QFE support for any of these clients but does provide support for service-side compatibility.



Apache Software Foundation

Non-Foundation, Community
/ Best Effort Support

Non-Foundation, Commercial
Support Option

Unifying Abstractions

Java JMS 2.0 (Jakarta Messaging)

An exceptionally successful Java/JVM API standard that works with most mature, queue-based pubsub message brokers on the market today.

```
Topic topic = session.createTopic(dest.toString());
MessageProducer sender = session.createProducer(dest);
TopicSubscriber topicSubscriber1 =
    session.createDurableSubscriber(topic, sub1Name);
TopicSubscriber topicSubscriber2 =
    session.createDurableSubscriber(topic, sub2Name, "JMSCorrelationID='5'", false);
MessageConsumer topicSubscriber3 =
    session.createSharedDurableConsumer(topic, sub3Name, "JMSCorrelationID='5'");
```

```
TextMessage message = session.createTextMessage("Text!");
message.setJMSCorrelationID(Integer.toString(i));
sender.send(message, DELIVERY_MODE, Message.DEFAULT_PRIORITY,
Message.DEFAULT_TIME_TO_LIVE);
```

- Queues
- Topics
- JMS Message Selectors
- Shared Durable Subscriptions
- Unshared Durable Subscriptions
- Shared Volatile Subscriptions
- Unshared Volatile Subscriptions
- Queue Browser
- Queue Message Selectors
- Temporary Queues
- Temporary Topics
- JMS Transactions

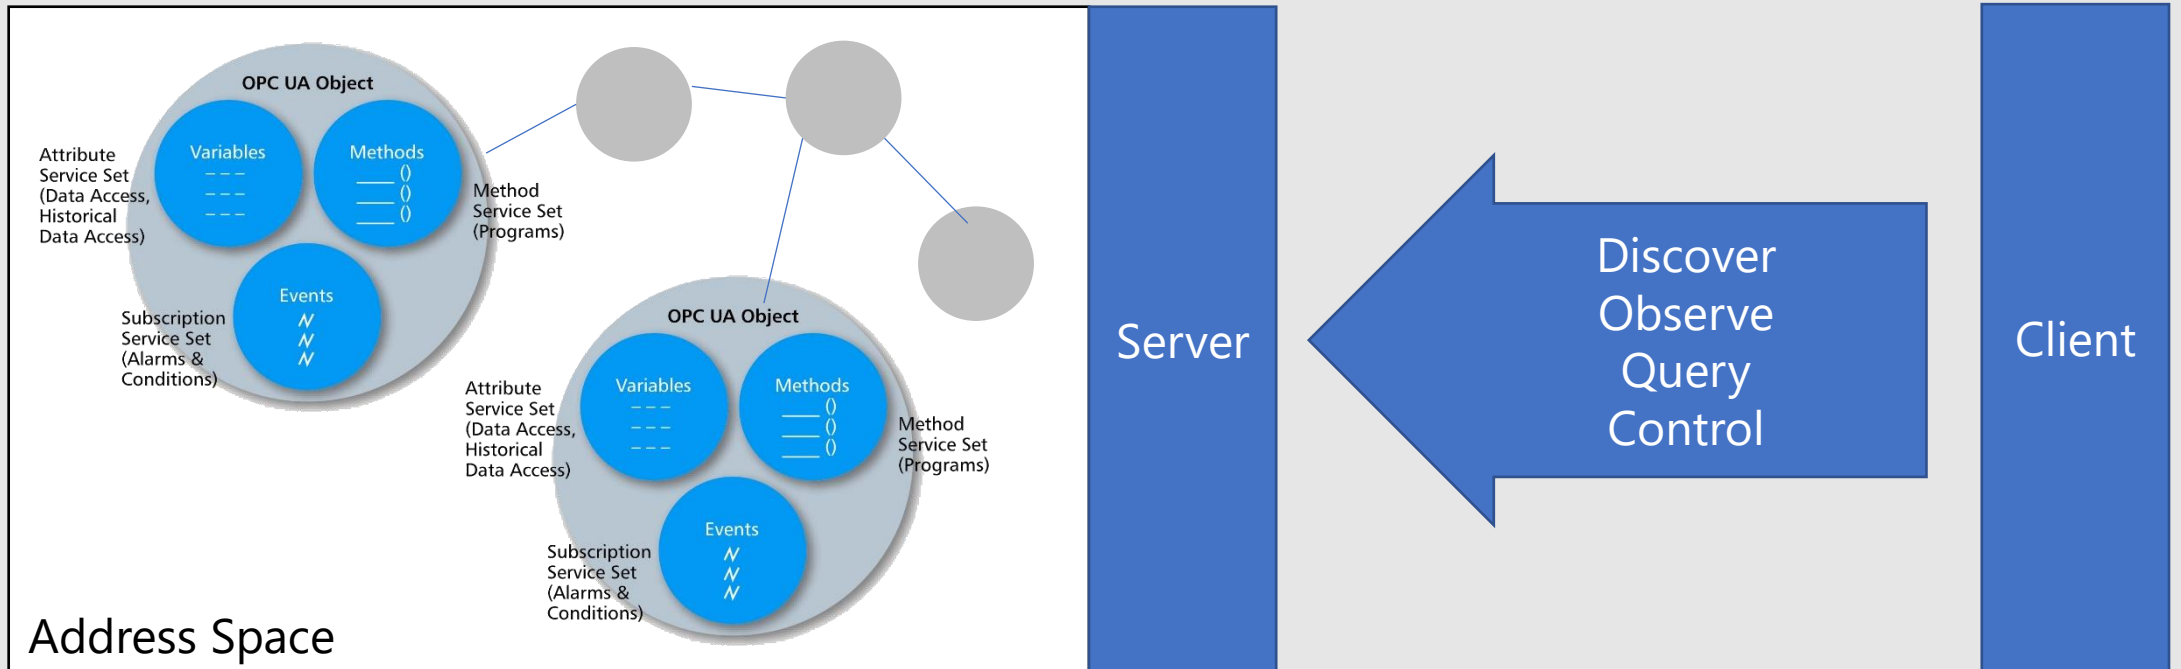
OPC UA

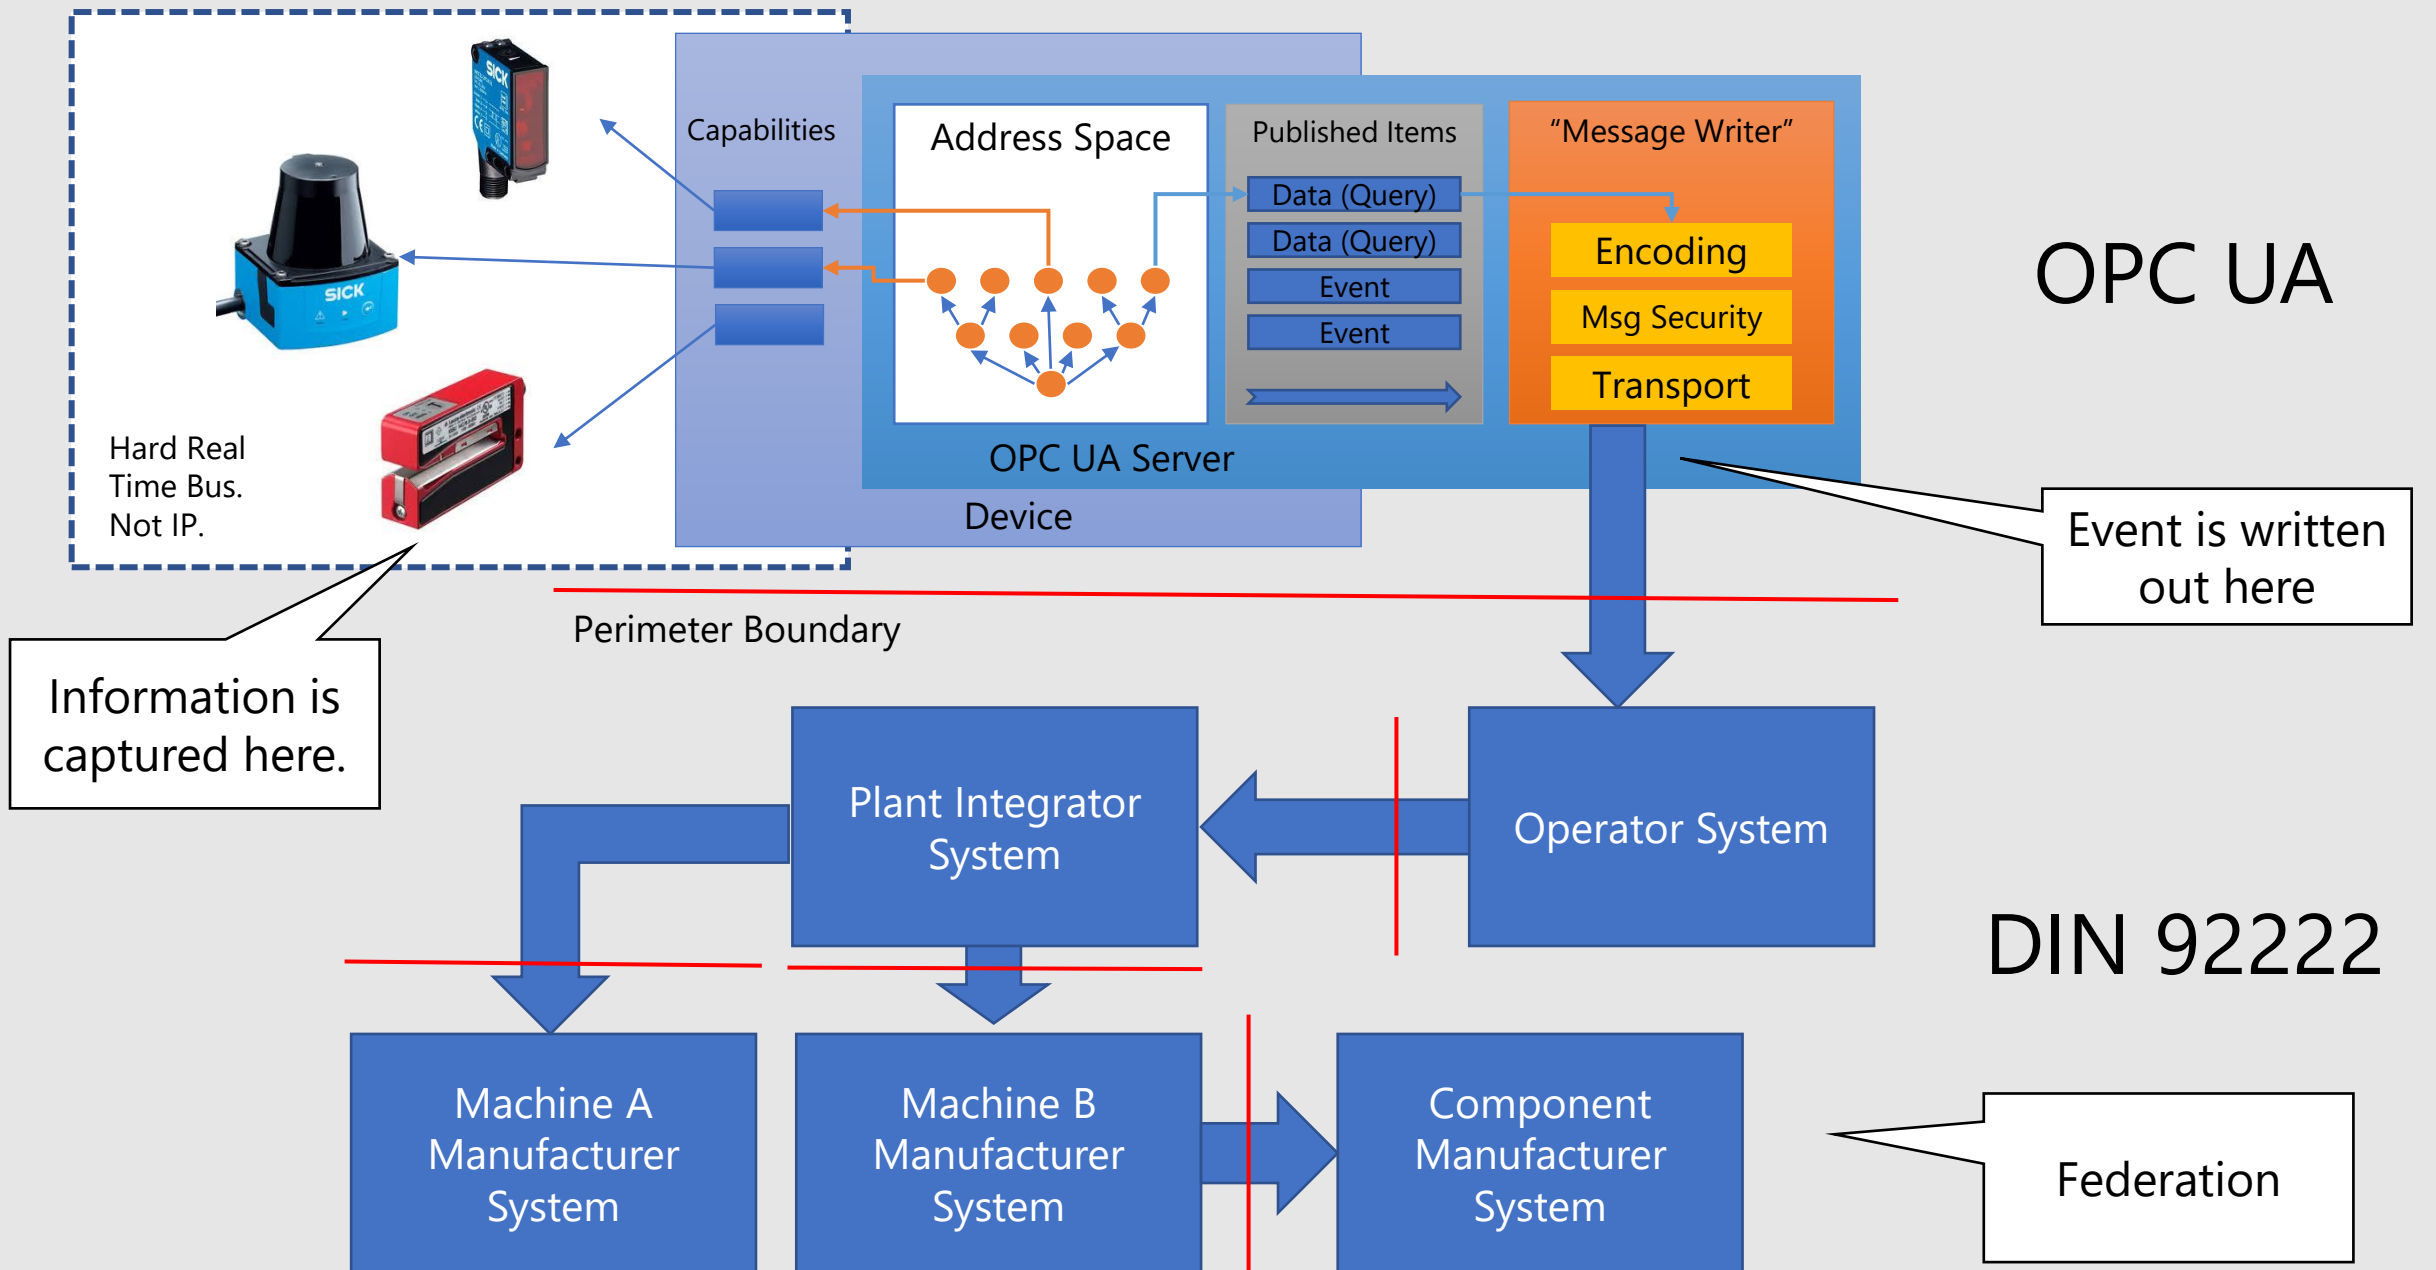
OPC Foundation Unified Architecture

OPC Foundation standard; IEC standard

Foundational architecture model for (industrial) device management and information flow

Cloud integration via Pub/Sub





CNCF CloudEvents

Event Protocol Suite developed in CNCF Serverless WG

Common metadata attributes for events

Flexibility to innovate on event semantics

Simple abstract type system mappable to different encodings

Transport options

HTTP(S) 1.1 Webhooks, also HTTP/2

MQTT 3.1.1 and 5.0

AMQP 1.0

Kafka, NATS

Encoding options

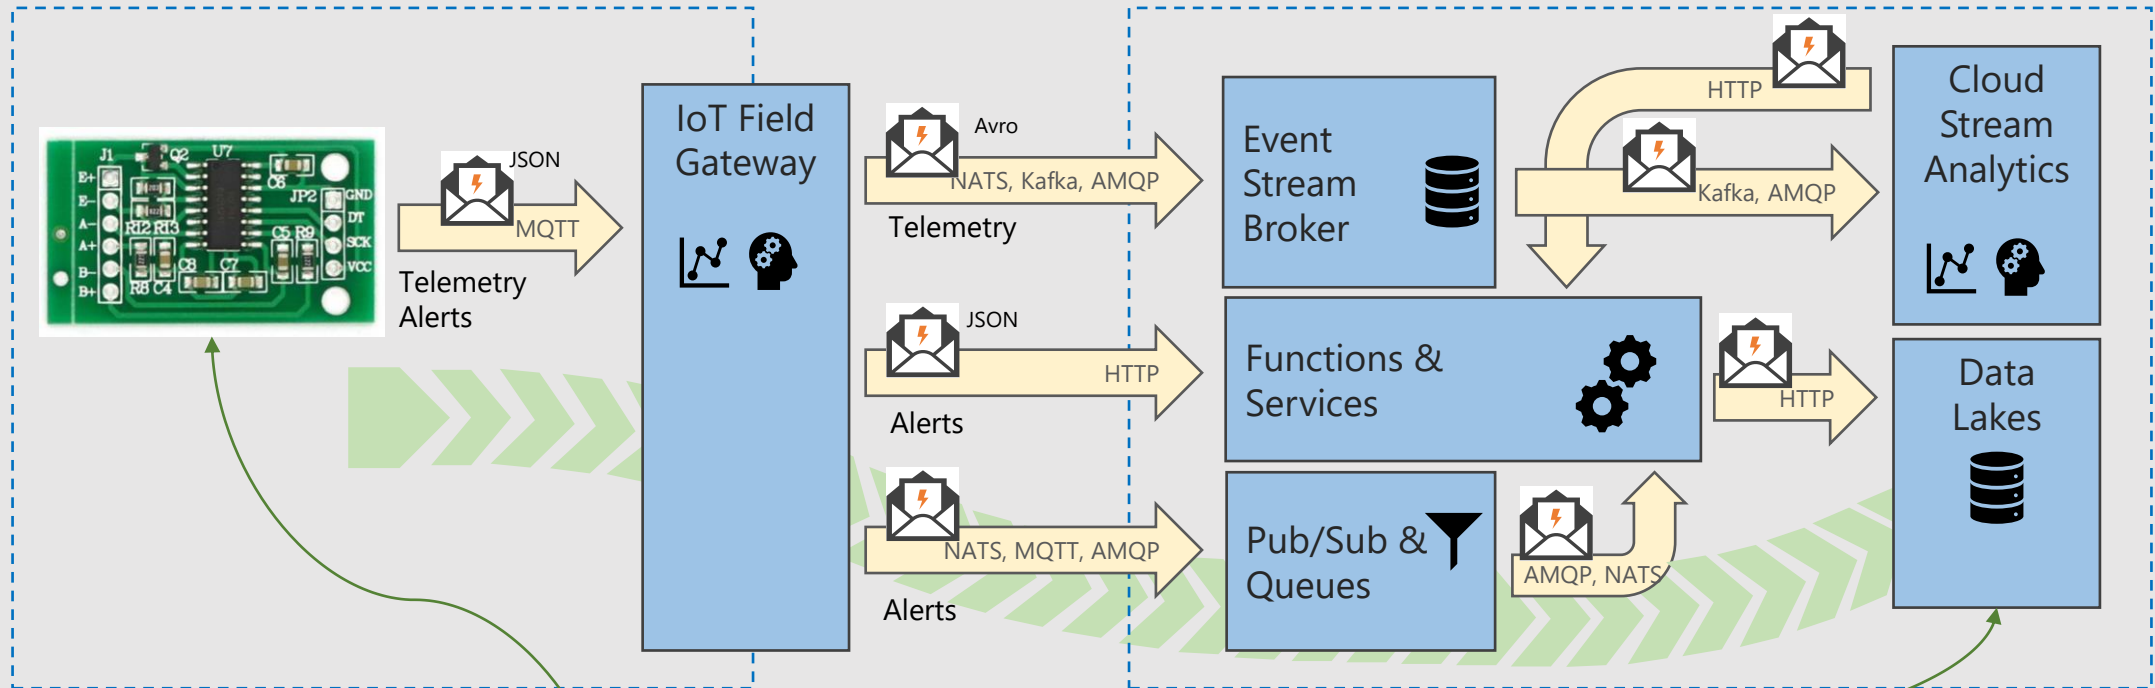
JSON (required for all implementations)

Extensible for binary encodings: Avro, AMQP, etc.



cloudevents

Why CloudEvents?



- Event data is often routed via multiple hops and often using different protocols
- How is what gets sent here easily routed to and stored here in hybrid edge/cloud and multi-cloud systems?

Why CloudEvents?

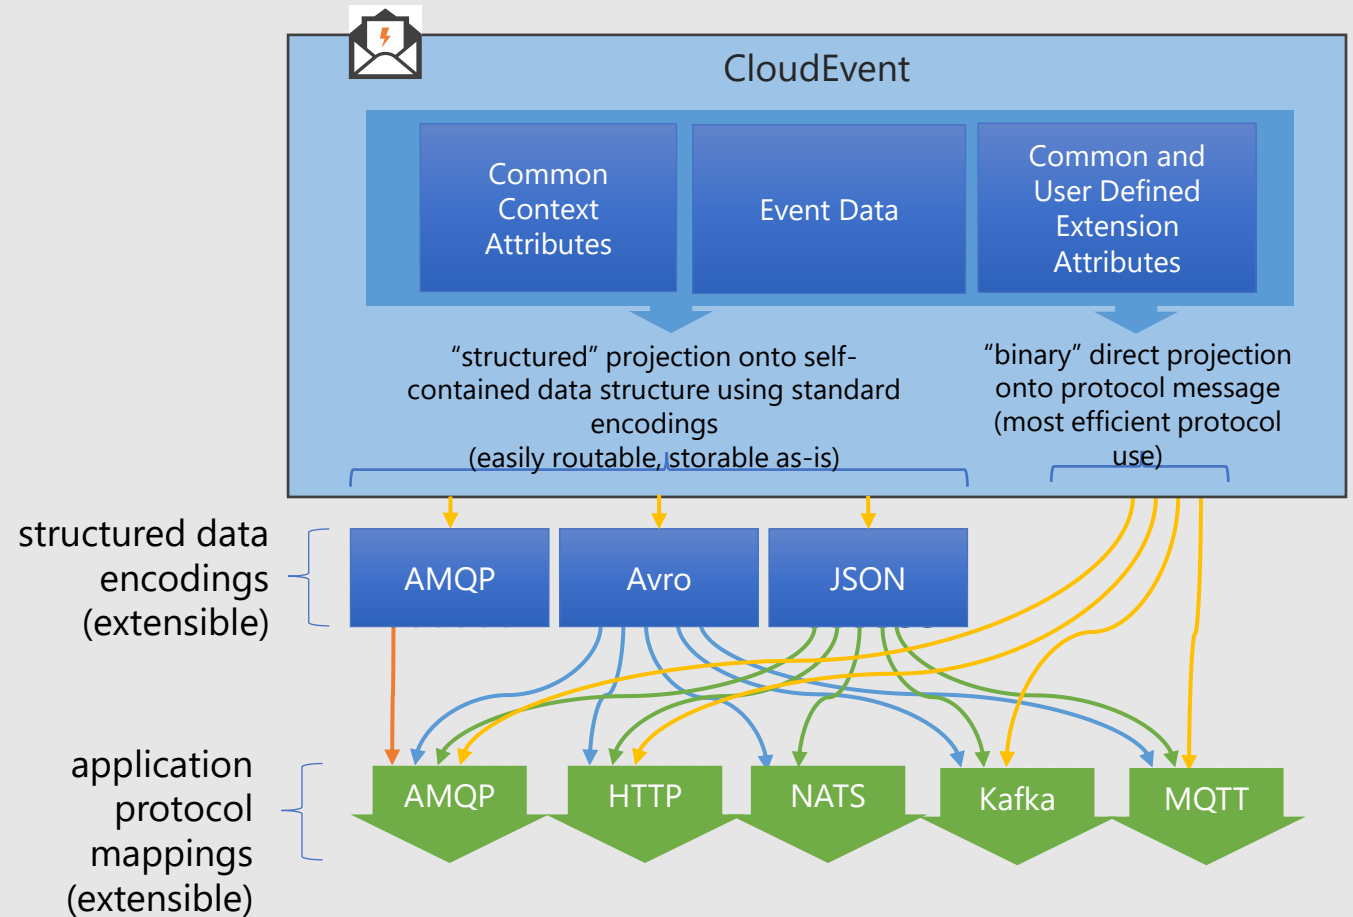
Binds to existing standard application protocols

Does not try to abstract away protocols but leverages each for its strengths

Integrates with existing messaging and eventing stacks

Leverages existing data encodings and is easy to adapt to new ones (Protobuf, CBOR, MsgPack, etc.)

Allows for protocol switching and transcoding on multi-hop routes



CloudEvents - Base Specification

CloudEvents is a lightweight common convention for events.

It's *intentionally* not a messaging model* to keep complexity low.

No reply-path indicators, no message-to-message correlation, no target address indicators, no command verbs/methods.

Metadata for handling of events by generic middleware and/or dispatchers

What kind of event is it? **type, specversion**

When was it sent? **time**

What context was it sent out of? **source, subject**

What is this event's unique identifier? **id**

What's the shape of the carried event data? **datacontenttype, schema**

Event data may be text-based (esp. JSON) or using some binary encoding



Event Formats

Event formats bind the abstract CloudEvents information model to specific wire encodings.

All implementation must support JSON. Avro is a supported binary format.

Further compact binary event format candidates might be CBOR, or Protobuf.

```
{  
  "specversion" : "1.0",  
  "type" : "myevent",  
  "source" : "uri:example-com:mydevice",  
  "id" : "A234-1234-1234",  
  "time" : "2018-04-05T17:31:00Z",  
  "datacontenttype" : "text/plain",  
  "data" : "Hello"  
}
```

JSON Representation

Example

HTTP - Binary

```
POST /event HTTP/1.0
Host: example.com
Content-Type: application/json
ce-specversion: 1.0
ce-type: com.bigco.newItem
ce-source: http://bigco.com/repo
ce-id: 610b6dd4-c85d-417b-b58f-3771e532
```

```
{
  "action": "newItem",
  "itemID": "93"
}
```

HTTP - Structured

```
POST /event HTTP/1.0
Host: example.com
Content-Type: application/cloudevents+json
```

```
{
  "specversion": "1.0",
  "type": "com.bigco.newItem",
  "source": "http://bigco.com/repo",
  "id": "610b6dd4-c85d-417b-b58f-3771e532",
  "datacontenttype": "application/json",
  "data": {
    "action": "newItem",
    "itemID": "93"
  }
}
```

Transport Bindings

HTTP 1.1, HTTP/2, HTTP/3:

Binds to the HTTP message

Binary and structured modes

AMQP:

Binds event to the AMQP message

Binary and structured modes

MQTT:

Binds event to MQTT PUBLISH frame.

Binary and Structured for MQTT v5

Structured mode only for MQTT v3.1.1

NATS:

Binds event to the NATS message.

Structured mode only

Apache Kafka:

Binds to the Kafka message

Structured and binary mode

Protocol bindings directly map onto the protocol's message structure and using protocol semantics. Accepts that protocols are different.

Binary mode: Event metadata projected onto the protocol message metadata, event data onto the protocol message payload

Structured mode: Event is self-contained as an encoded byte stream, metadata may be promoted (duplicated) into protocol message metadata.

Discovery, Catalog, Subscriptions

Cloud Events Futures





Service Bus & Azure Queues

Cloud messaging

- AMQP
- HTTP
- Cloud-Events
- JMS 2.0



Event Hubs

Telemetry stream ingestion


- AMQP
- HTTP
- Kafka
- Cloud-Events



Event Grid

Event distribution

- AMQP
- HTTP
- Cloud-Events



IoT Hub

IoT messaging and management

- AMQP
- MQTT
- OPCUA
- Cloud-Events



Relay

Discovery, Firewall/NAT Traversal

- HTTP
- Web-Sockets
- Cloud-Events

