Microsoft

Azure

Azure Event Hubs

The Architect's Cut

Clemens Vasters

Chief Messenger, **Twitter: @clemensv**

Patterns

Features

Architecture

Internals

Agenda

A *very* brief overview: “What is Azure Event Hubs?”

The Event Hubs Deep Dive

Clusters and Namespaces
Networking Capabilities
Authorization Models
Partitions and Log Storage
Message Broker Compute Model
Federation and Replication
Throughput and Latency
Observed Availability and Reliability

Q&A

What is Azure Event Hubs?

Platform-as-a-Service Event Stream Broker

Use the Apache Kafka® API, but with far lower cost and better performance.

Fully managed: You use the features, Azure deals with everything else

AMQP 1.0 standards compliant, Apache Kafka® wire-compatible

Polyglot Azure SDK and cross-platform client support

Industry-leading reliability and availability

Fast.

Azure Event Hubs vs. Apache Kafka®

Similar yet very different

	Azure Event Hubs	Apache Kafka
User Model	Partitioned event stream broker with high-availability replication	Partitioned event stream broker with high-availability replication
Architecture	Multi-tenant, 3-Tier Gateway/Broker/Storage cluster model, with tenant-isolation, all tiers independently scalable	Single-tenant monolith. Need to increase broker instances in a cluster to scale any dimension.
Implementation Language	C# and Native (C/C++)	Java
Cluster Manager	Azure Service Fabric (inline)	Apache Zookeeper (external); KRaft (inline, experimental)
Partition Mapping	Key hashing, client or server-side mapping of events	Key hashing, client-side mapping of events
Consumer Partition Ownership Coordination	Server-coordinated partition ownership (Kafka), client-coordinated ownership with external leader election. Parallel, direct partition reads.	Server-coordinated partition ownership
Server Workload Balancing	Dynamic and fully automated (100% hands-off). Broker resource allocation independent of partition count or ownership, flexible scaling.	Static assignment of partitions to broker instances requiring operator intervention for rebalancing.
Storage Model	Replicated log store, synchronous per-message flush-to-disk on all replicas	Replicated log store, asynchronous flush-to-disk controlled by host file system write cache settings.
Networking	Single endpoint access to all partitions, Public IP/DNS or Virtual Networking, Firewall.	Endpoint per broker instance. Multiple IPs required. Complex network management required.
Access Control	Token-based access policy model with unlimited publisher policies, Azure Active Directory role-based access control	Local accounts, federation extensibility.
Protocols	AMQP 1.0 (optional: WebSockets) HTTPS 1.1 Apache Kafka RPC	Apache Kafka RPC
Batching / Archives	Avro-packaged batch-packaging and archival to blob store	
Schema Registry	Schema Registry based on open CNCF Schema Registry API	(Proprietary from commercial vendors)

SKUs and Pricing

Azure Event Hubs *Standard*:

Shared-resources model

Blended capacity and consumption-based pricing.

Up to 40 MB/sec ingress:

- Throughput Units: 1TU = min(1000 msg/s, 1MB/s in), min(2000 msg/s, 2MB/s out)
- Capture
- First 12M events free, 12-25M $x/1M, 25-50M $x/1M, 50+ $x/1M

Azure Event Hubs *Premium*:

Isolated-resources model

Capacity-based pricing model. No consumption charges.

Up to 120 MB/sec ingress:

- Premium Units: 1PU = ~5-10MB/sec throughput, no per-event charge, includes Capture

Azure Event Hubs *Dedicated*:

Dedicated-resources model

Capacity-based pricing model. No consumption charges. Full isolated cluster.

Beyond 120MB/sec ingress (>4 GB/sec in prod)

- Capacity Unit: 1CU $4995/mo.
- 8CU required for availbility zones.

We will discuss all these versions today and what the units mean

You're running Apache Kafka® in your own datacenter?

Azure Event Hubs

Azure Event Hubs is

more reliable and dependable with higher 24/365 uptime
(no service windows or weekend downtime)

far less costly to own
(no initial/recurring licensing fees as with commercial Kafka distributions)

far less costly to operate
(fully managed, no hardware)

just a quick network hop away from your existing workloads

easy to integrate into your network via VPN Gateway or Express Route

Microsoft Azure

Event Hubs Namespaces: Virtual Brokers

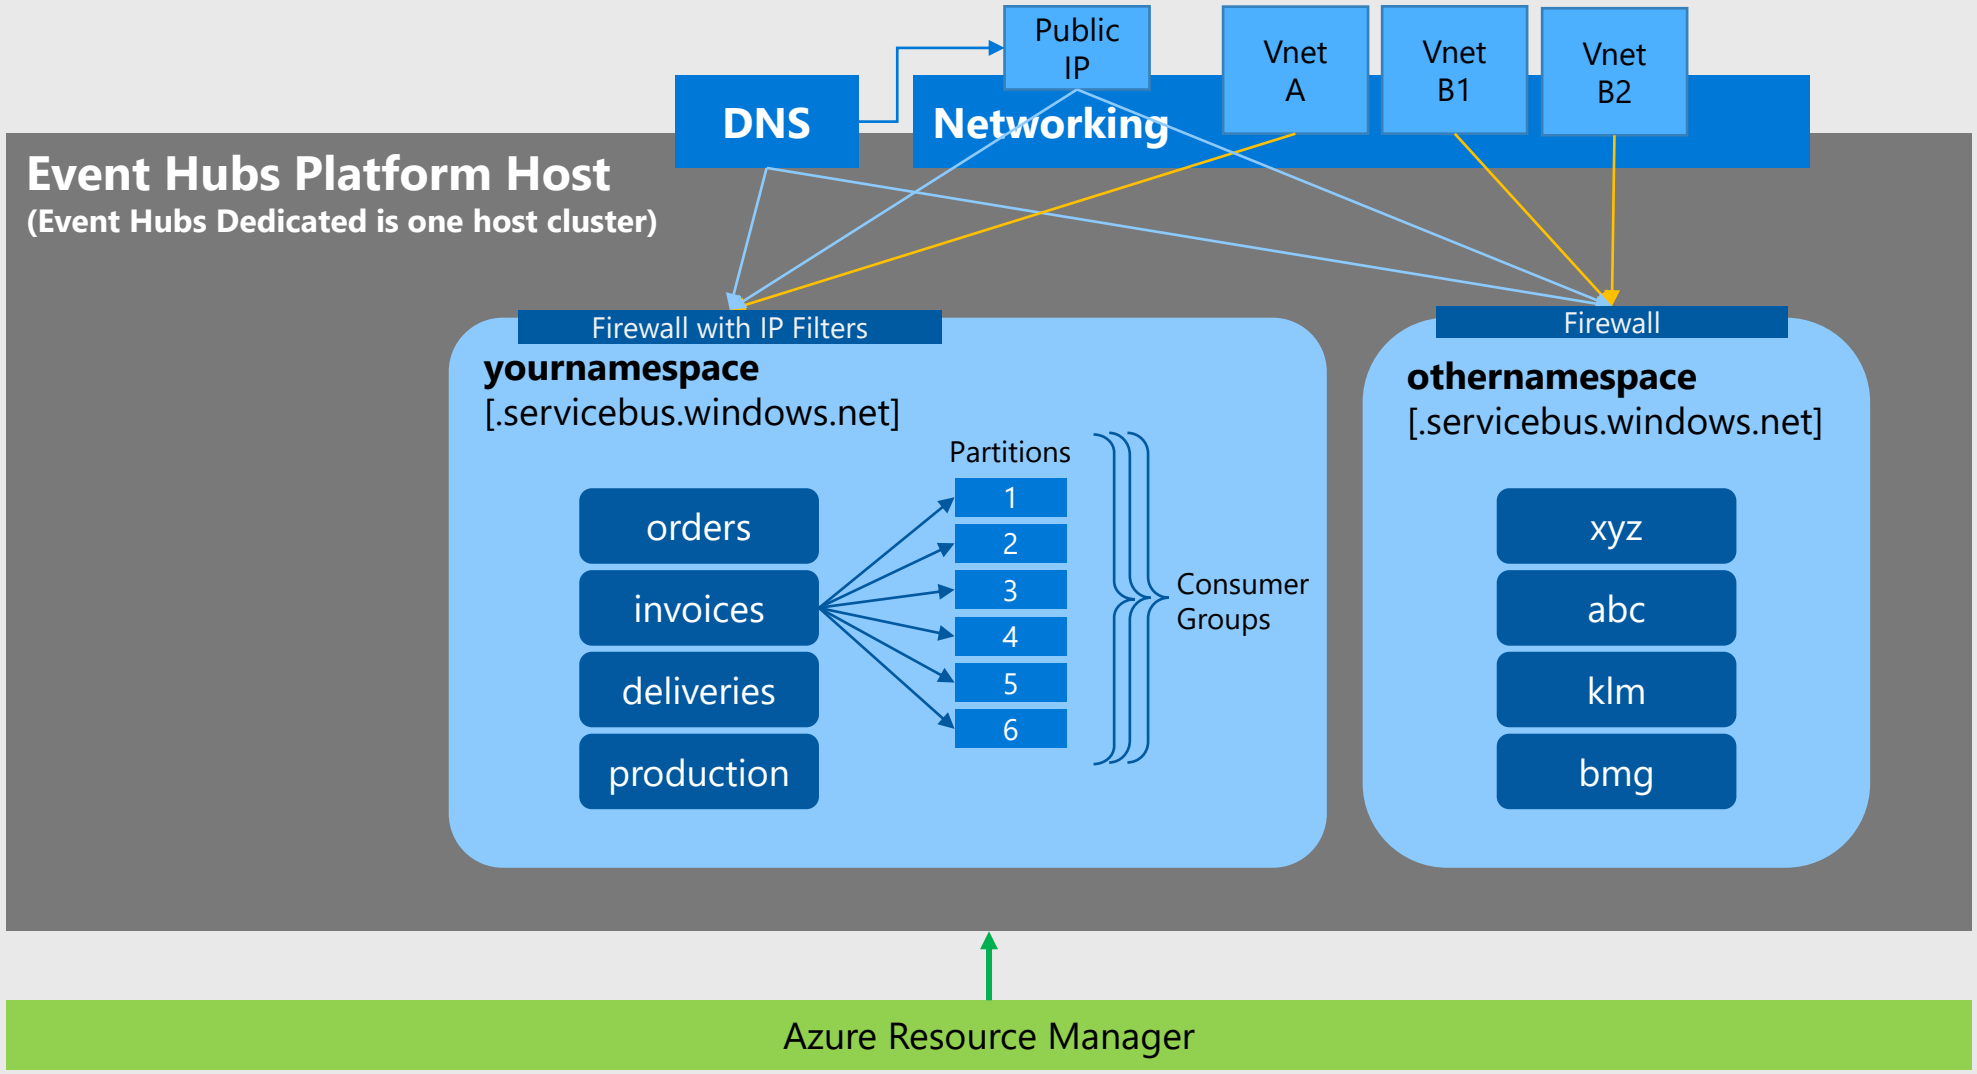

Event Hubs

A log partition is a durable, triple-replicated, append-only storage structure. Each partition has a throughput ceiling imposed by IO and compute limits. More partitions provide more throughput capacity, but also make consumption more complex.

A consumer group is an optional mechanism to allow multiple parallel consumers to coordinate who owns which partition.

Kafka RPC: Broker assigns partitions and checkpoints can be stored in the CG.

AMQP: Clients choose leader using an external mechanism and can break leases using an epoch model. Checkpoints are external. AMQP with server-assigned partitions and CG checkpoints on roadmap.

Producers send using a protocol of their choosing, with or without setting a partition key. Partitions can also be targeted directly.

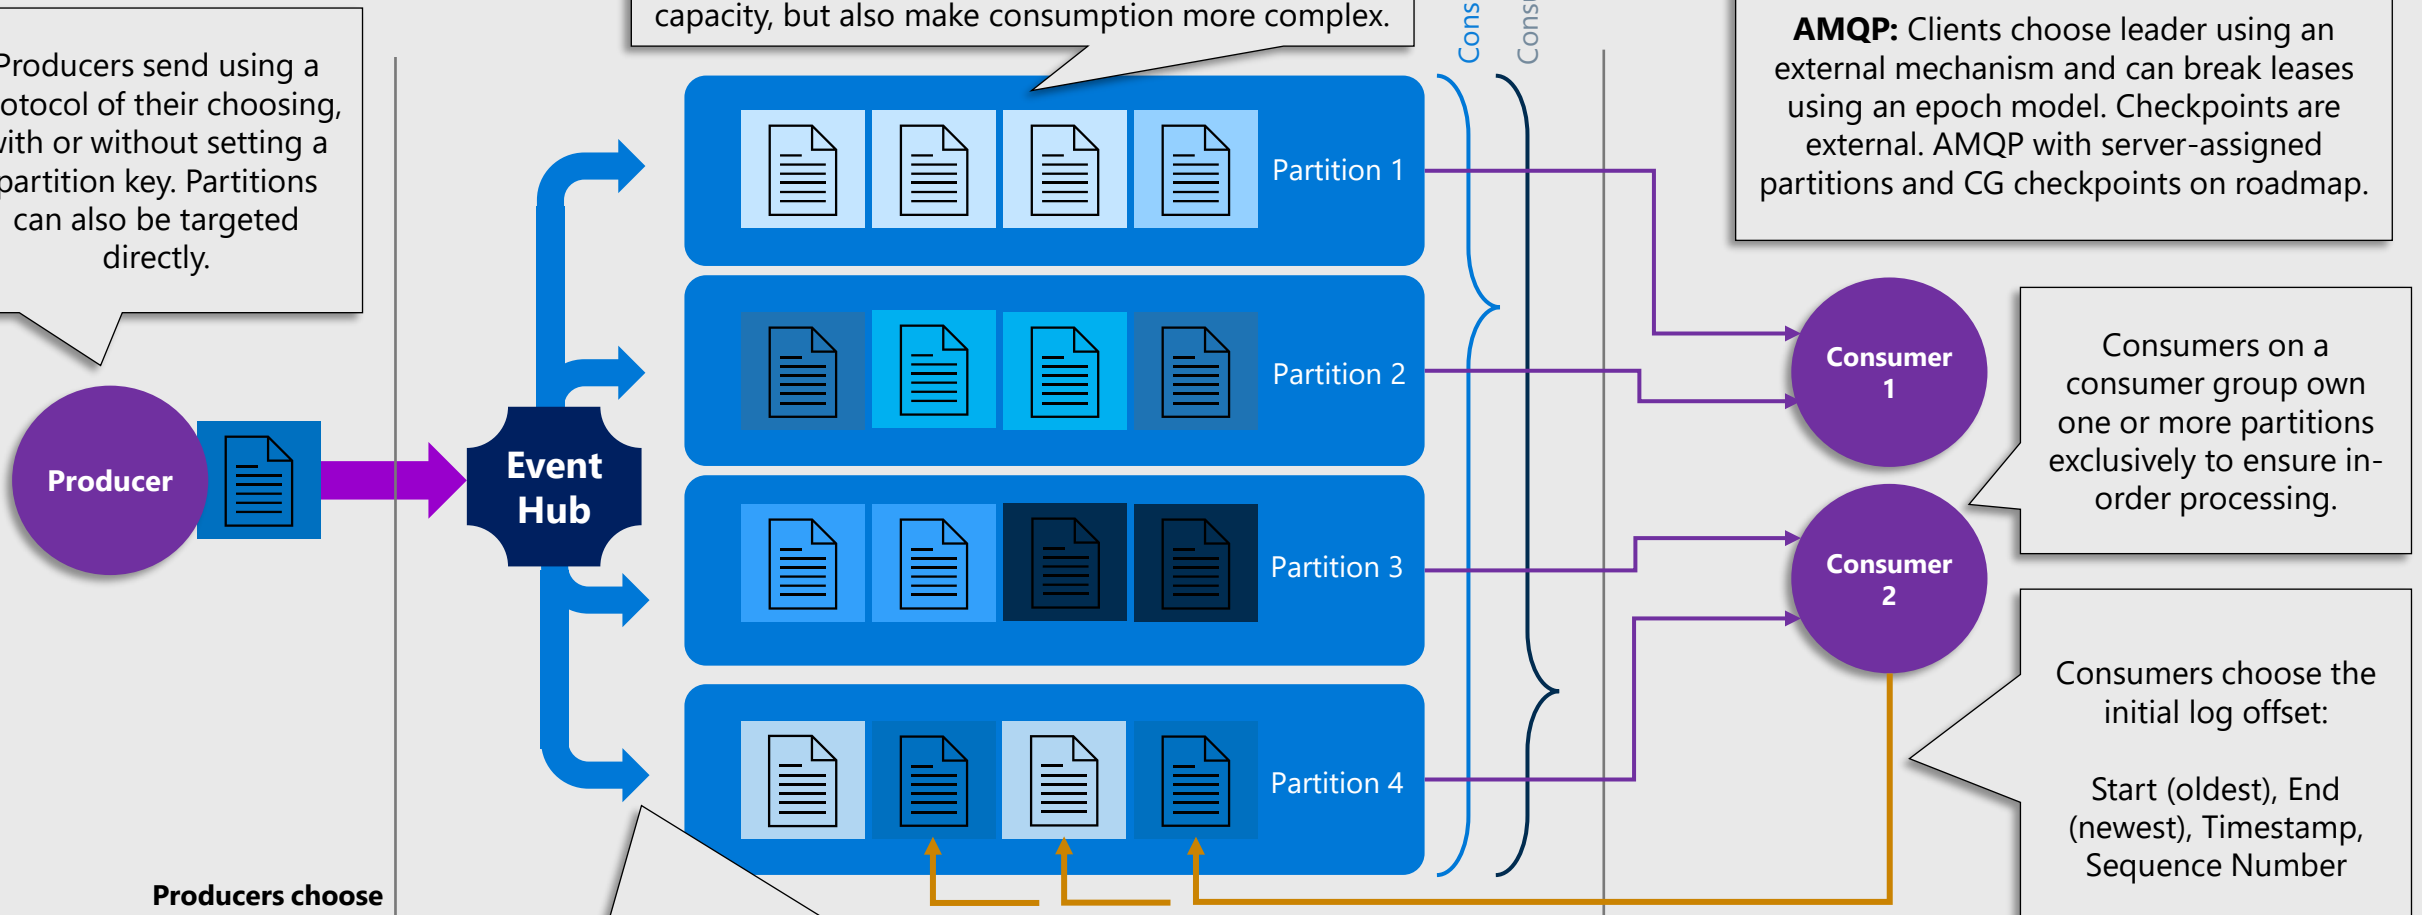

Consumers on a consumer group own one or more partitions exclusively to ensure in-order processing.

Consumers choose the initial log offset:

Start (oldest), End (newest), Timestamp, Sequence Number

Read progress is noted in checkpoints that can be stored anywhere.

Producers choose their preferred protocol

OASIS AMQP 1.0
Apache Kafka® RPC
HTTPS 1.1

An Event Hub is an event stream manager handling one or more independent log partitions. Different from Apache Kafka®, Event Hubs can optionally assign events to partitions on the server-side.

Consumers choose their own preferred protocol

OASIS AMQP 1.0
Apache Kafka® RPC

Event Hubs Capture

{Namespace}/**{EventHub}**/**{PartitionId}**/{Year}/{Month}/{Day}/{Hour}/{Minute}/{Second}

https://mystorageaccount.blob.core.windows.net/mycontainer/mynamespace/myeventhub/0/2017/12/08/03/03/17.avro

Event Hubs Capture allows organizing torrents of telemetry in batches that can be easily consumed by batch-oriented analytics frameworks.

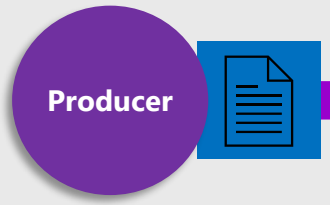

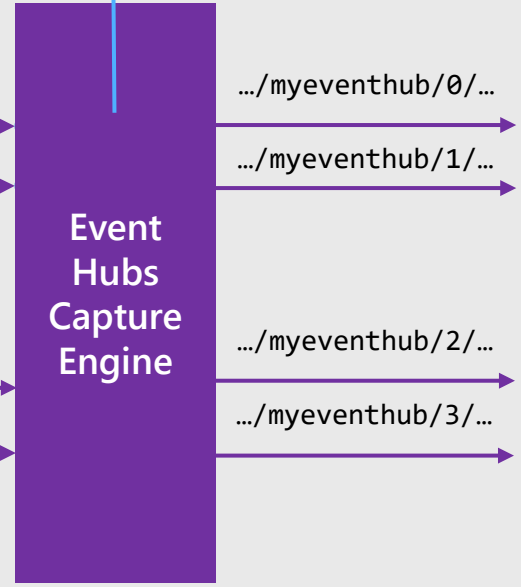

Event is raised when capture file has been written

.../myeventhub/0/...

.../myeventhub/1/...

.../myeventhub/2/...

.../myeventhub/3/...

Azure Blob Storage

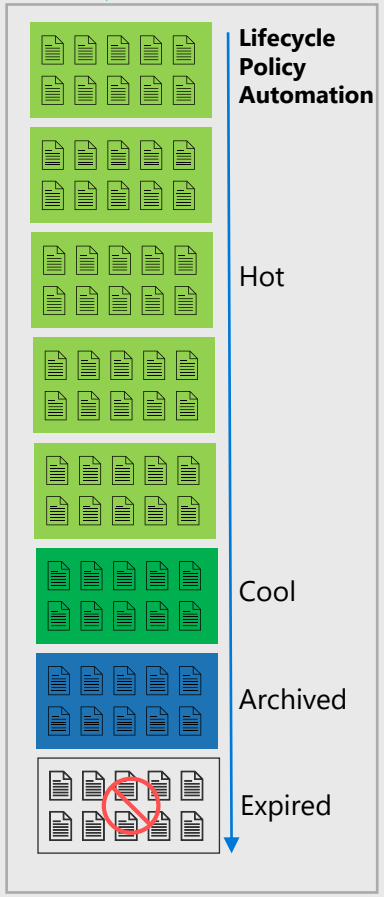

Azure Data Lake Gen2

Producers choose their preferred protocol

OASIS AMQP 1.0
Apache Kafka® RPC
HTTPS 1.1

Capture throughput keeps up with the Event Hubs' ingress throughput capacity and provides automated low latency capture to batch storage.

Capture policy allows balancing latency and batch size.

Event Hubs Capture Avro Archives

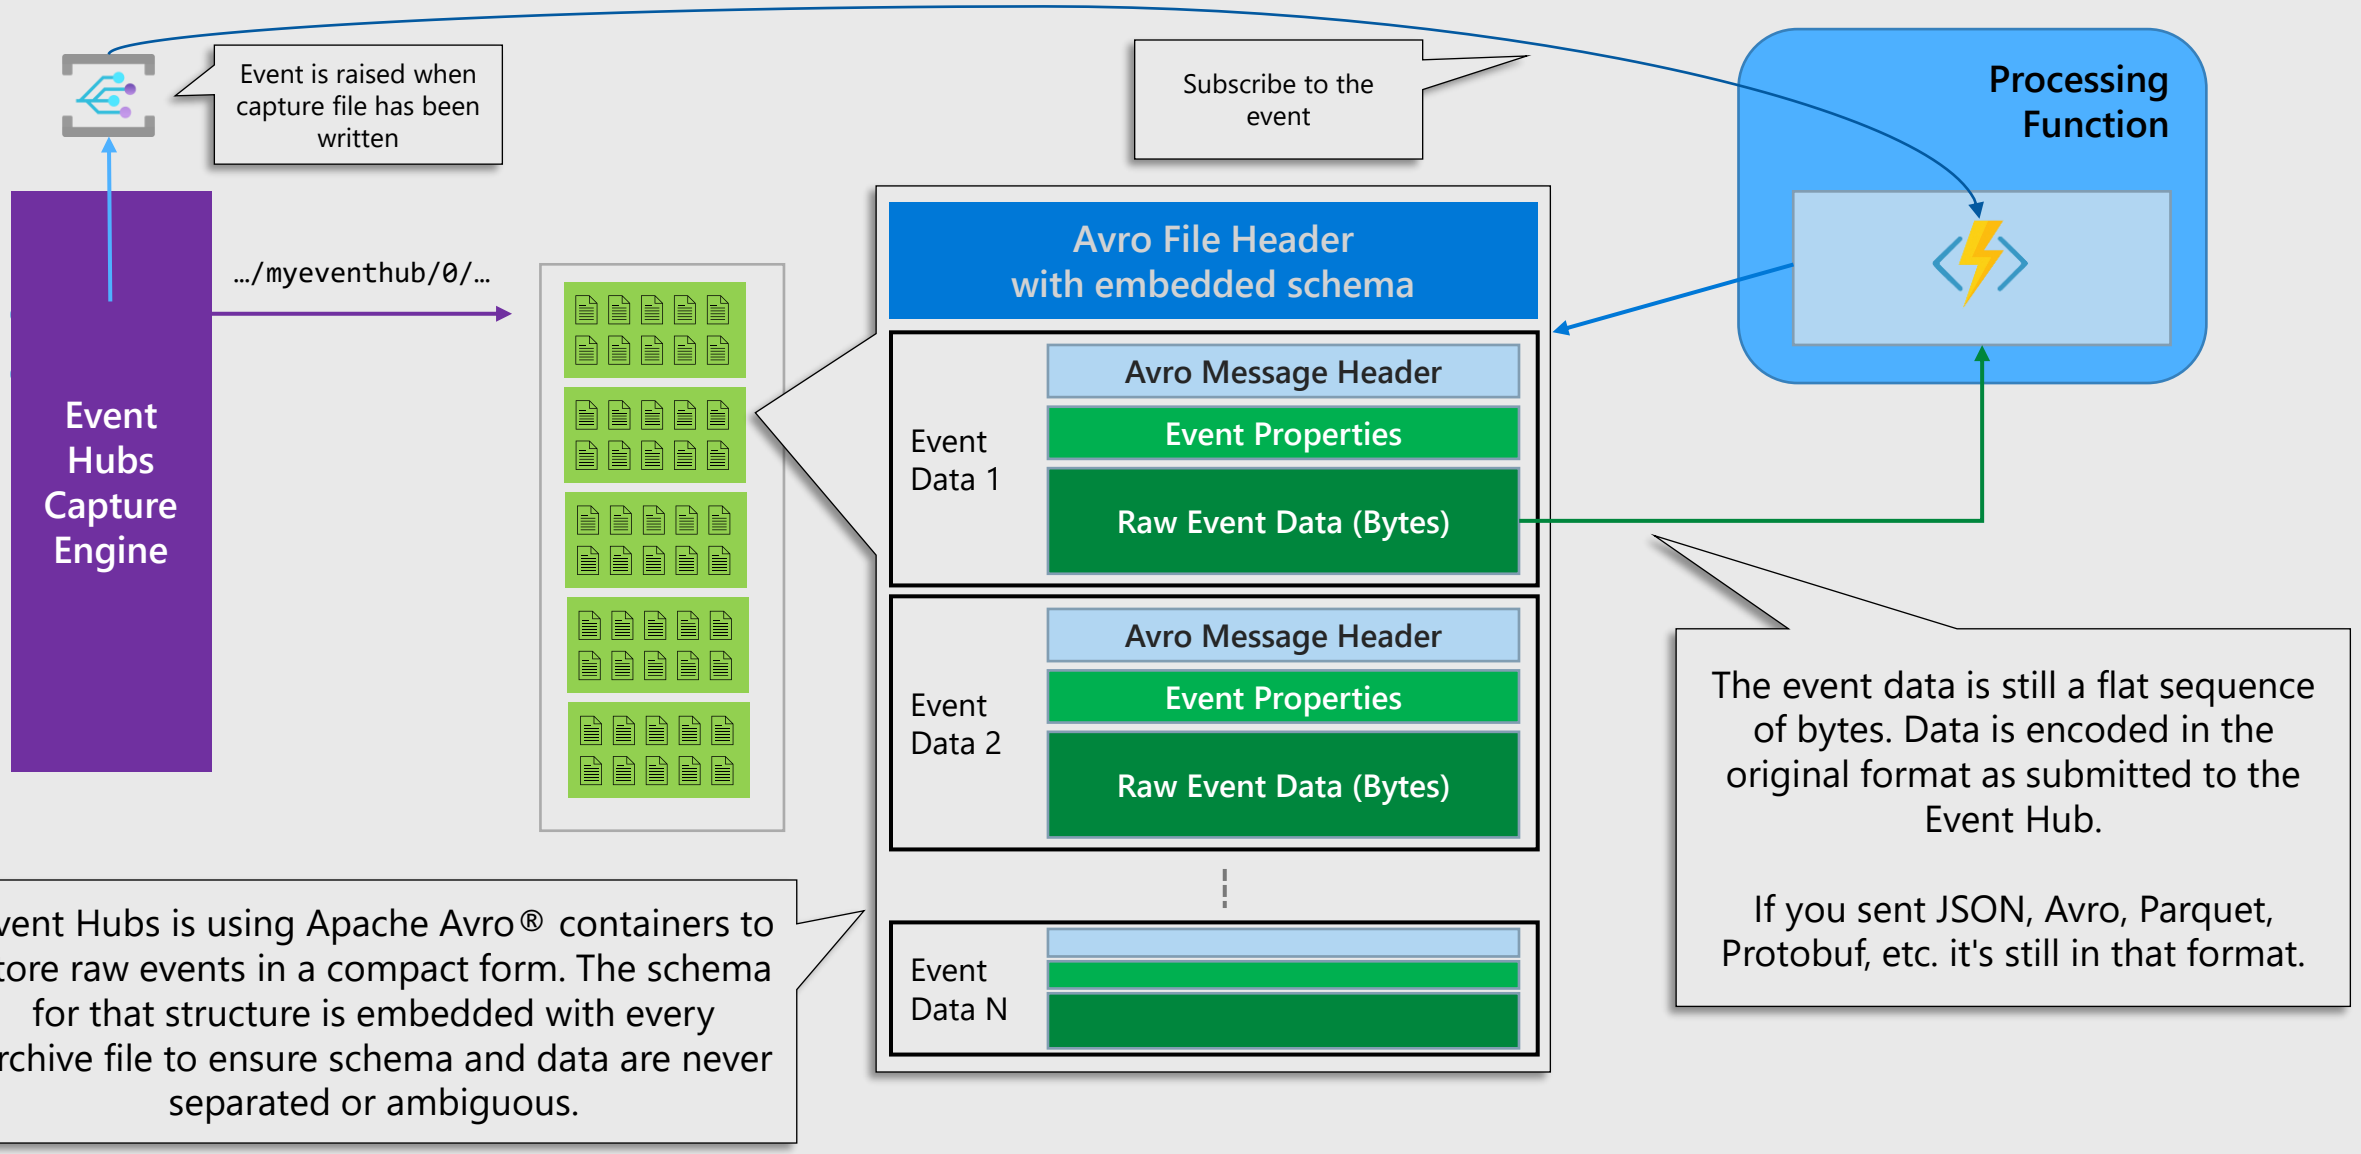

* Apache Kafka isn't any of these things, either.

Event Hubs* is not…

… a publish/subscribe broker. Partitions are not subscriptions. They are chosen by the producer or the broker on ingress. There is also no server-side filtering.

→ Azure Service Bus, Azure Event Grid

… a queue broker. Read progress over the log is handled by the client and there is no event-level ownership and delivery state handling.

→ Azure Service Bus

… a discrete event distribution engine. Event Hubs does not do push deliveries, and delivery failures need to be tracked individually.

→ Azure Event Grid

… a database or long-term event store. Event Hubs exists to catch, store and provide fast access to event data organized around time axis. As data ages (days, not months), you need better indexing.

→ Azure Cosmos DB, Azure SQL, Azure Synapse, Azure Data Explorer

Apache Kafka® Clients

https://cwiki.apache.org/confluence/pages/viewpage.action?pageId=30736606

Azure Event Hubs is compatible with the producer and consumer APIs of all current Apache Kafka clients.

Only the Java/JVM client is part of the Apache Kafka project. All other clients are under various non-foundation ownerships and licenses. Azure does not provide QFE support for any of these clients but does provide support for service-side compatibility.

Apache Software Foundation

Non-Foundation, Community / Best Effort Support

Non-Foundation, Commercial Support Option

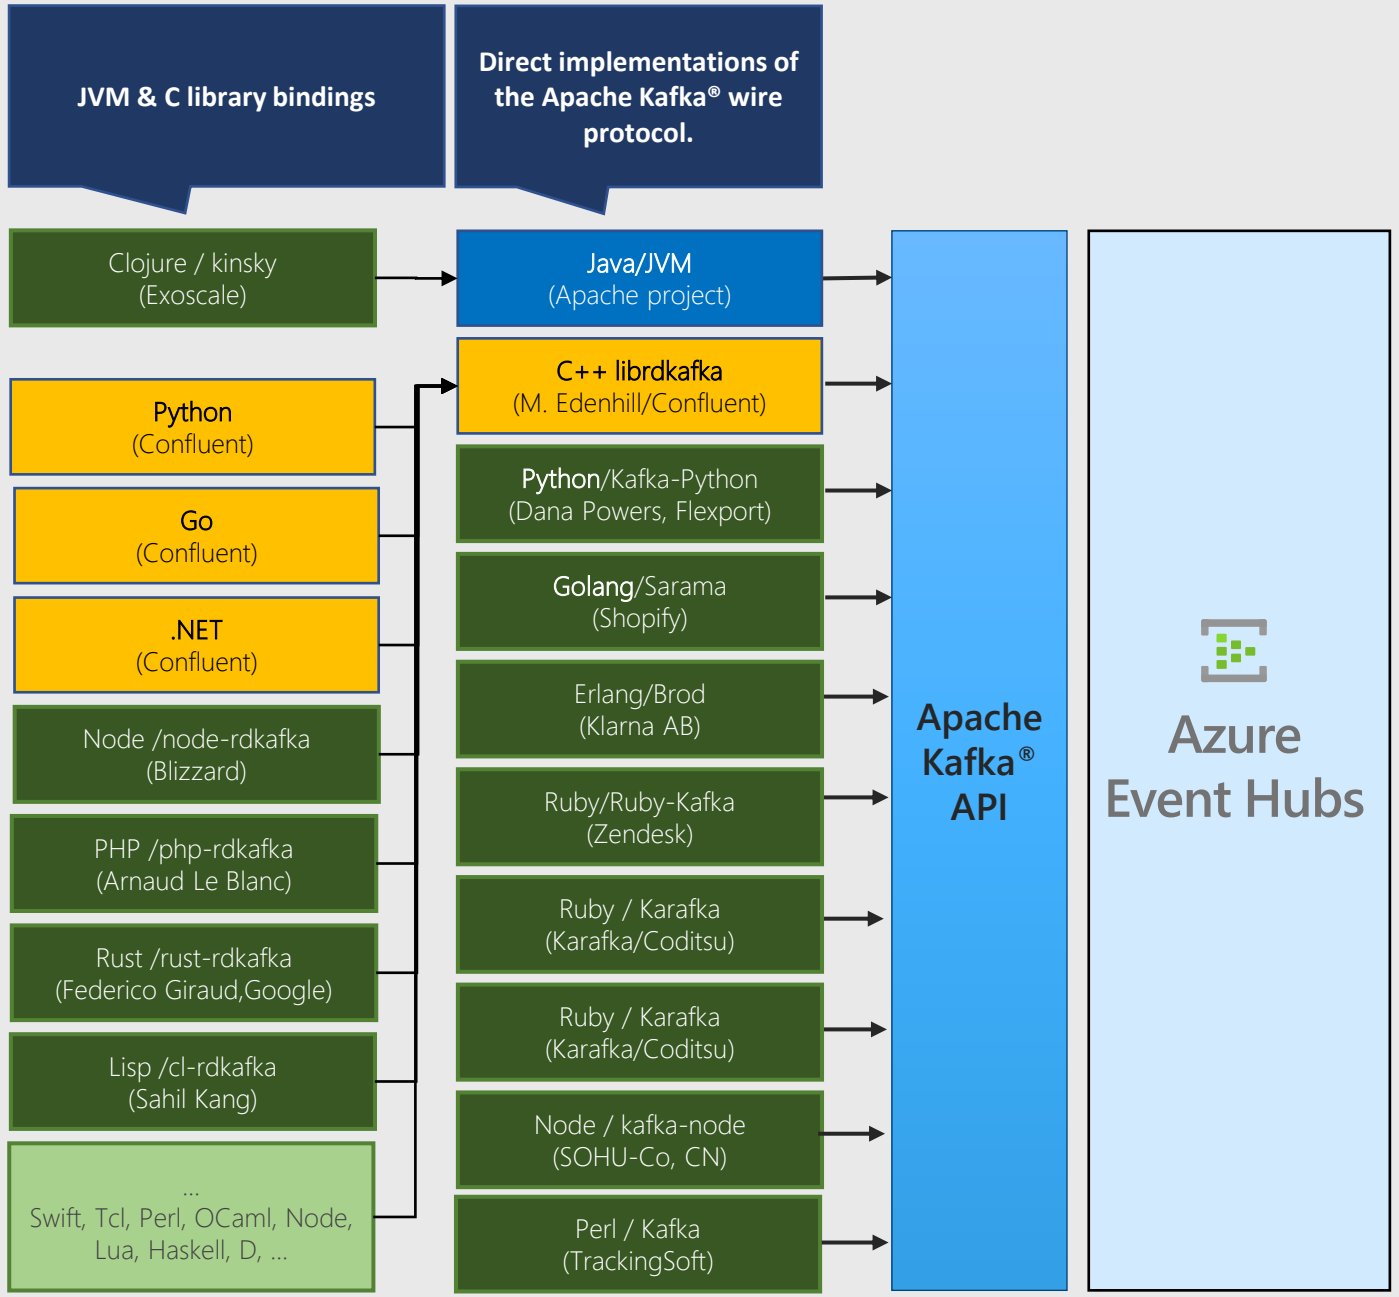

Event Hubs SDK Support: Azure SDKs

.NET

Java

Java Spring Boot

Python

JavaScript/Typescript

Go

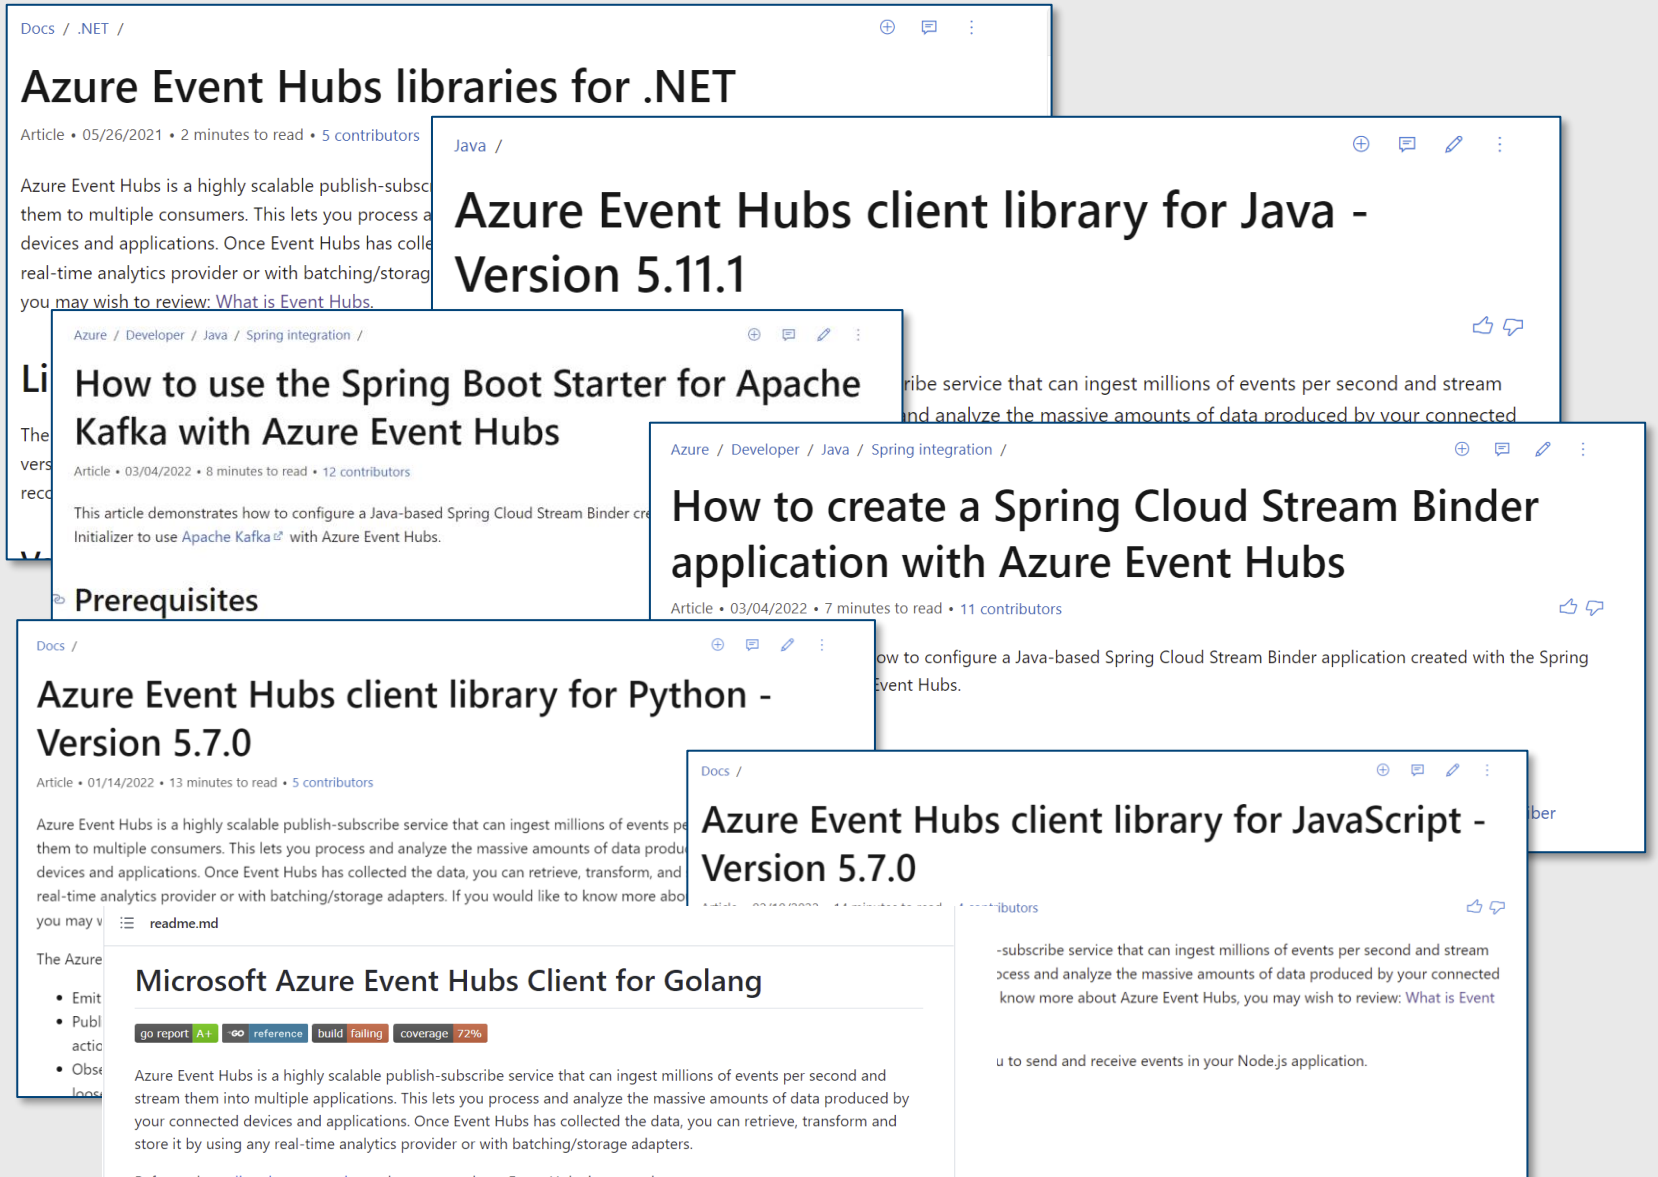

Compatible Generic AMQP 1.0 Clients

Client Stack	URL
Apache Qpid Proton-C (Go, Python, Ruby, C++)	https://qpid.apache.org/proton/
Apache Qpid Messaging API	https://qpid.apache.org/components/messaging-api
Apache Qpid Proton-J	https://qpid.apache.org/proton/
Apache NMS AMQP	https://activemq.apache.org/components/nms/providers/amqp/
AMQP .NET Lite (all variants of .NET, incl. Nano & Micro)	https://github.com/Azure/amqpnetlite
Azure AMQP (our own server stack)	https://github.com/Azure/azure-amqp
Azure uAMQP C (Python, PHP)	https://github.com/Azure/azure-uamqp-c https://github.com/Azure/azure-uamqp-python
Rhea (NodeJS)	https://github.com/amqp/rhea
Go AMQP	https://github.com/Azure/go-amqp
Vert.X AMQP Client	https://vertx.io/docs/vertx-amqp-client/java/

Azure Event Hubs Deep Dive

Architecture

Internal Architecture

Protocols

Performance Metrics

Logical Architecture

3-tier architecture: Networking, Messaging, Storage

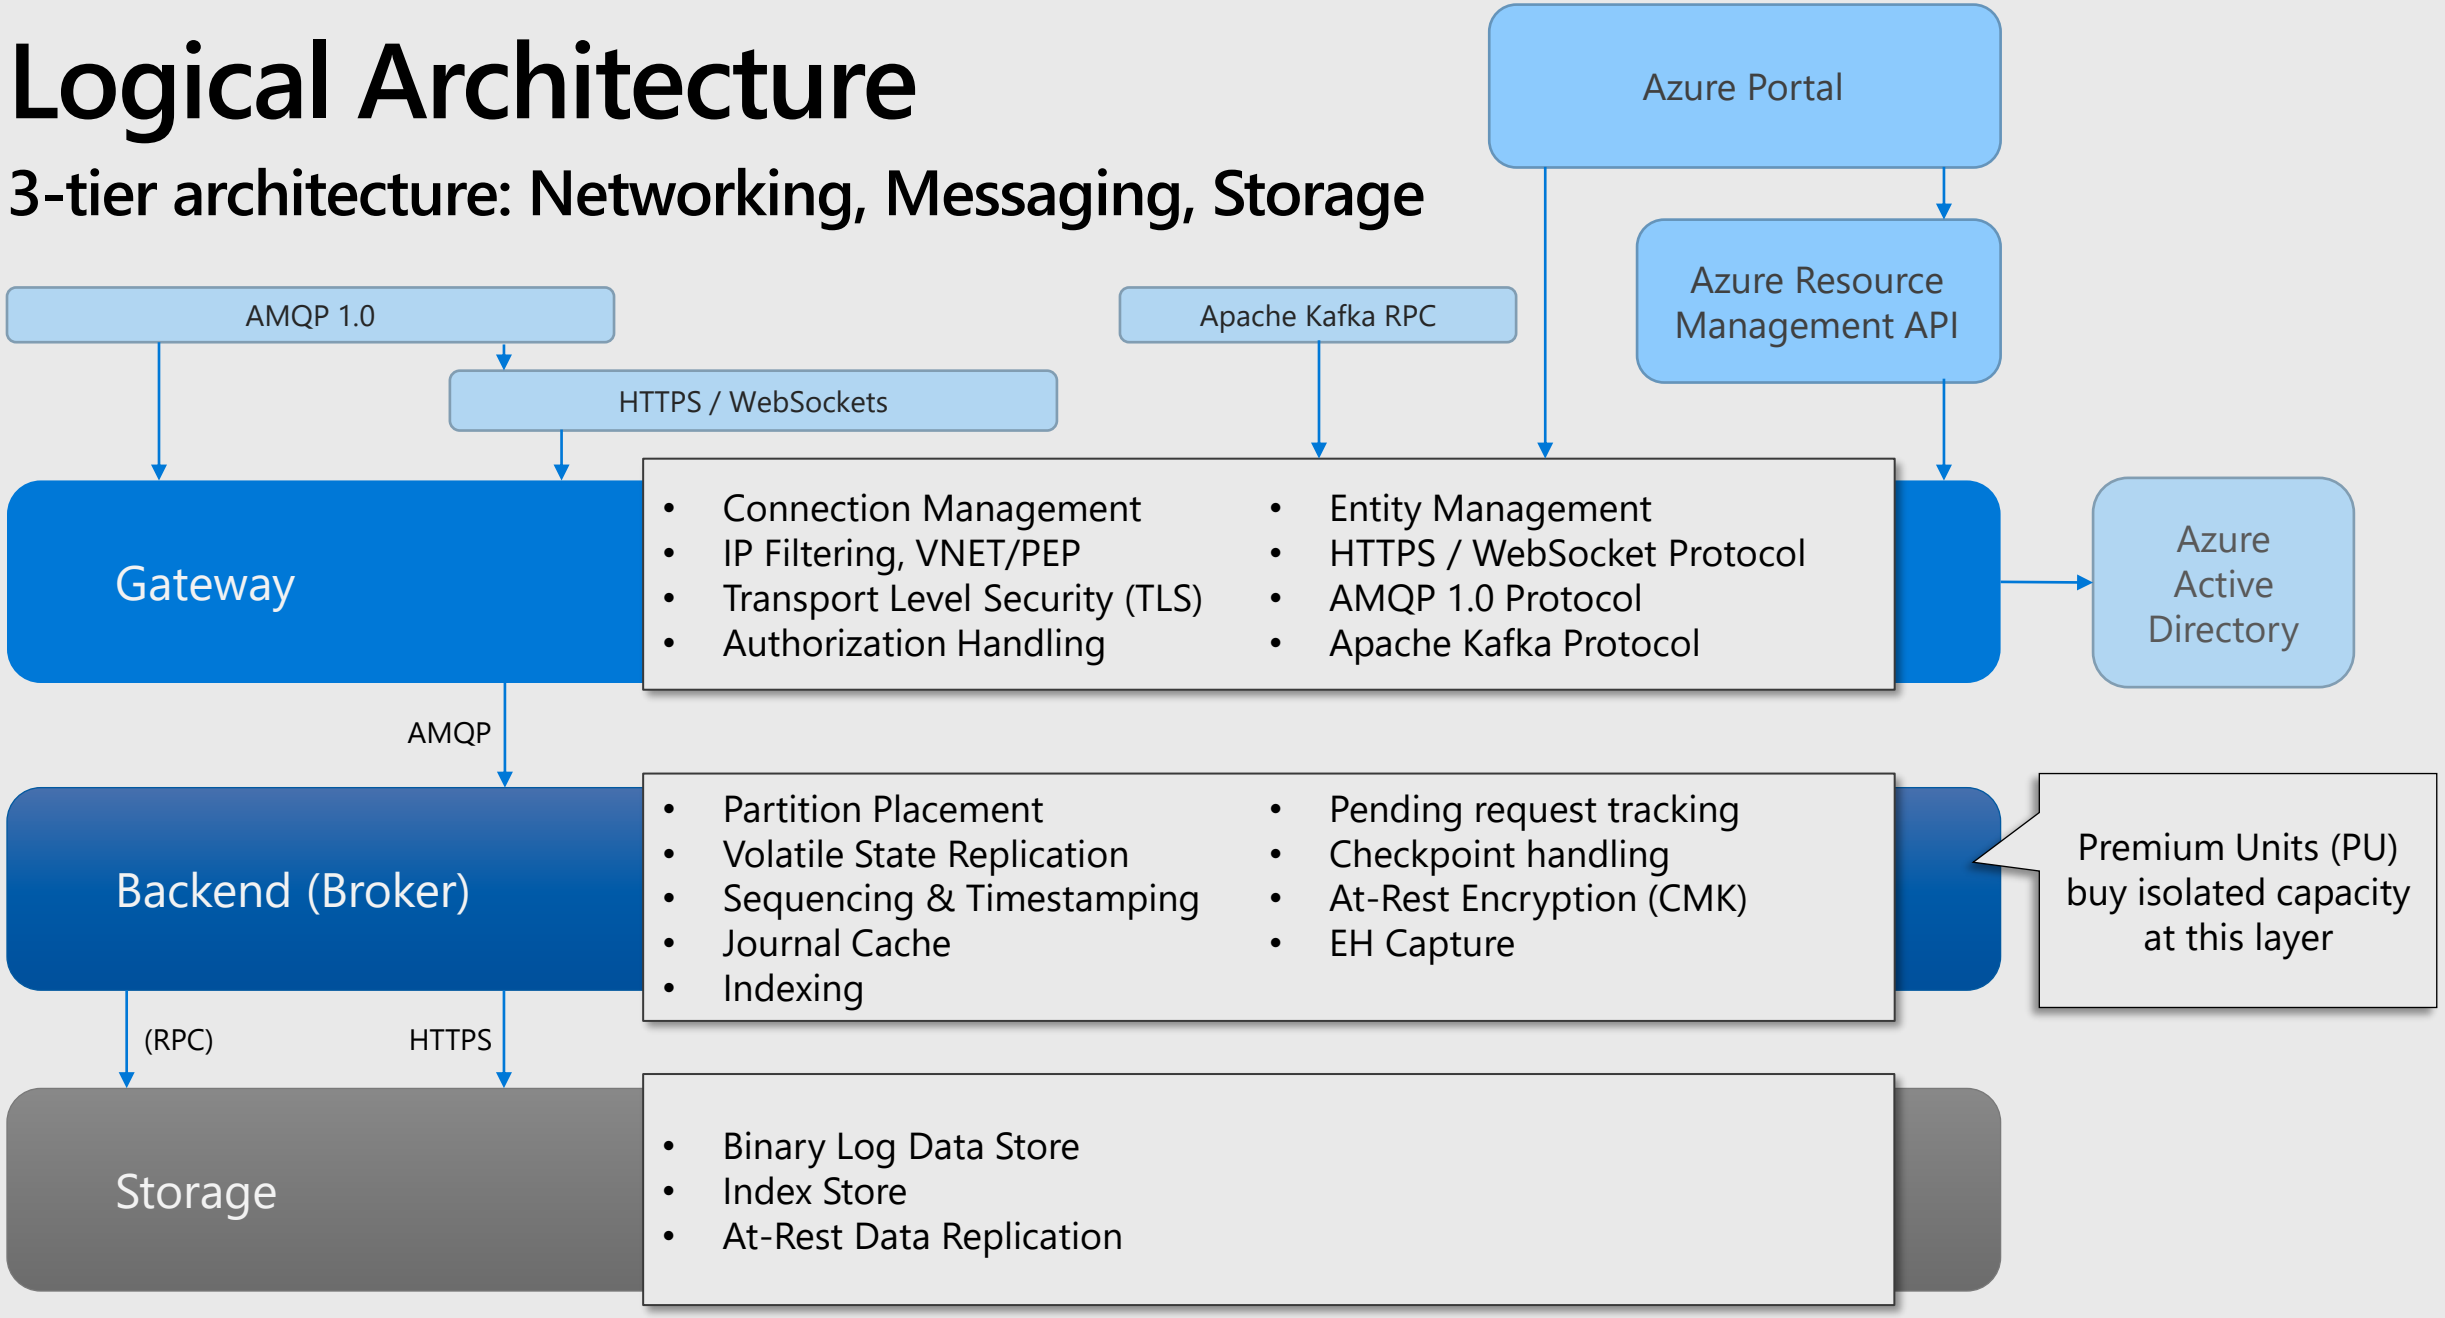

Fault Domains & Availability Zones

High-Availability starts at the physical reality

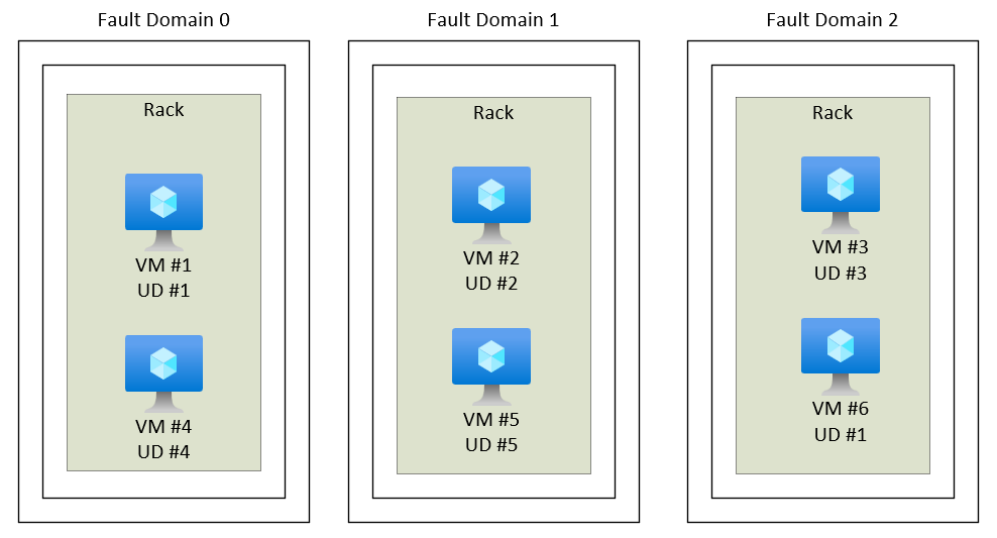

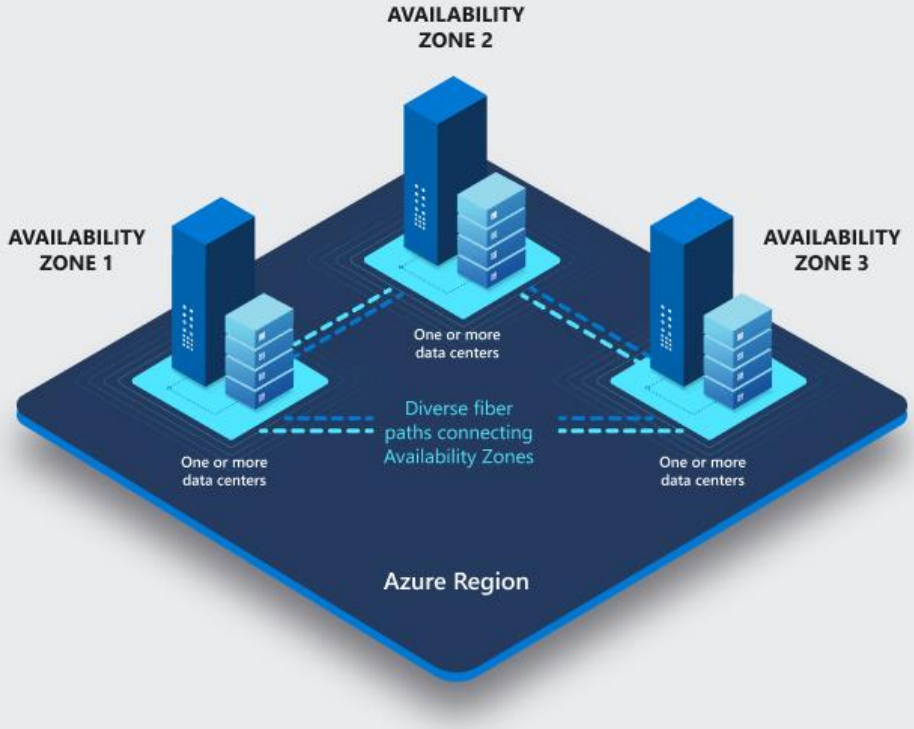

Fault Domain Placement:
Cluster VMs are spread across at least 3 fault domains such that the loss of a rack or network poses no availability risk. Recovery from a fault domain failure is fully automated and the system maintains SLA.

Availability Zones Placement:
Each cluster spans three availability zones and maintains SLA without any tolerance for data loss when one or two zones fail.

Backend and Gateway Clusters

Logical Architecture meets Placement

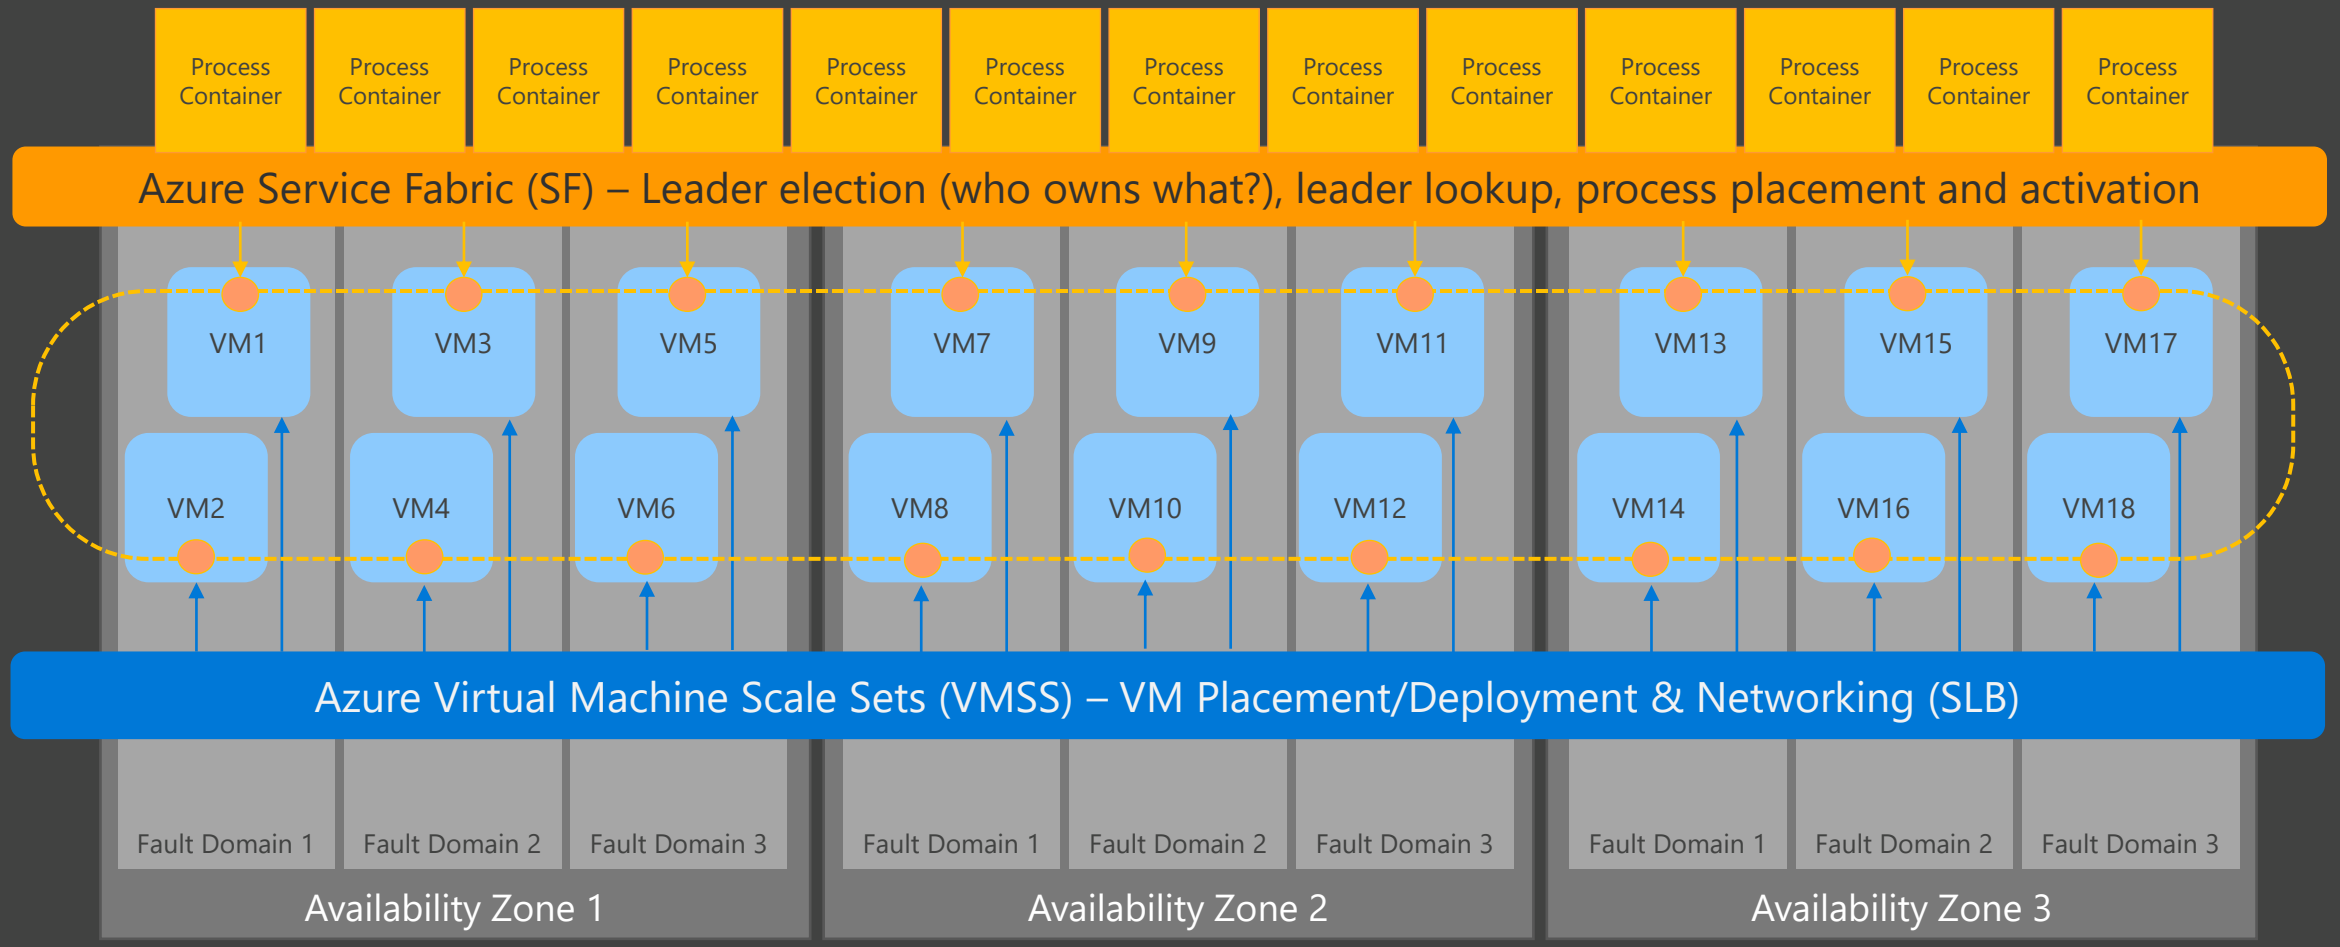

Backend and Gateway Clusters

Principles

Everything is automated. There's no manual intervention in placement of VMs or processes or workloads.

Process placement and servicing/upgrade cycles are controlled by Azure Service Fabric.

Each cluster has a well-known number of virtual machines. We don't use auto-scaling with VMs. We may (rarely) choose to initiate scale up/down to a different level.

Allocation of namespaces to clusters are based on well-established heuristics (Standard) or deterministic capacity limits (Premium). Customer-defined allocation in Dedicated.

Gateway and Backend

Workload placement

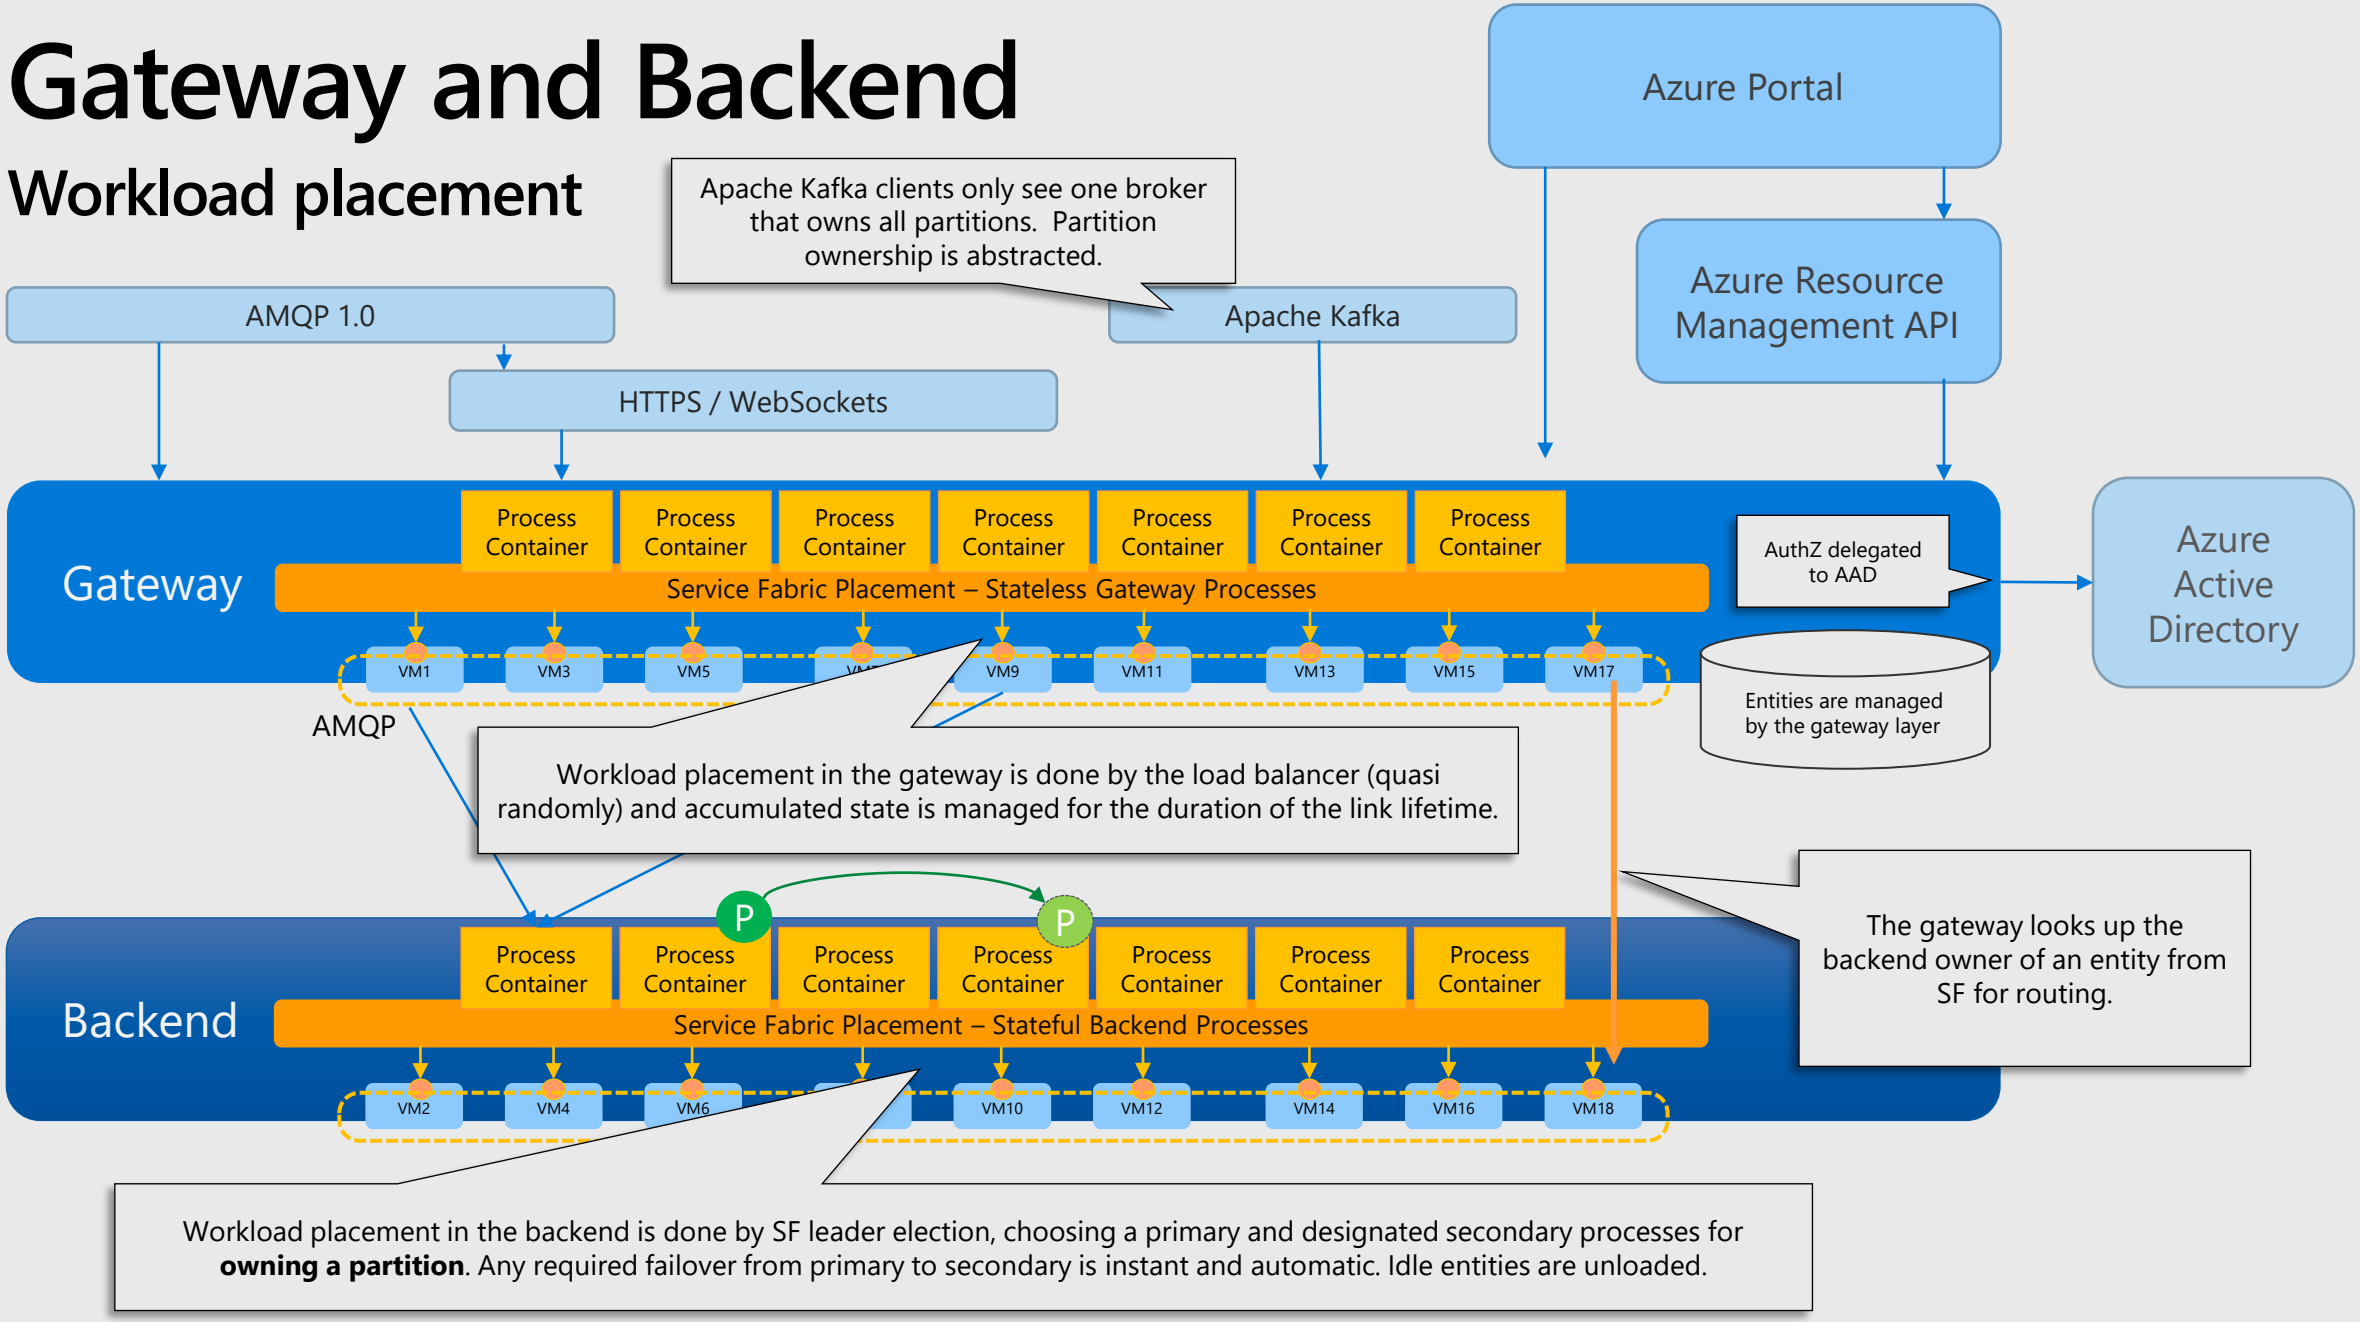

Backend: Event Hubs Premium.

Workload placement

Namespace PU are split across processes:
1 PU = 2 Proc (8GB Mem), 1 Core/Proc (2C)
2 PU = 2 Proc, 1 Core/Proc + 1 Core (3C)
4 PU = 2 Proc, 2 Core/Proc + 1 Core (5C)
8 PU = 4 Proc, 2 Core/Proc + 1 Core (9C)
16 PU = 4 Proc, 4 Core/Proc + 1 Core (17C)

Partition ownership is dynamically mapped to the process(es) associated with a namespace via SF

Host 2

Namespace Broker Process 2 (**4PU**)

P P P P P P P P

Host 1

Namespace Broker Process 1 (**4PU**)

P P P P P P P P

Cores are exclusively mapped to a broker process.

>=2PU: 1 Core extra for utility tasks.

Isolated VMs / L88is

CPU Cores L-Series

1 2 +

Memory Allocations

8GB

VM6

Backend

Process Container | Process Container | Process Container | Process Container | Process Container | Process Container | Process Container

Service Fabric Placement – Stateful Backend Processes

VM2 VM4 VM6 VM8 VM10 VM12 VM14 VM16 VM18

Designated secondaries in different fault domains and AZ reserved for failover

Storage Layer – Event Hubs Standard

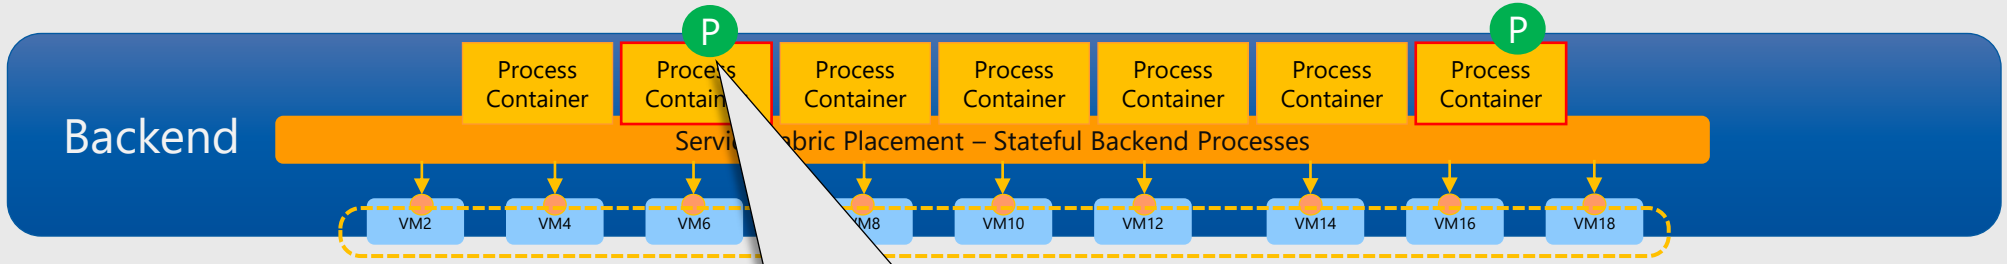

Events are written to append blobs ("extents"). All writes are synchronously triple replicated across availability zones.

Extents are sealed after hitting a size threshold.

Expiry removes extents.

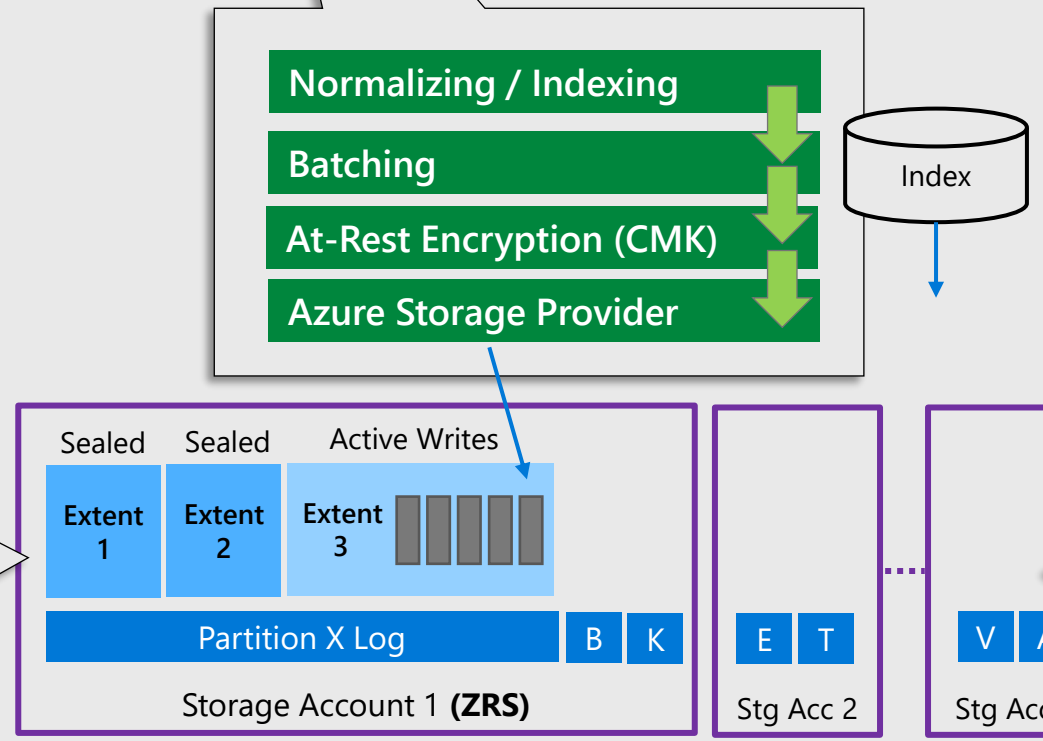

Each Event Hub Cluster uses numerous storage accounts to spread the load and overcome storage quota limitations.

Partition storage is allocated to those storage accounts.

Storage Layer – Event Hubs Premium

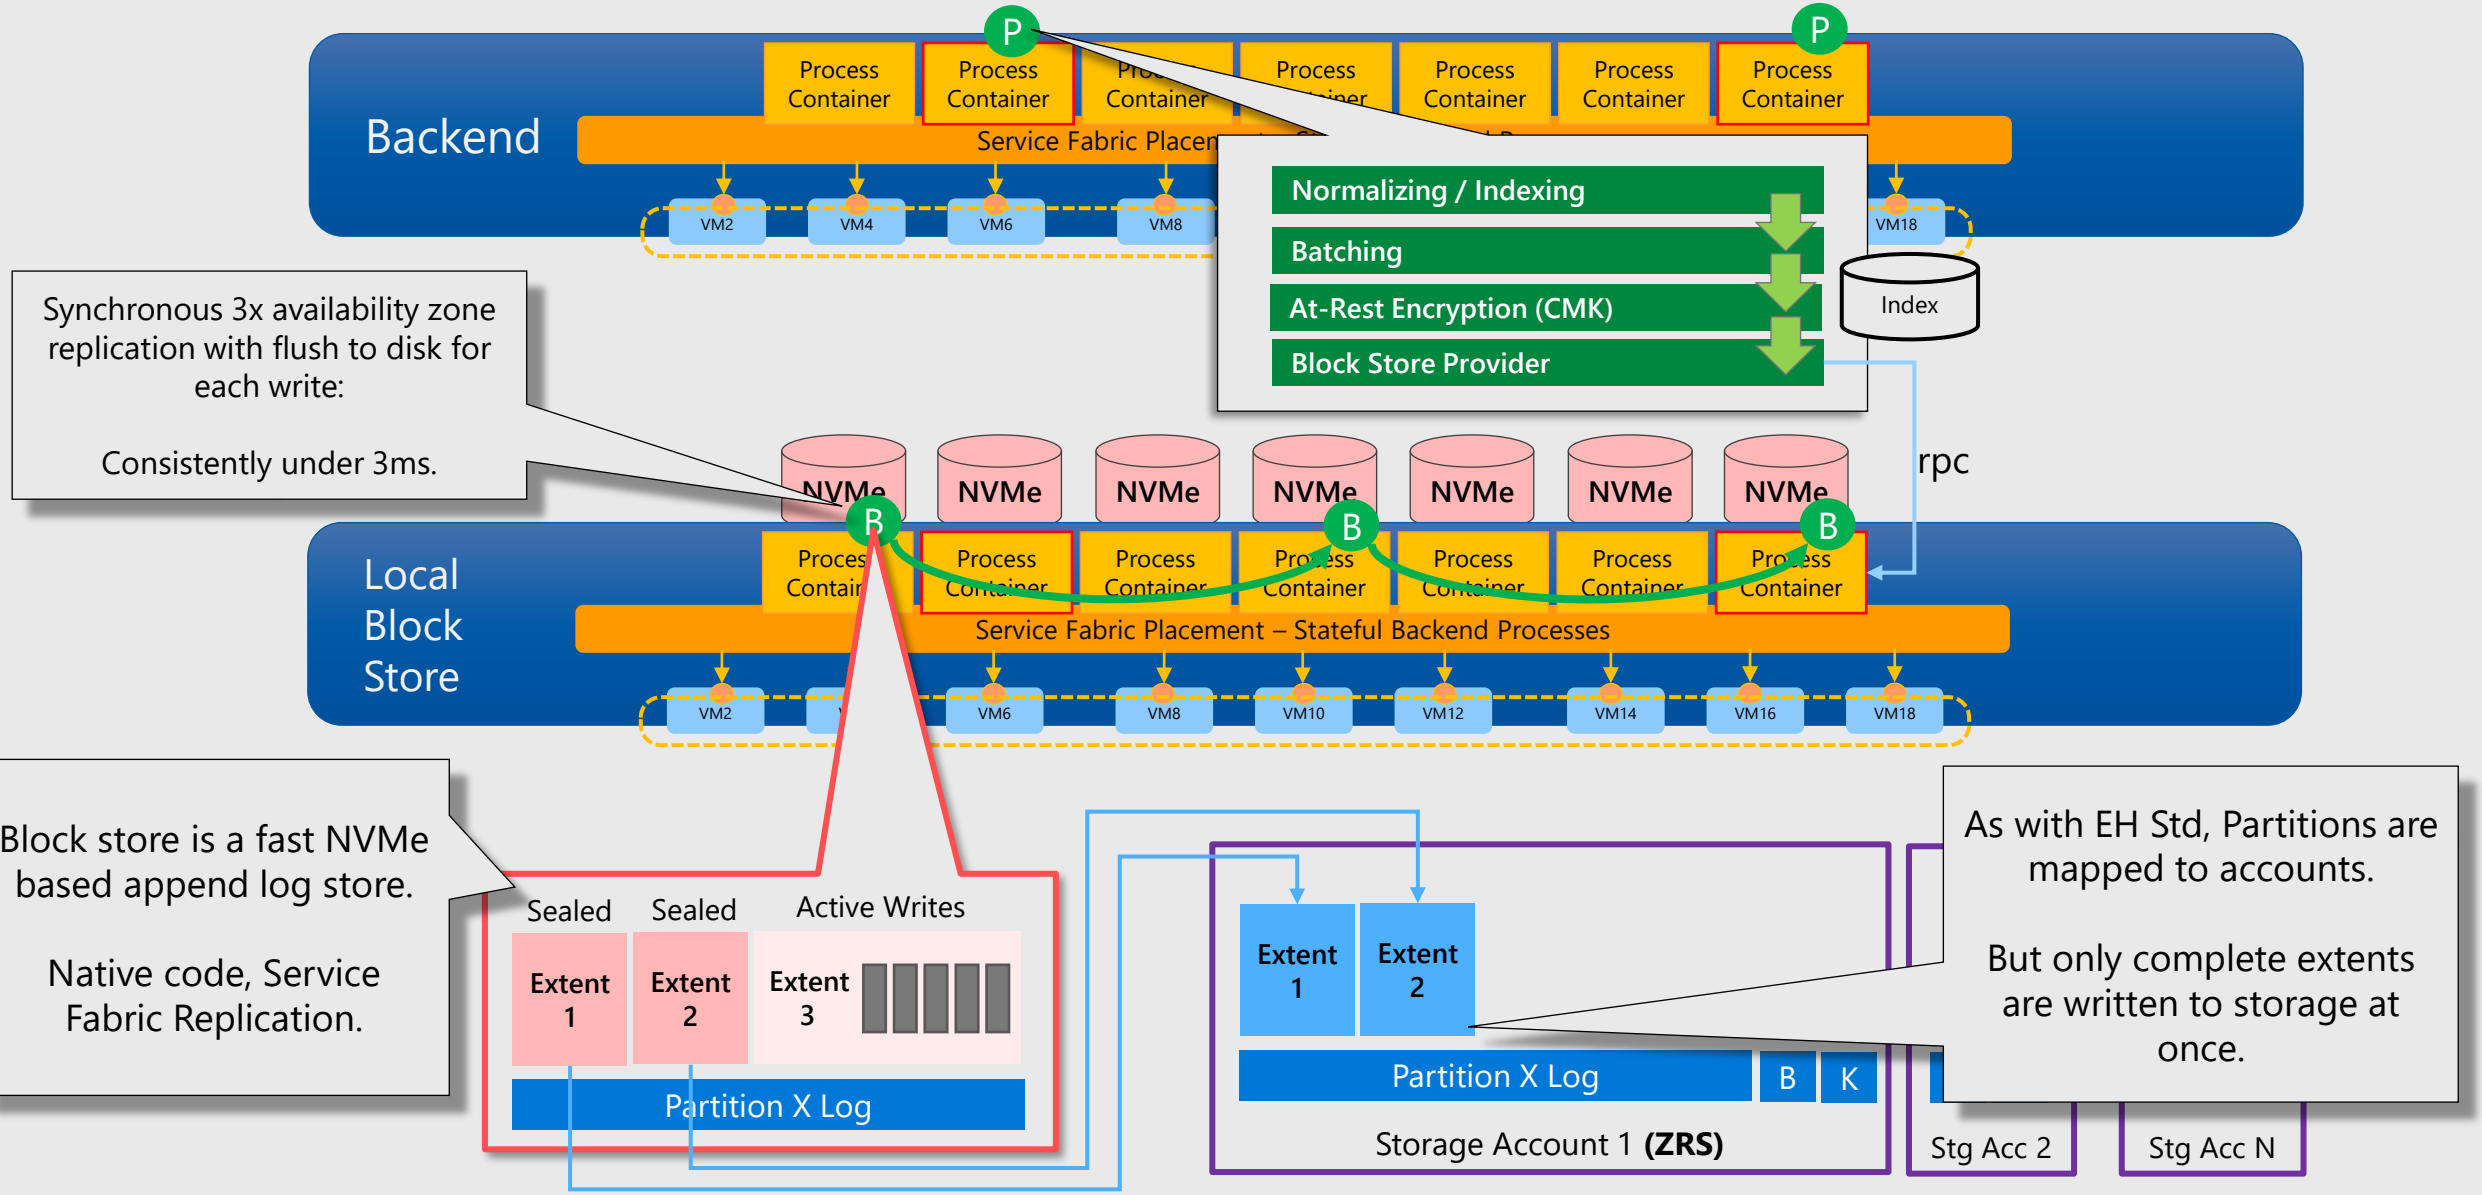

Networking Features

Firewall, Virtual Network Integration with Private Endpoints

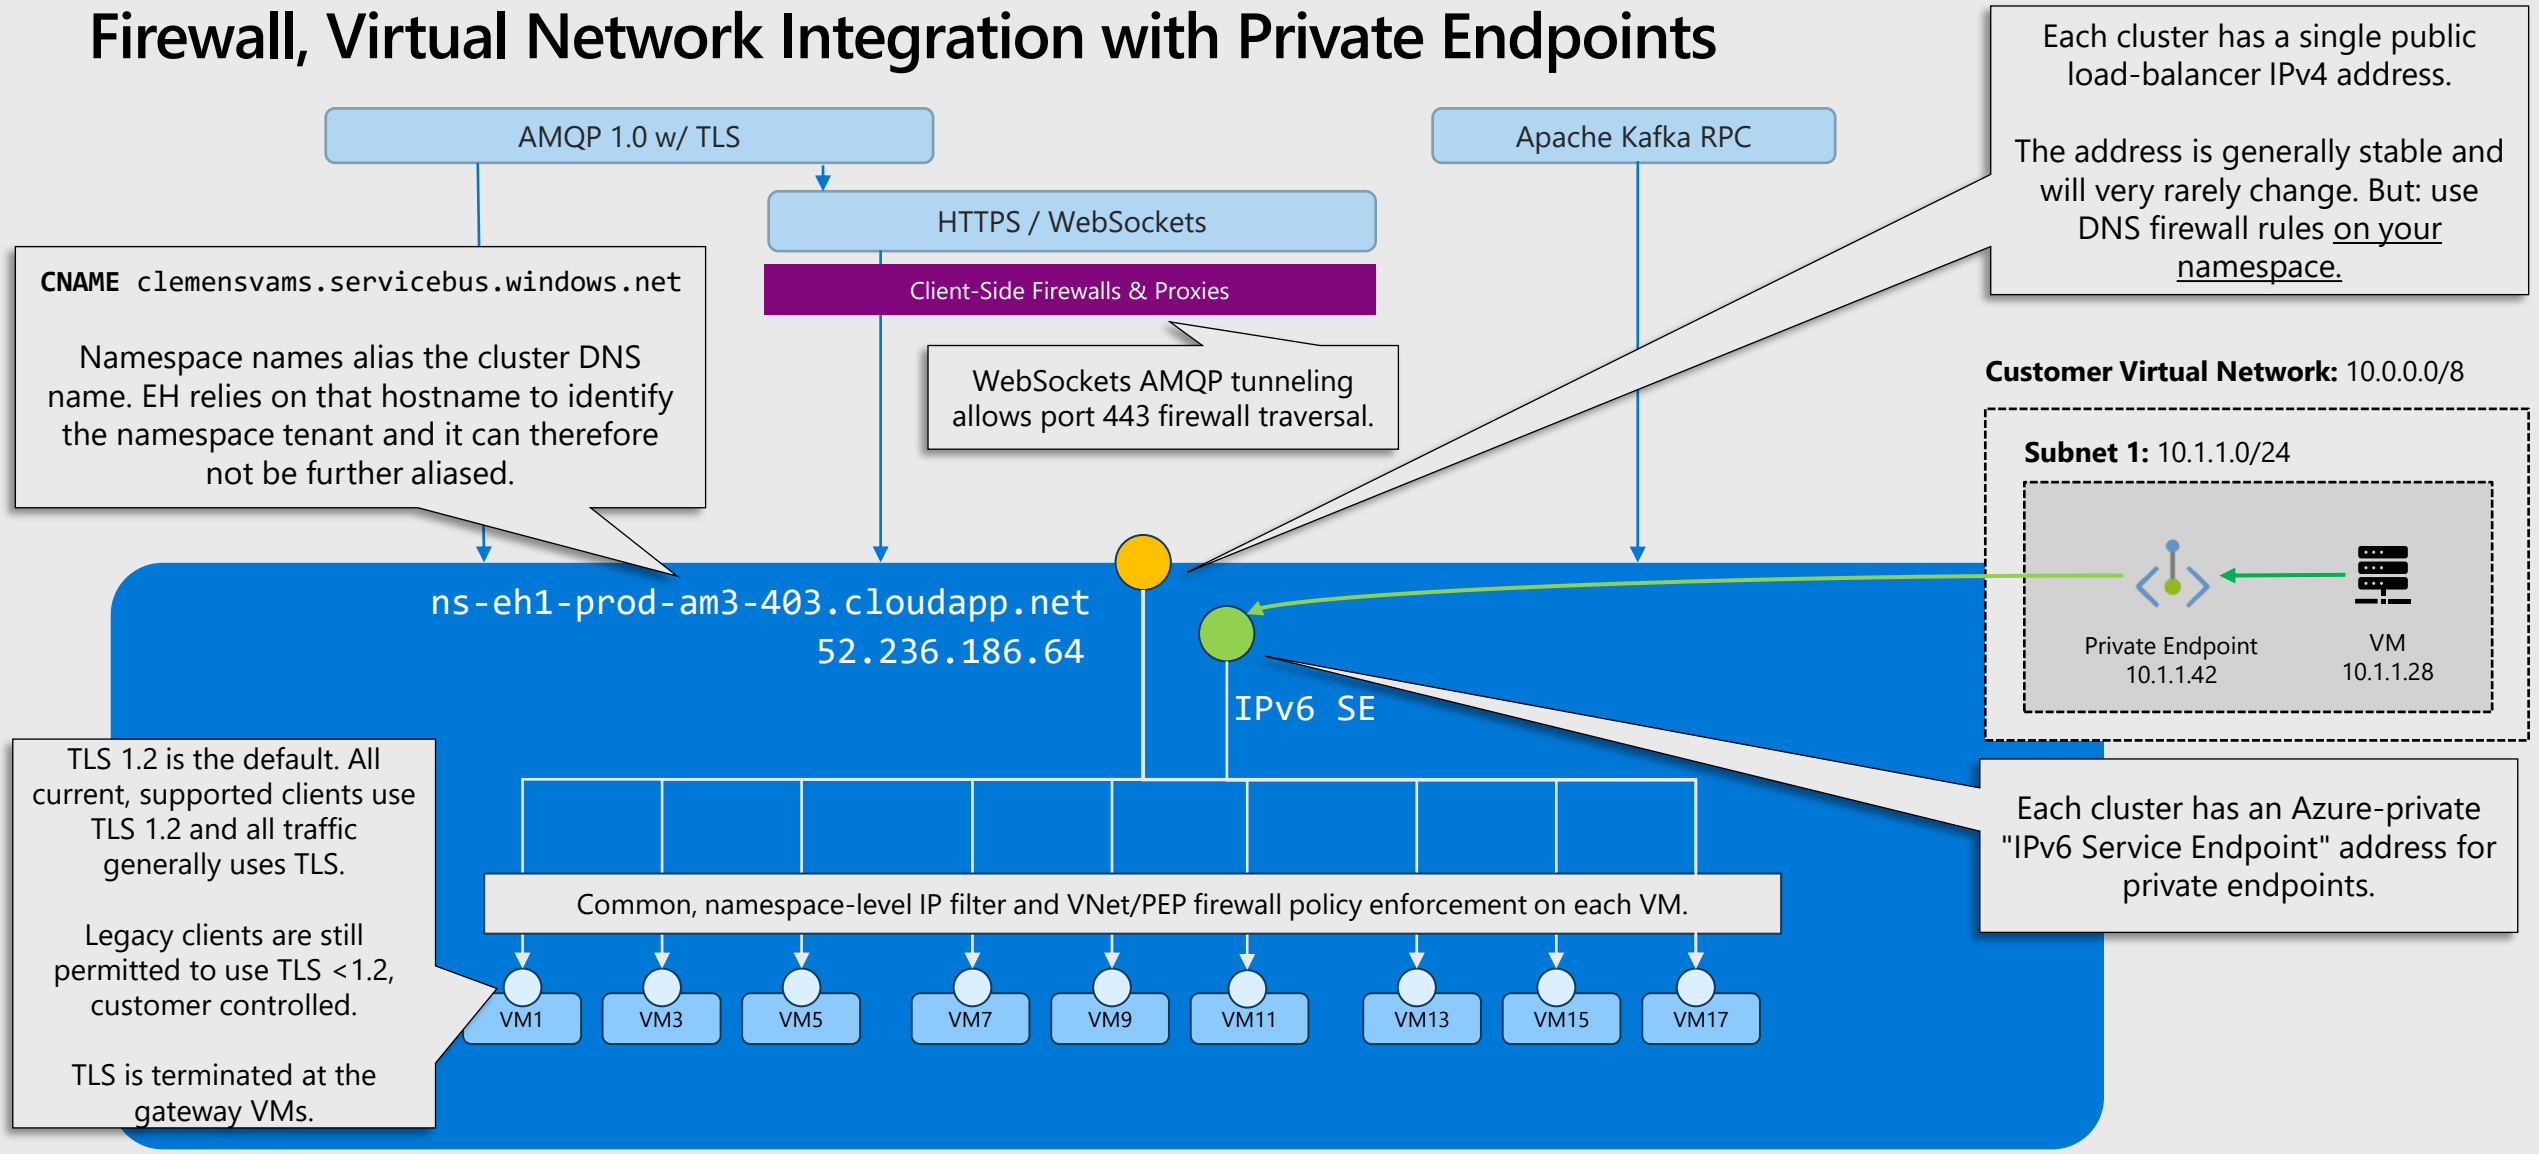

Networking Features

Event Hubs namespaces can be attached to one or more virtual networks and the public IP address space concurrently and act as safe "Layer 7" stream bridges.

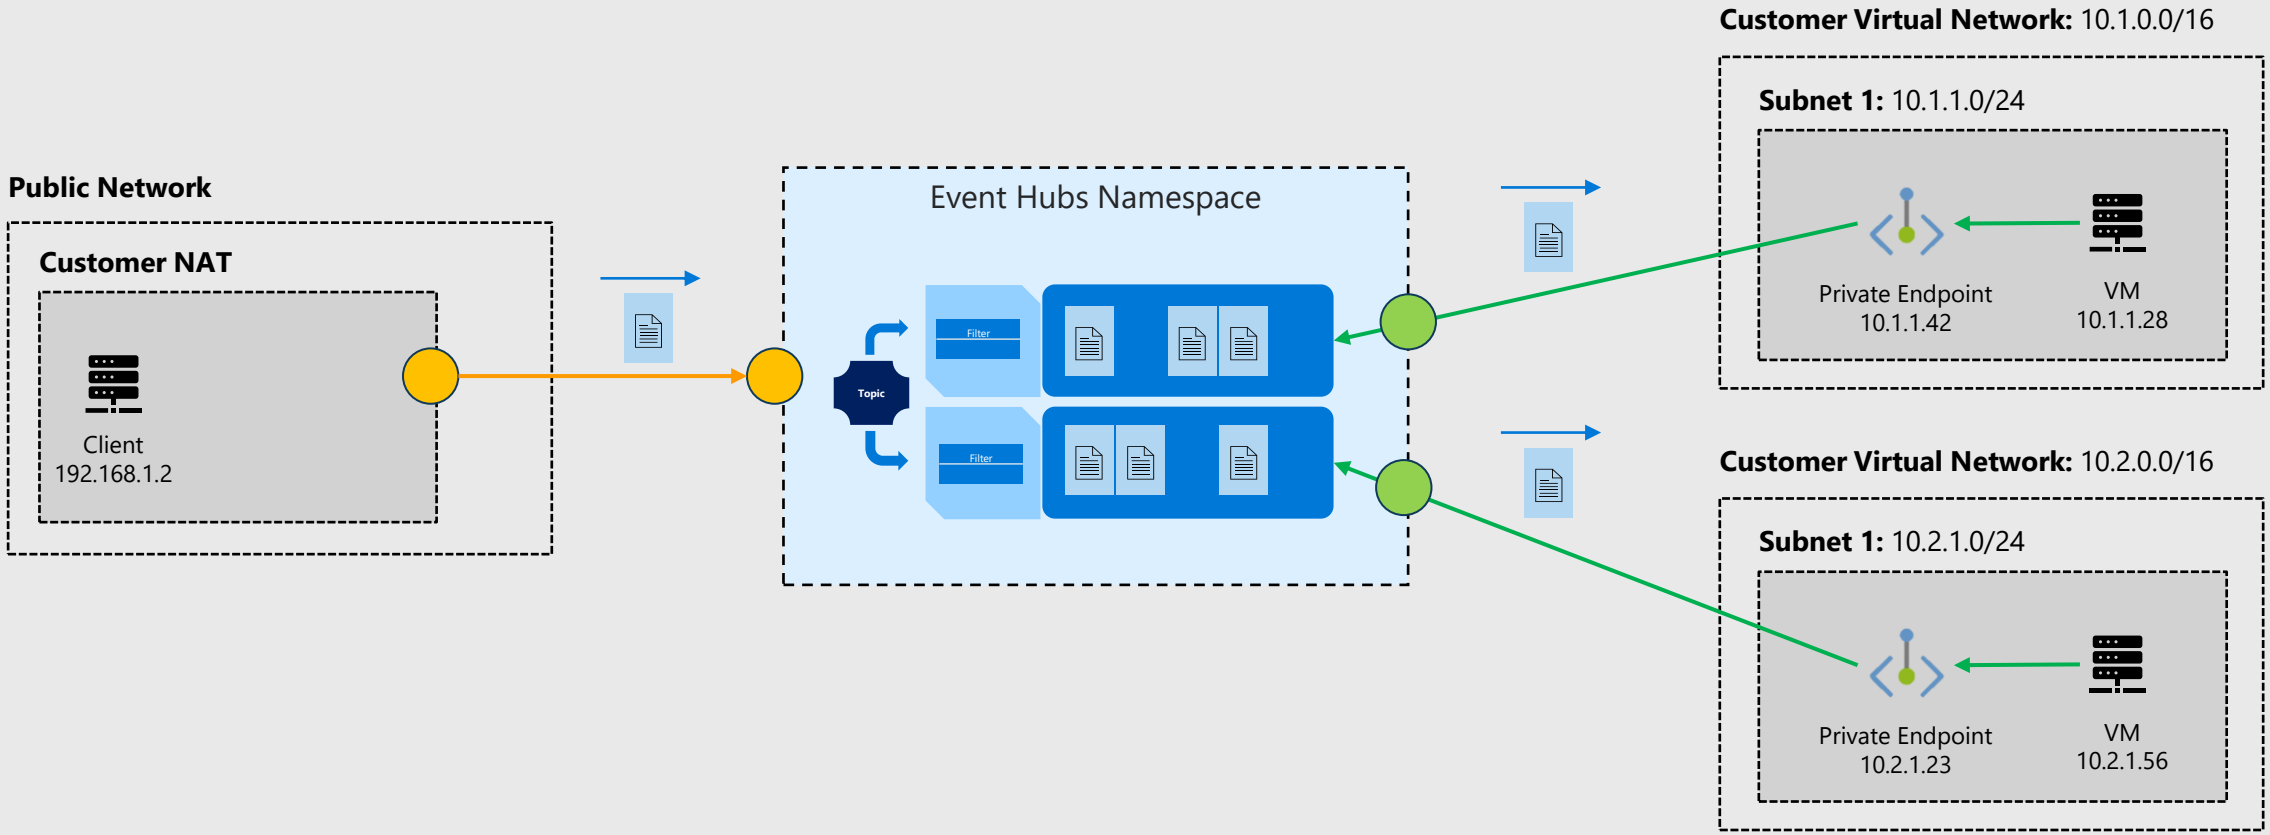

Authorization

Azure Active Directory RBAC and Shared Access Signatures (SAS)

Event Hubs has two authorization models.

SAS (Local): Simple model for external clients that cannot use Azure Active Directory for any reason.

- Very simple. Rights at namespace or entity scope: "Send", "Listen" (Receive/Subscribe), "Manage"
- No accounts. 12 named rules per scope, each rule has a key. Rules combine a set of rights.
- HMACSHA256 signed **tokens** are issued using name and key and passed to the service.
- Tokens can be issued in the client or by some security token service holding the key
- *Can be turned off via the portal on the overview page under "Local Authentication"*

AAD RBAC: Integrated Azure identity model, also for service-level integration (managed identity)

- Very detailed, operation-level permission set that can be assigned to roles
- Several standard control- and data-plane roles (like Data Owner, Data Sender, Data Receiver)
- Roles can be assigned at the namespace and entity level

Authorization

Azure Active Directory RBAC and Shared Access Signatures (SAS)

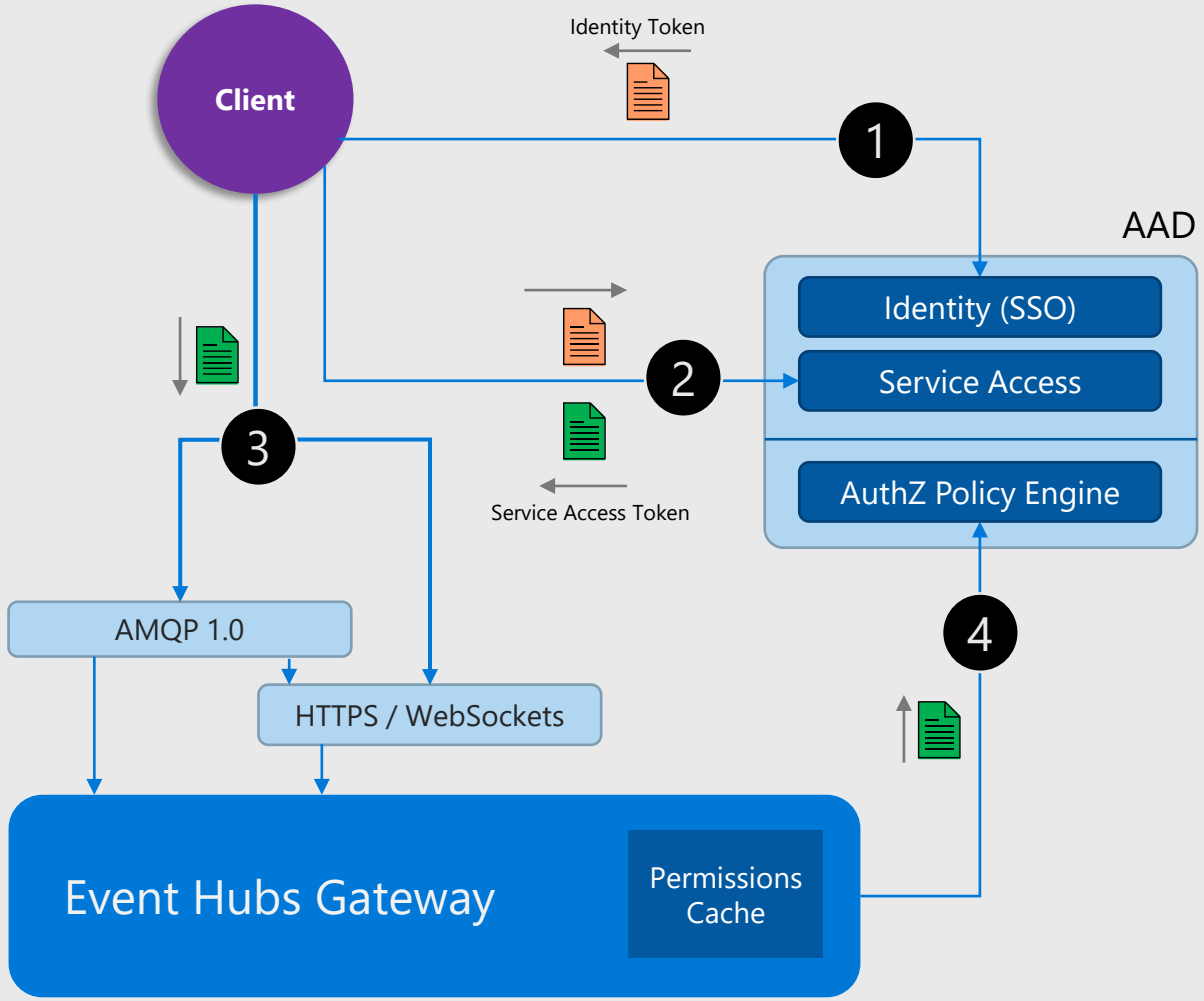

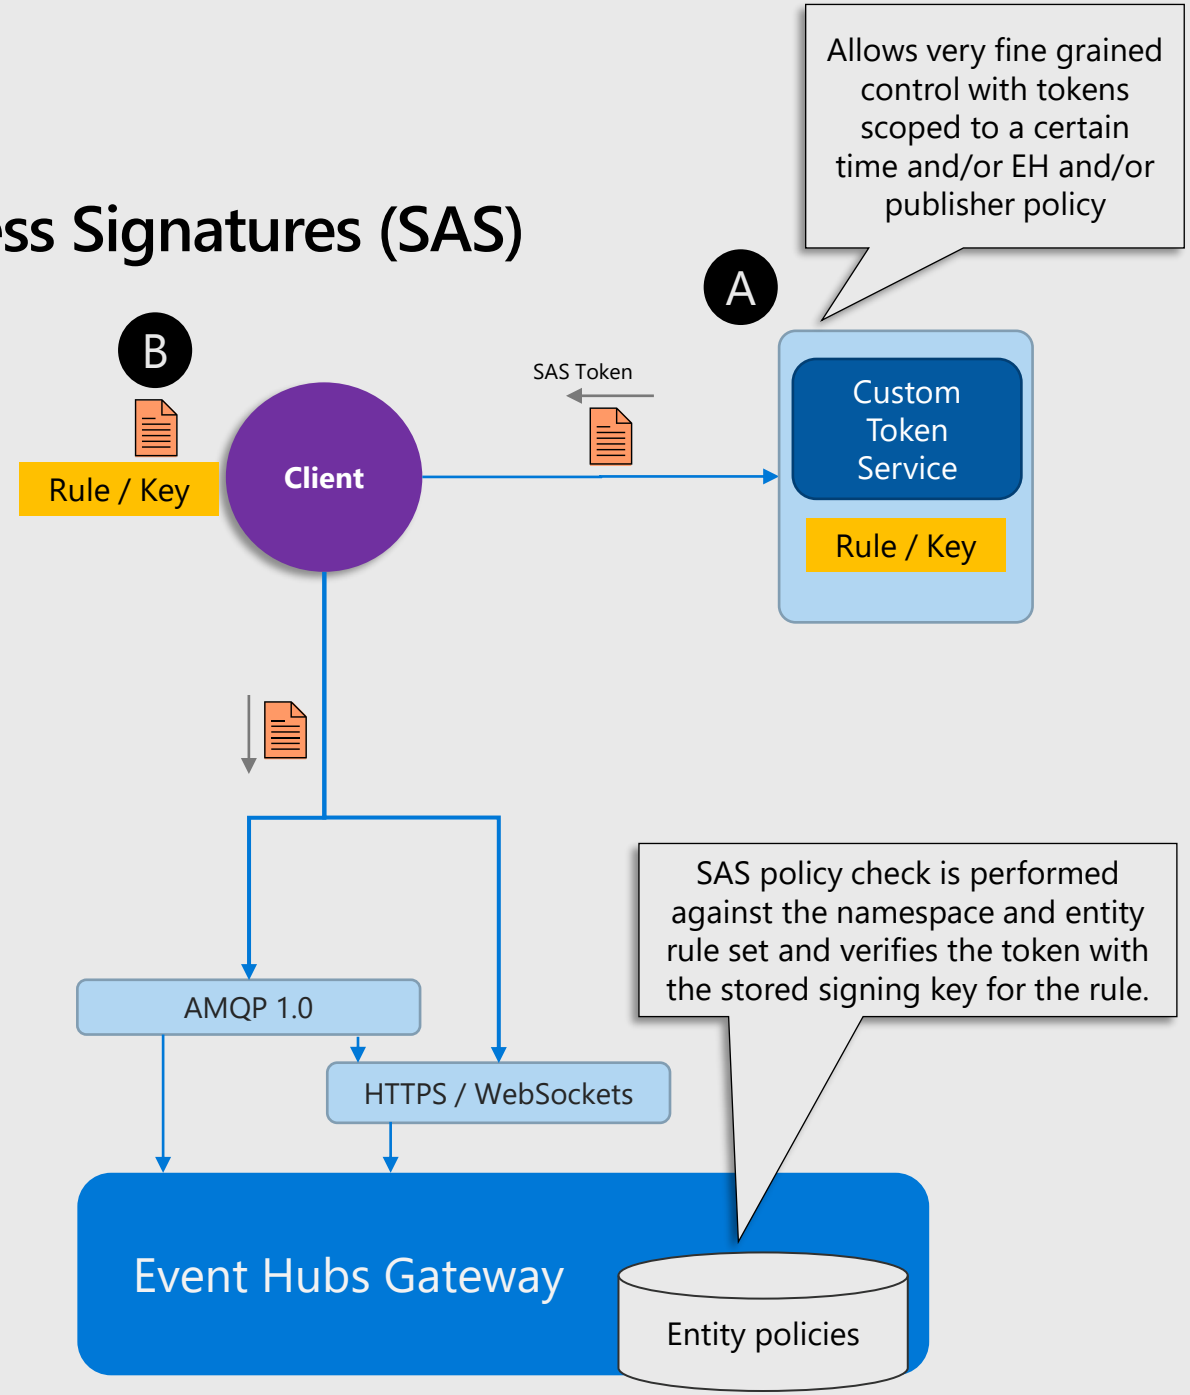

Event Hubs Premium

Replication & Federation

Replication Tasks

Providing code-free replication of data between Service Bus and Event Hubs

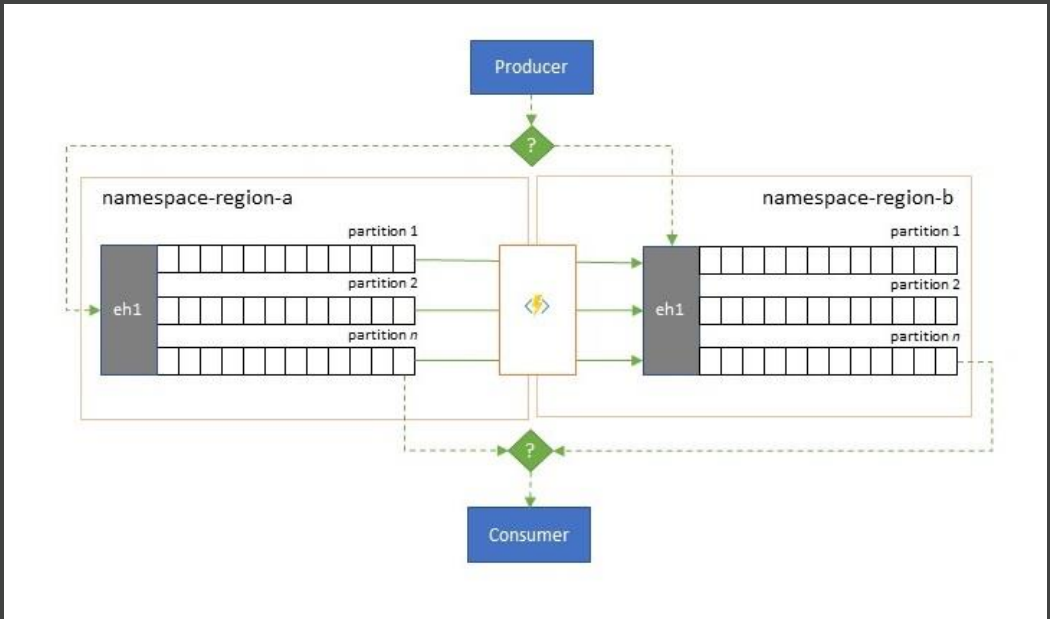

Problem

In many scenarios, event data needs to be moved or merged across regions.

Solution

Quick and simple replication between namespaces through templates available on Azure Portal's Tasks.

Automation Tasks to the rescue

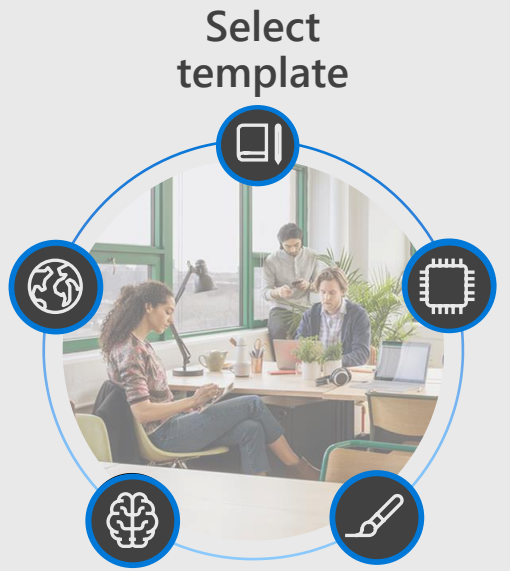

EventHub replication visualization

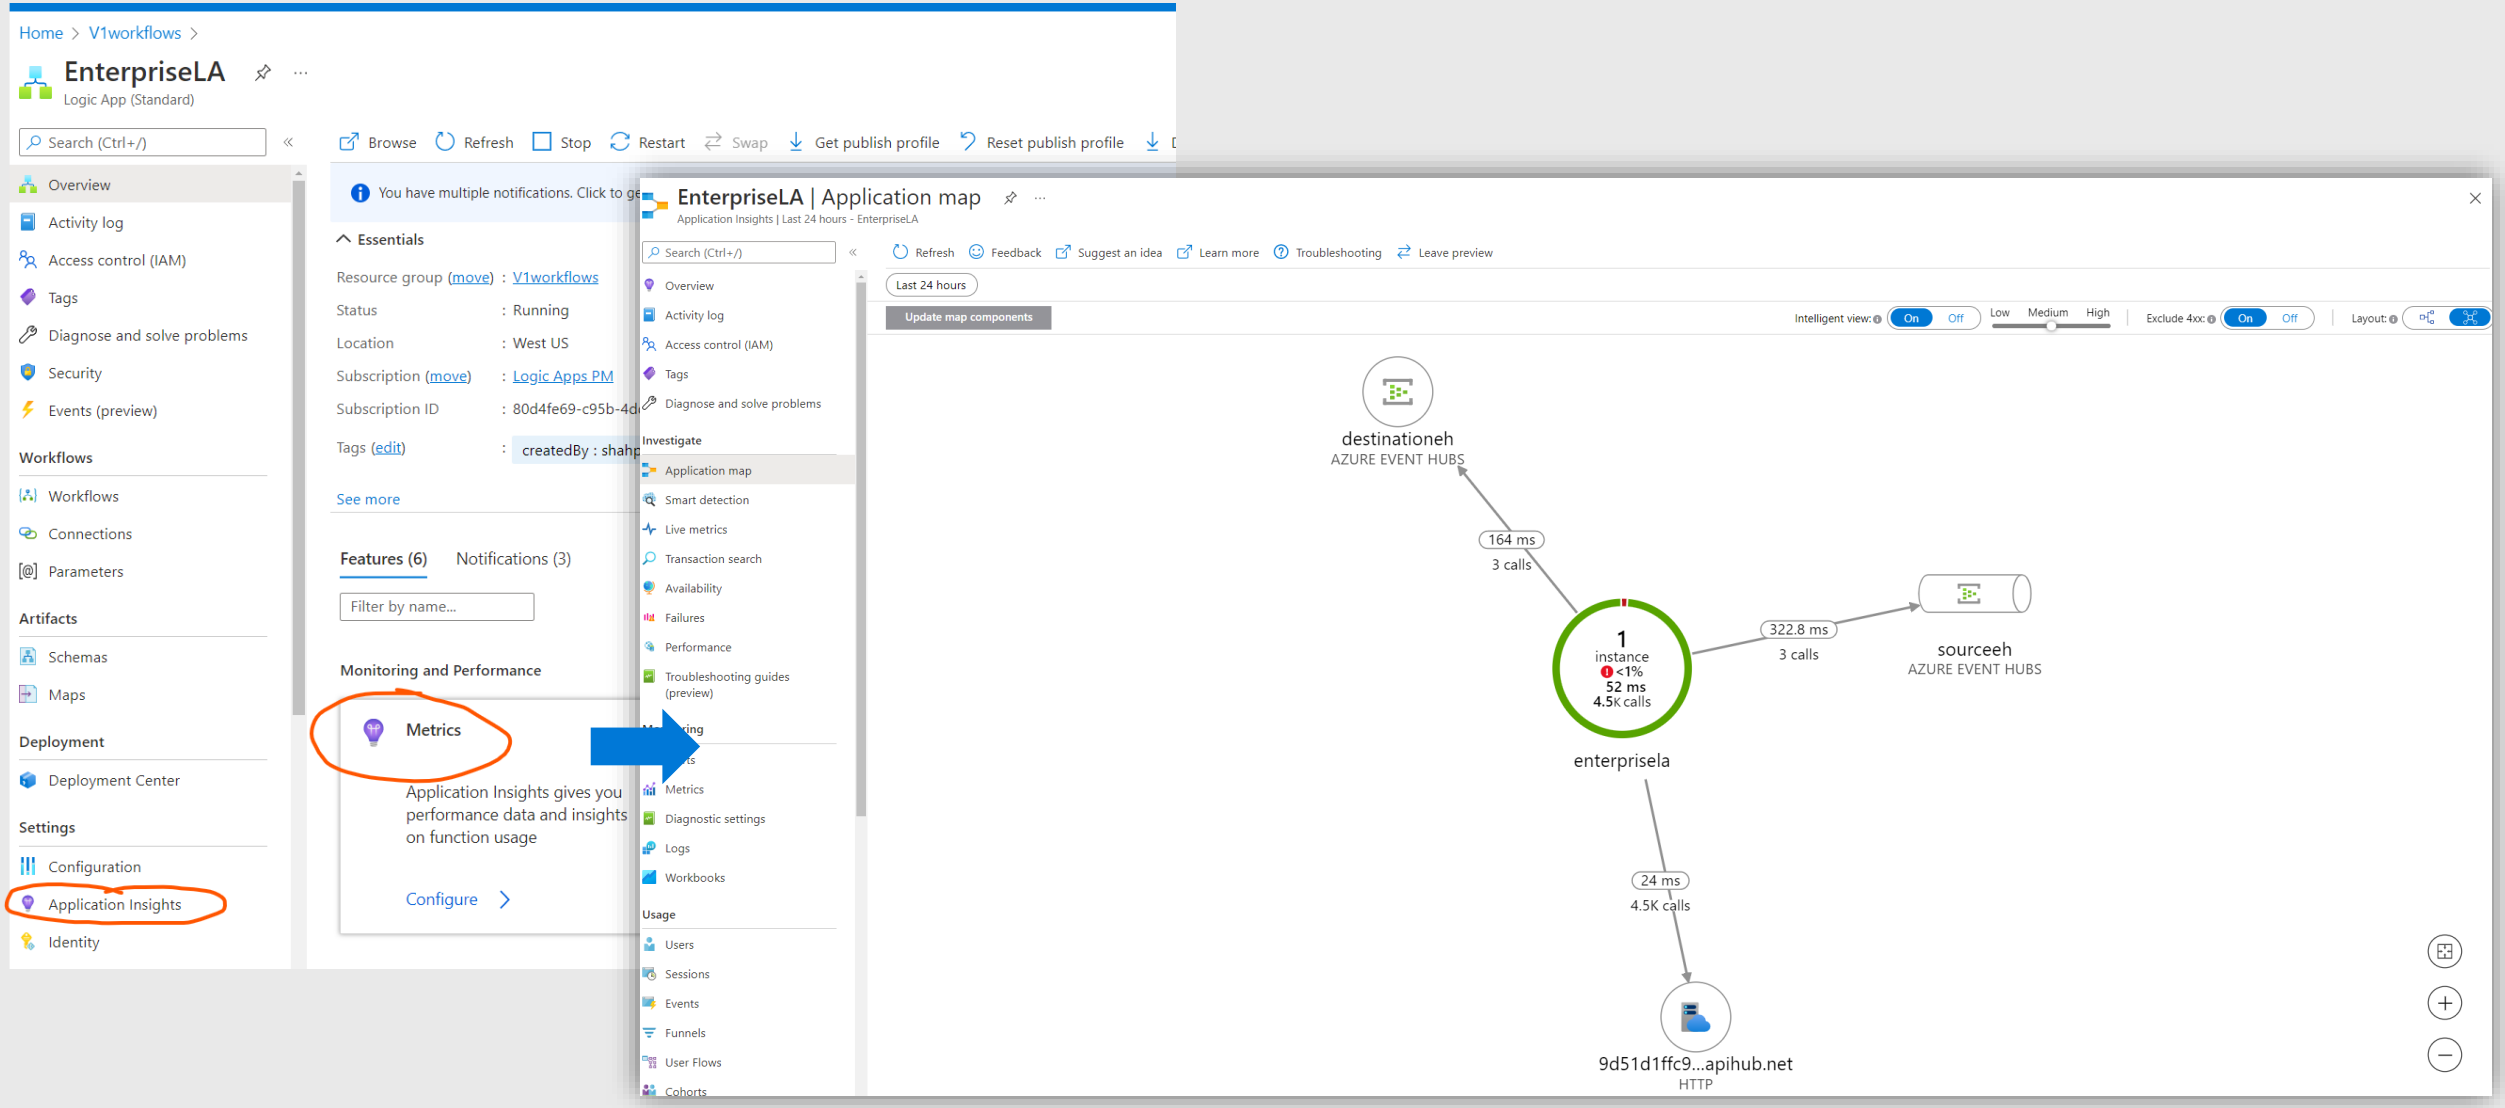

Event Hubs Premium

Performance!

Event Hubs Premium is *Fast*!

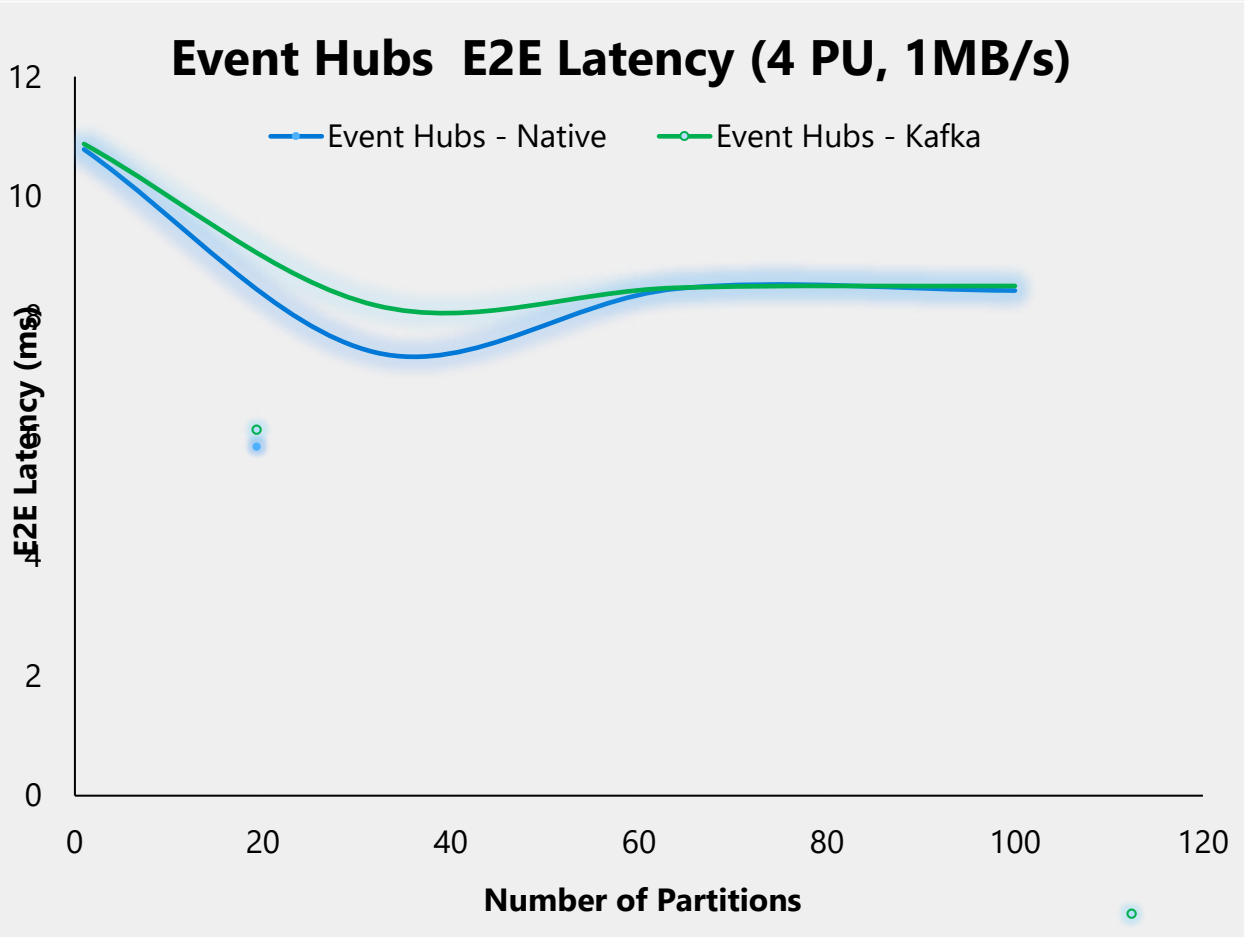

Event Hubs Premium end-to-end latency < 10ms for most Event Streaming workloads.

Performance Metric	Expected Range	Comments
Throughput	~5-10 MB/s per 1 PU	Maximum throughput can be achieved by scaling partitions For both Kafka and AMQP
Latency	<10 ms	With 1 MB/s load, 4 PU namespace For both Kafka and AMQP

Event Hubs Premium – Predictability

Event Hubs Premium offers consistent end-to-end latency of < 10ms.

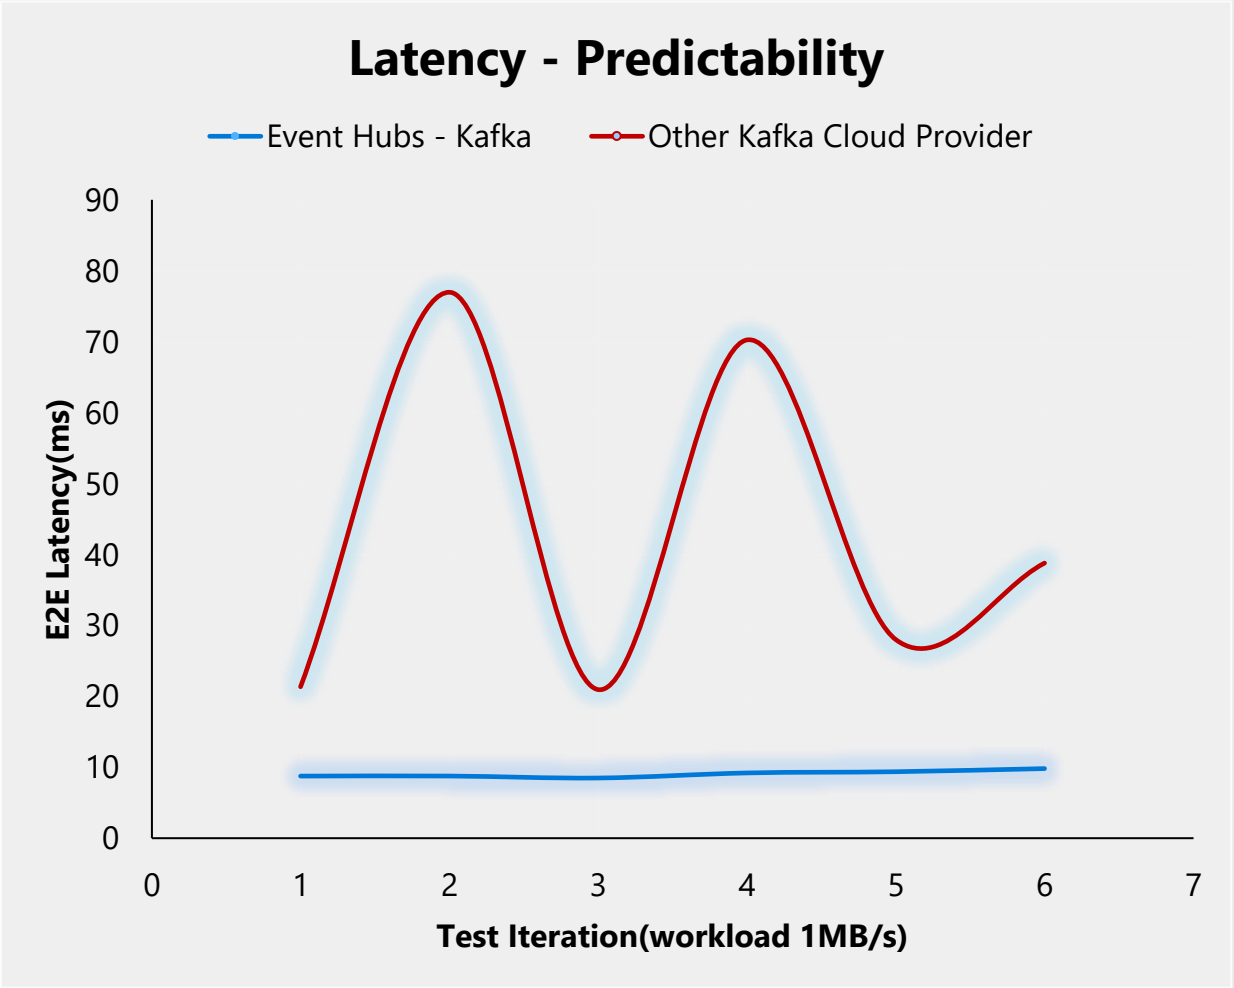

Effective Global Reliability

Uptime – Is the cluster available to accept connections?

100%

Reliability – Are individual operations suceeding?

>99,9999% weekly average

Azure Event Hubs

Platform-as-a-Service Event Stream Broker

Use the Apache Kafka® API, but with far lower cost and better performance.

Fully managed: You use the features, Azure deals with everything else

AMQP 1.0 standards compliant, Apache Kafka® wire-compatible

Polyglot Azure SDK and cross-platform client support

Industry-leading reliability and availability

Fast.

That was a lot, wasn't it?

Q&A

© Copyright Microsoft Corporation. All rights reserved.