cdc.dev

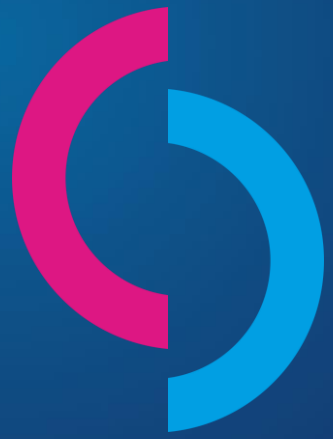

CARIBBEAN DEVELOPERS CONFERENCE

A Top-Notch tech conference
right in the middle of paradise!

@caribbeandevconf

Clemens Vasters
Microsoft Corporation
Chief Messenger

@clemensv

CARIBBEAN DEVELOPERS CONFERENCE

Flow Architectures with Azure Event Hubs

https://github.com/clemensv/messaging

Event-Driven Flow

Observe, Decide, Act, Tell

Imperative Flow

Listen, Obey, Execute, Respond

Imperative

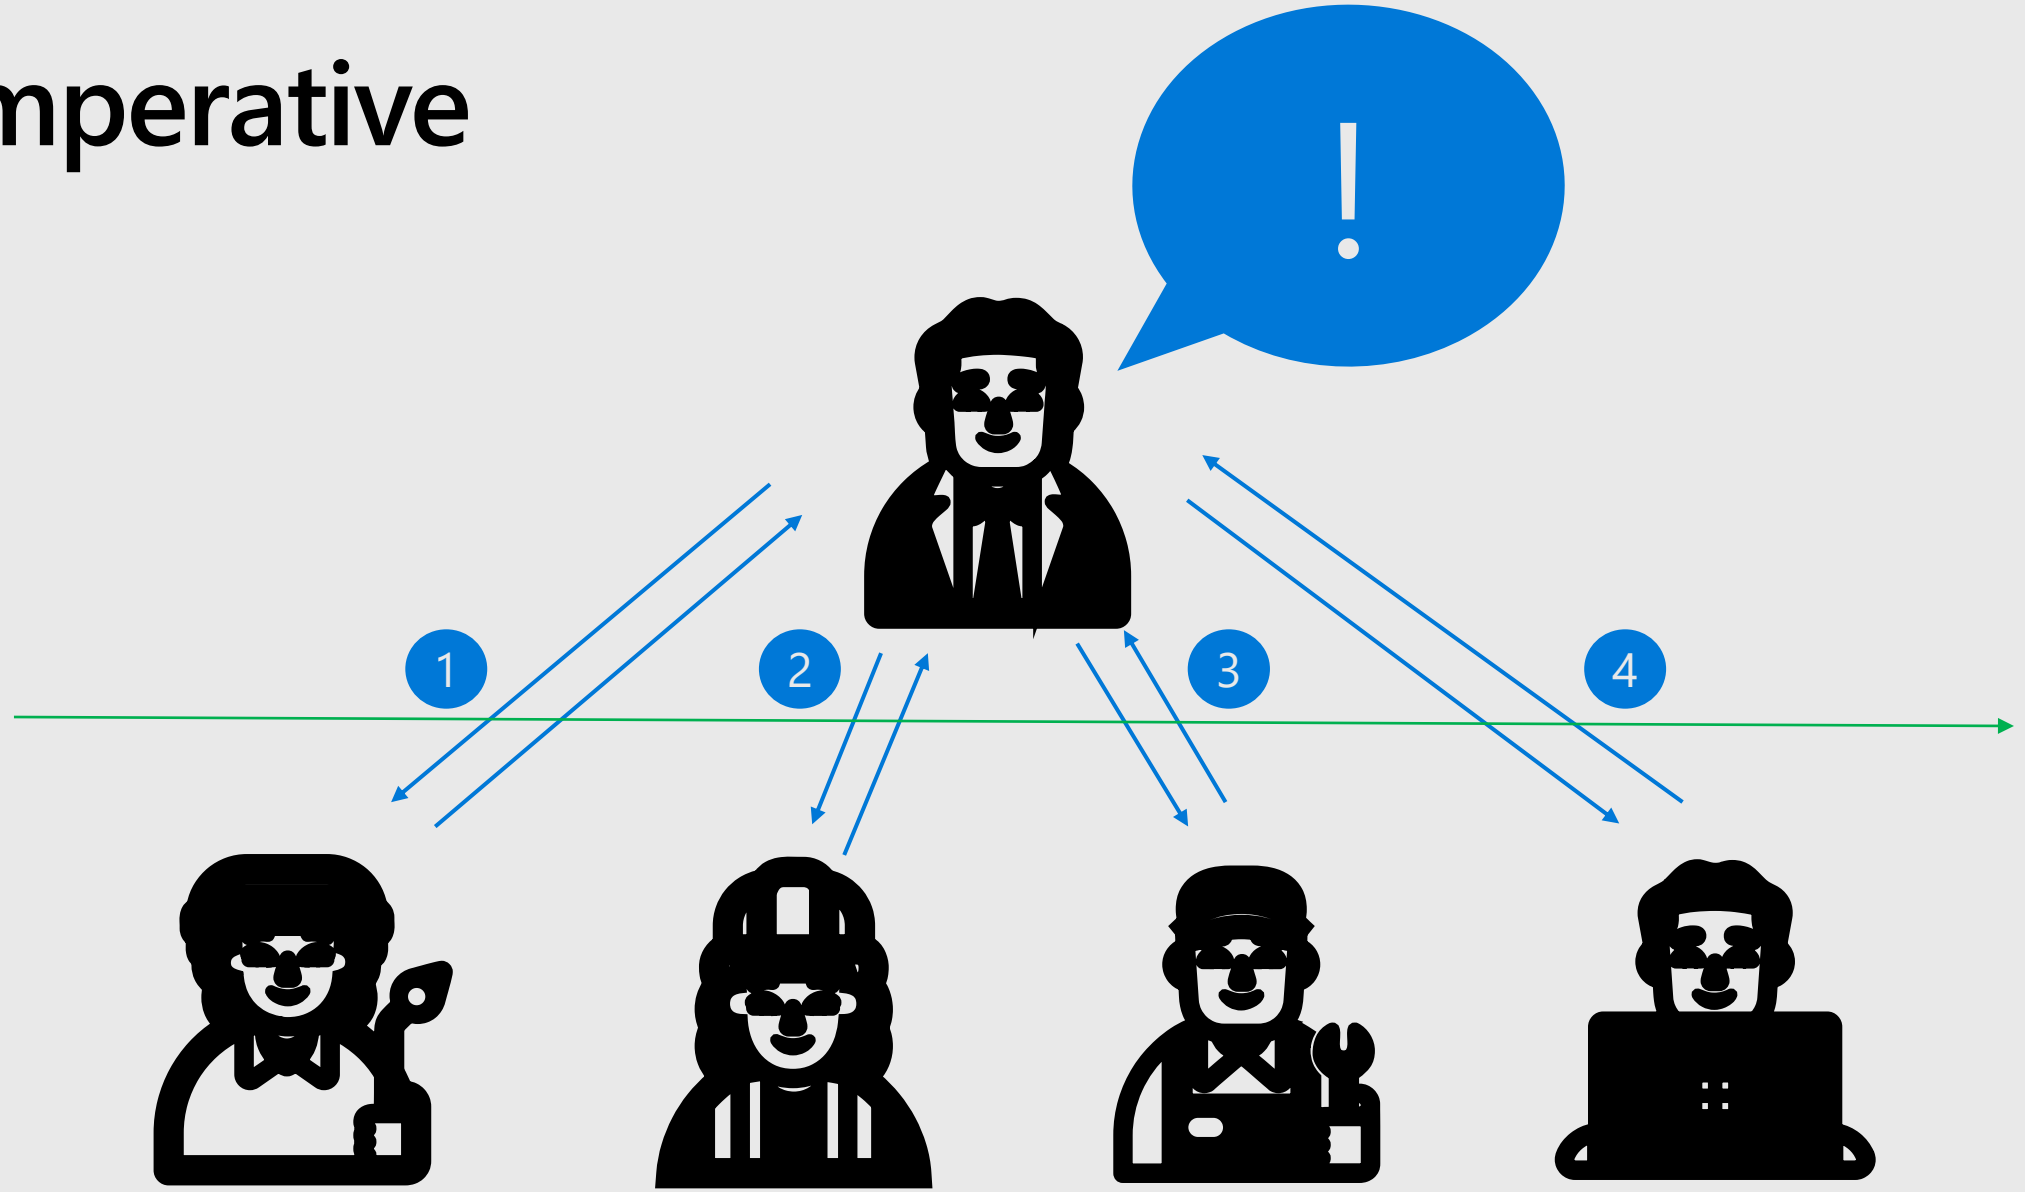

Event-Driven

Imperative

Execution plan with predetermined actor roles

Work is assigned, executed, and completion reported.

Changes and extensions require modifying the execution plan.

Command & Control

Event-Driven

No predetermined execution plan or actors.

Work is volunteered reacting to events and completion is shared

Changes and extensions are local to actors or add/remove actors

Dynamic Collaboration

Most complex systems require both models in different areas

TICKETING INTEGRATION FLOW

BORUSSIA-PARK, MÖNCHENGLADBACH

DIE FOHLEN

TICKETING INTEGRATION FLOW

BORUSSIA / EVENTIM / AXESS

Automated access control: Axess

ticket codes

Ticket portal: Eventim

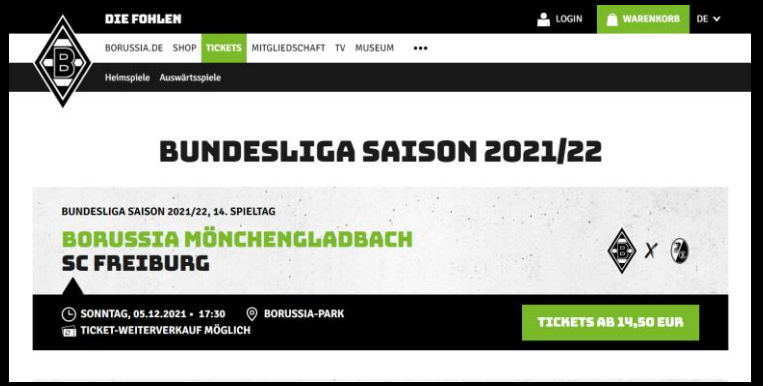

ticket sales

customer info

ERP

Live Reporting

customer info

Single-Sign-On
Profile

Signals, Streams, and Jobs

***Signal*:** The capture of an **occurrence** (statement of fact) during the operation of a software system

***Event*:** A data record expressing a signal and its context**.** The context is expressed in metadata annotating the signal.

***Event Stream*:** A chronological sequence of events belonging to the same context.

Job: *Not an event*. A description of a task that needs to be performed by some party. Preferably just once.

Signals


Calendar entry created
Meeting scheduled
Meeting attendee accepted
Inventory item added to shelf
Payment collected
Calendar appointment reminder
Meeting attendee invited
Inventory item removed from shelf
Inventory item added
Application deployment initiated
Change committed
Change proposal verification succeeded
Change proposed
Change proposal commented
Inventory item scanned for sale
Aircraft rotated
Application deployment succeeded
Mouse cursor moved
Sales lead added
Pizza ordered
Aircraft position report
Application deployment failed
Vehicle passed induction loop
Eye focus tracked
Delivery address changed
Pizza prepared
Aircraft squawk code change
CO PPM over threshold
Room humidity read
Advertisement clicked
Database record inserted
Database record updated
Room temperature read
Smoke detected
Virtual machine started
Engine oil pressure read
Package scanned
Pizza delivered
Aircraft landed
Stock trade order issued
Venue seat ticket reserved
CPU utilization read
Package delivered
File changed
Object detected ahead
Infrared movement detected
Engine coolant temperature read
Engine RPM read
Stock trade order executed
Account transfer scheduled
Venue seat ticket released
Package picked up
File created
Venue seat ticket purchased
Object velocity differential calculated
Visual object classified
Engine oil pressure threshold alert
Tank pressure read
Package loaded
Package routed
Stock trade order settled
Account transfer completed

Events: Observations, Signals, Jobs

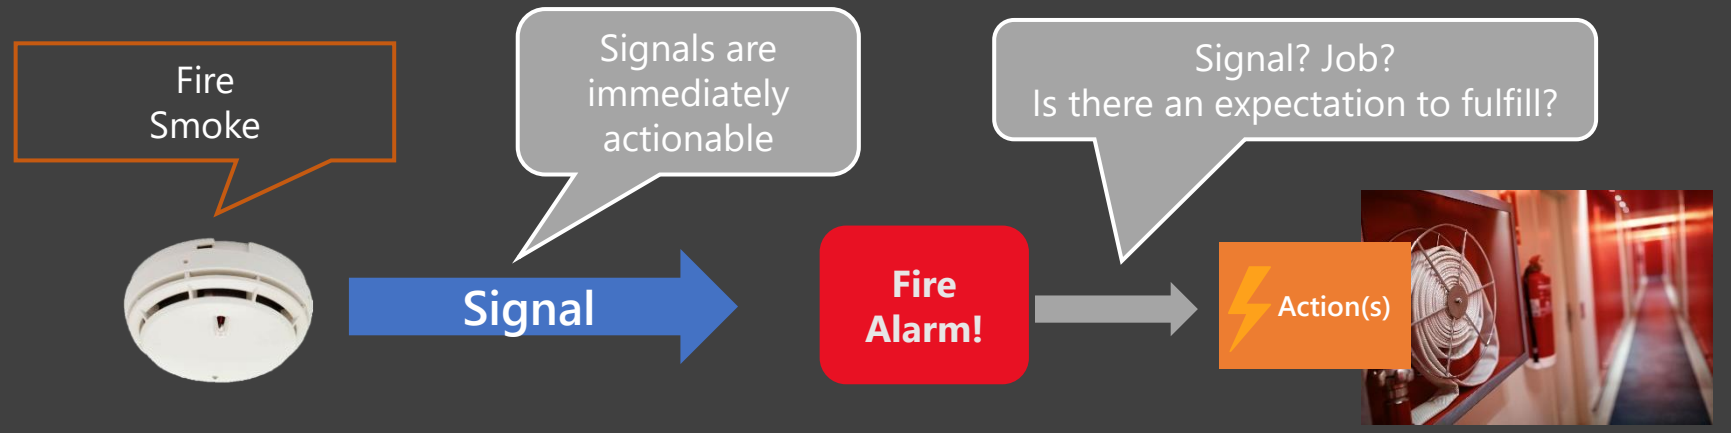

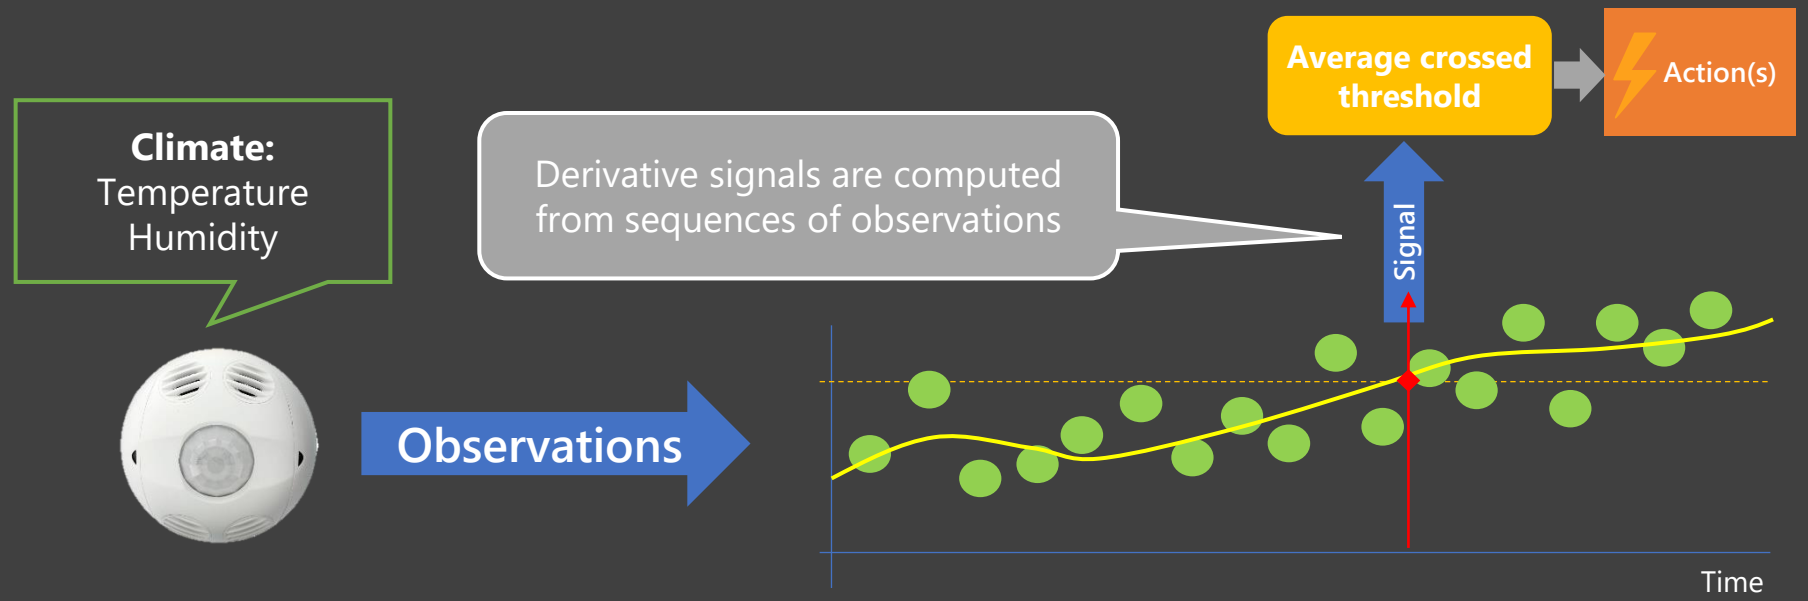

Discrete Event Router

Azure Event Grid, AWS Event Bridge, Knative Eventing

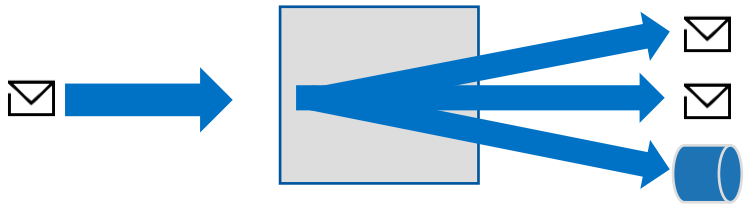

Push-style distribution of discrete events to serverless workloads or other messaging infrastructures

Queue Pub/Sub Broker

Azure Service Bus, AWS SQS/SNS, Google PubSub, Apache ActiveMQ, RabbitMQ, IBM MQ

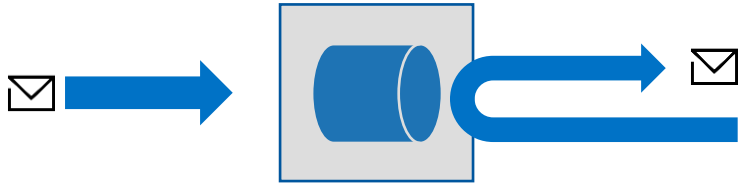

Pull-style, queue-based transfer of jobs and control via message queues and topics

Event Stream Engine

Azure Event Hubs, AWS Kinesis, Apache Kafka, Apache Pulsar, CNCF Pravega

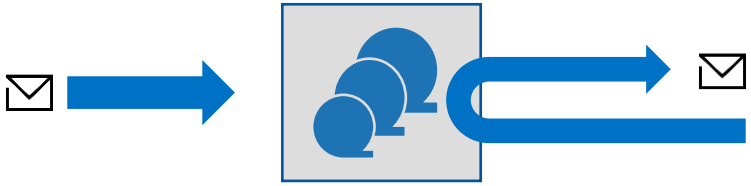

Partitioned, high-volume, tape-drive-style sequential recording and unlimited, pull-style re-reads of event streams.

Event Stream Aggregator

Azure Stream Analytics, AWS Kinesis Analytics, Apache Samza, Apache Flink, etc.

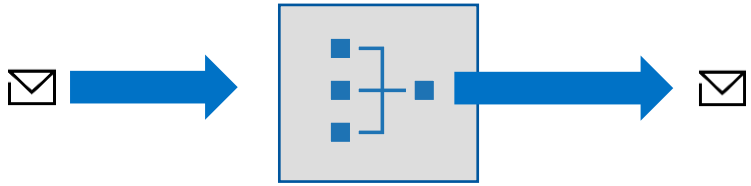

Stateful processing of event streams yielding event streams and discrete events as continuous output

Event Streams and Time(-liness)

Event Data Value – Securities Markets

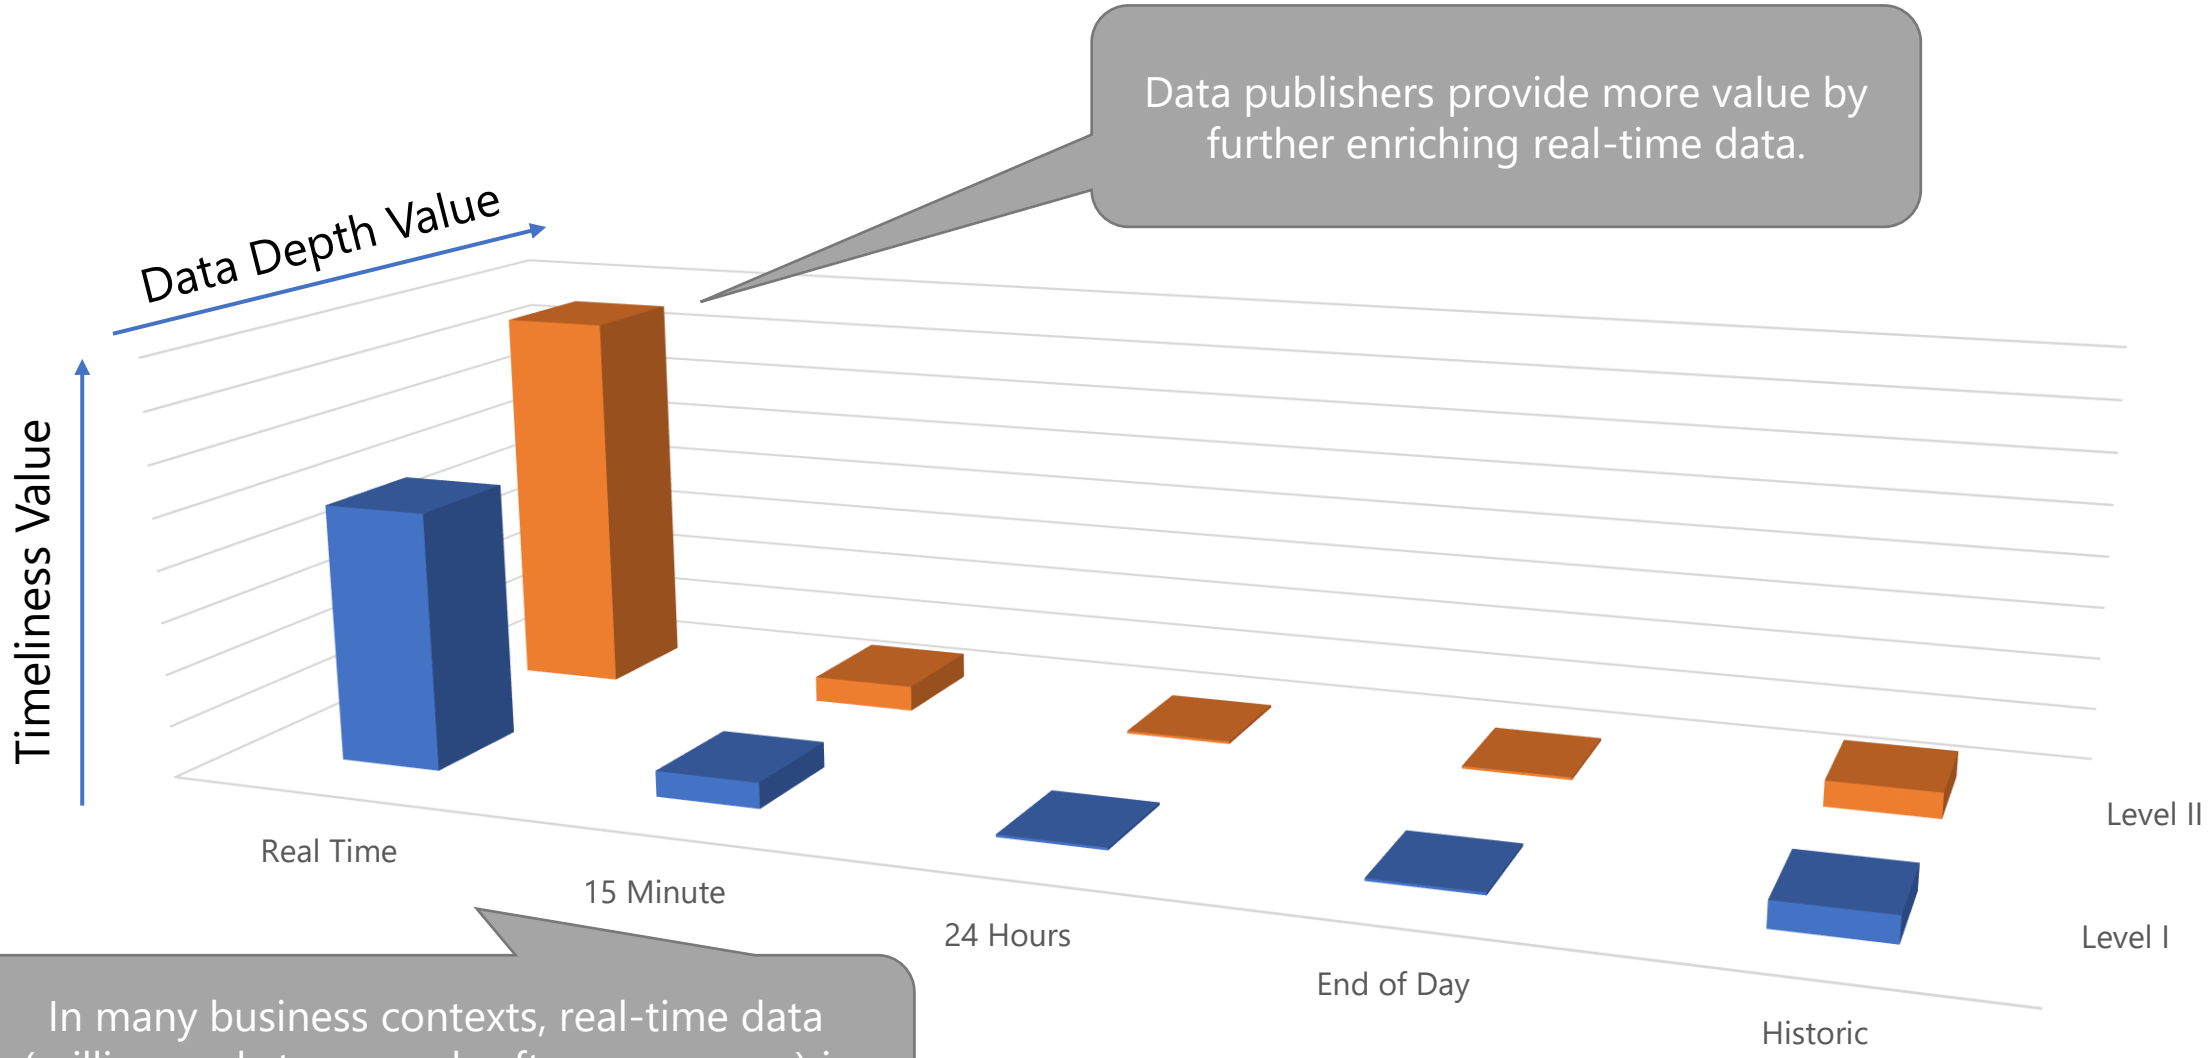

Velocity Matters ➔ Parallelization Matters

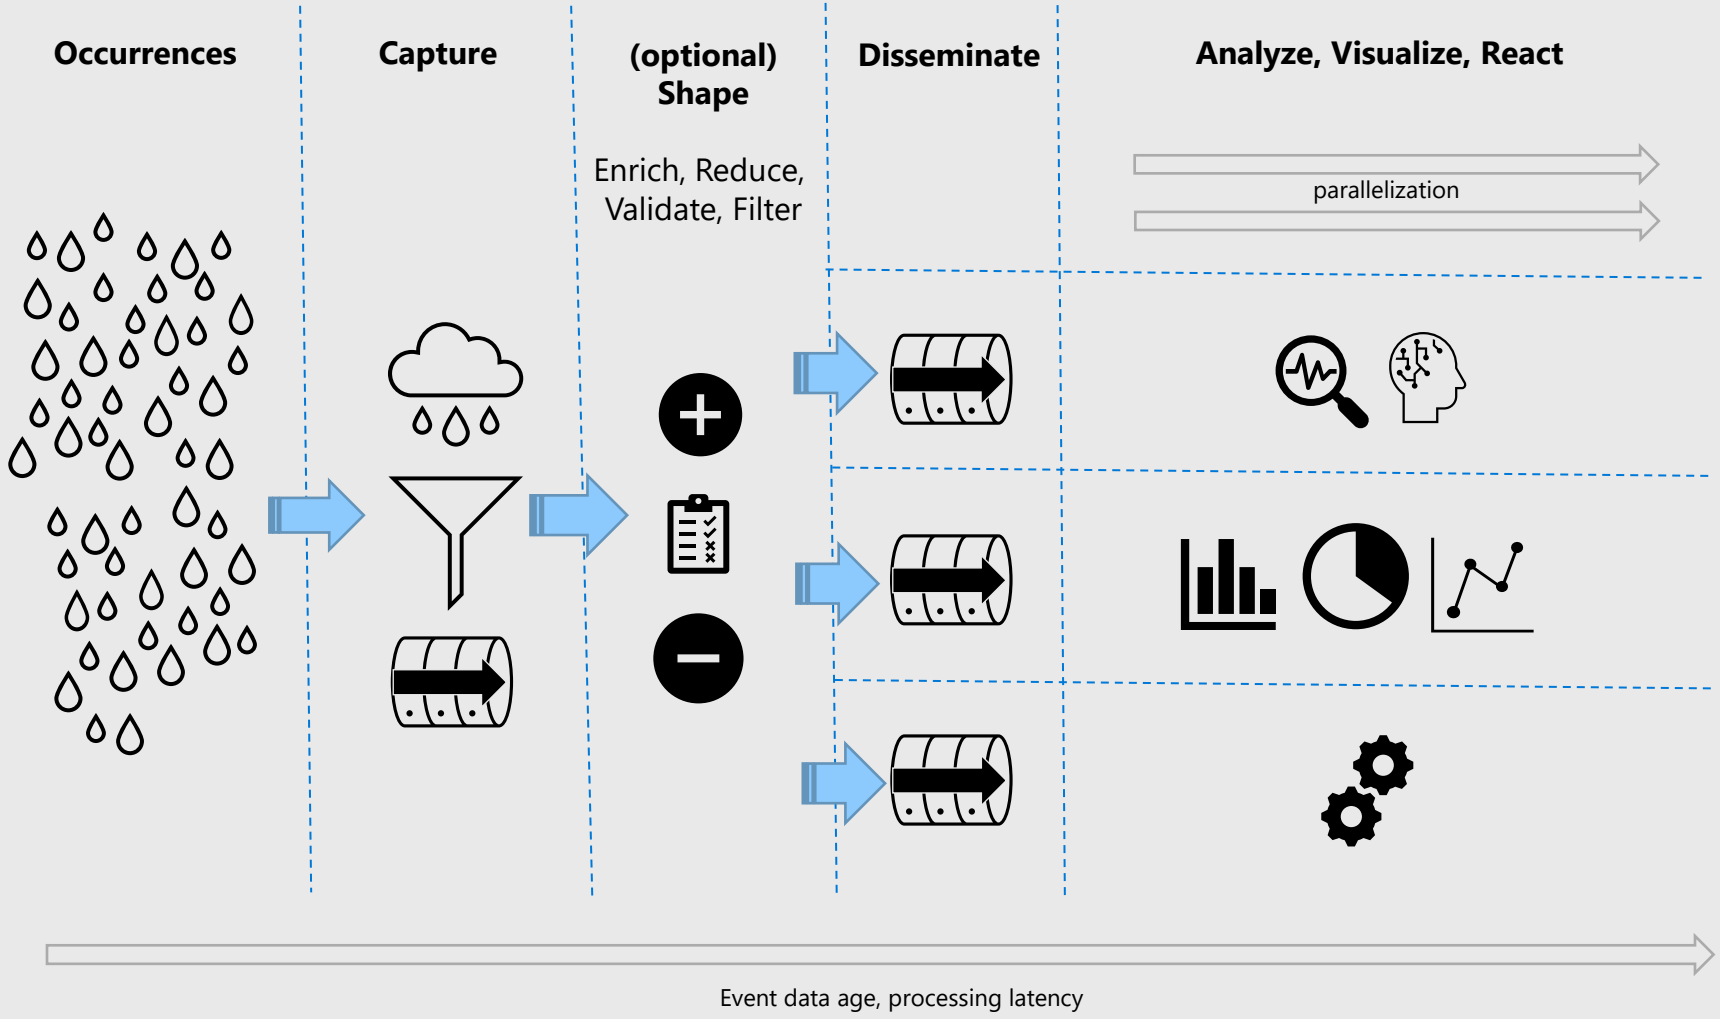

What is Azure Event Hubs?

Platform-as-a-Service Event Stream Broker

Use the Apache Kafka® API, but with far lower cost and better performance.

Fully managed: You use the features, Azure deals with everything else

AMQP 1.0 standards compliant, Apache Kafka® wire-compatible

Polyglot Azure SDK and cross-platform client support

Industry-leading reliability and availability

Fast.

Azure Event Hubs vs. Apache Kafka®

Similar yet very different

	Azure Event Hubs	Apache Kafka
User Model	Partitioned event stream broker with high-availability replication	Partitioned event stream broker with high-availability replication
Architecture	Multi-tenant, 3-Tier Gateway/Broker/Storage cluster model, with tenant-isolation, all tiers independently scalable	Single-tenant monolith. Need to increase broker instances in a cluster to scale any dimension.
Implementation Language	C# and Native (C/C++)	Java
Cluster Manager	Azure Service Fabric (inline)	Apache Zookeeper (external); KRaft (inline, experimental)
Partition Mapping	Key hashing, client or server-side mapping of events	Key hashing, client-side mapping of events
Consumer Partition Ownership Coordination	Server-coordinated partition ownership (Kafka), client-coordinated ownership with external leader election. Parallel, direct partition reads.	Server-coordinated partition ownership
Server Workload Balancing	Dynamic and fully automated (100% hands-off). Broker resource allocation independent of partition count or ownership, flexible scaling.	Static assignment of partitions to broker instances requiring operator intervention for rebalancing.
Storage Model	Replicated log store, synchronous per-message flush-to-disk on all replicas	Replicated log store, asynchronous flush-to-disk controlled by host file system write cache settings.
Networking	Single endpoint access to all partitions, Public IP/DNS or Virtual Networking, Firewall.	Endpoint per broker instance. Multiple IPs required. Complex network management required.
Access Control	Token-based access policy model with unlimited publisher policies, Azure Active Directory role-based access control	Local accounts, federation extensibility.
Protocols	AMQP 1.0 (optional: WebSockets) HTTPS 1.1 Apache Kafka RPC	Apache Kafka RPC
Batching / Archives	Avro-packaged batch-packaging and archival to blob store	
Schema Registry	Schema Registry based on open CNCF Schema Registry API	(Proprietary from commercial vendors)

Event Hubs

A log partition is a durable, triple-replicated, append-only storage structure. Each partition has a throughput ceiling imposed by IO and compute limits. More partitions provide more throughput capacity, but also make consumption more complex.

A consumer group is an optional mechanism to allow multiple parallel consumers to coordinate who owns which partition.

Kafka RPC: Broker assigns partitions and checkpoints can be stored in the CG.

AMQP: Clients choose leader using an external mechanism and can break leases using an epoch model. Checkpoints are external. AMQP with server-assigned partitions and CG checkpoints on roadmap.

Producers send using a protocol of their choosing, with or without setting a partition key. Partitions can also be targeted directly.

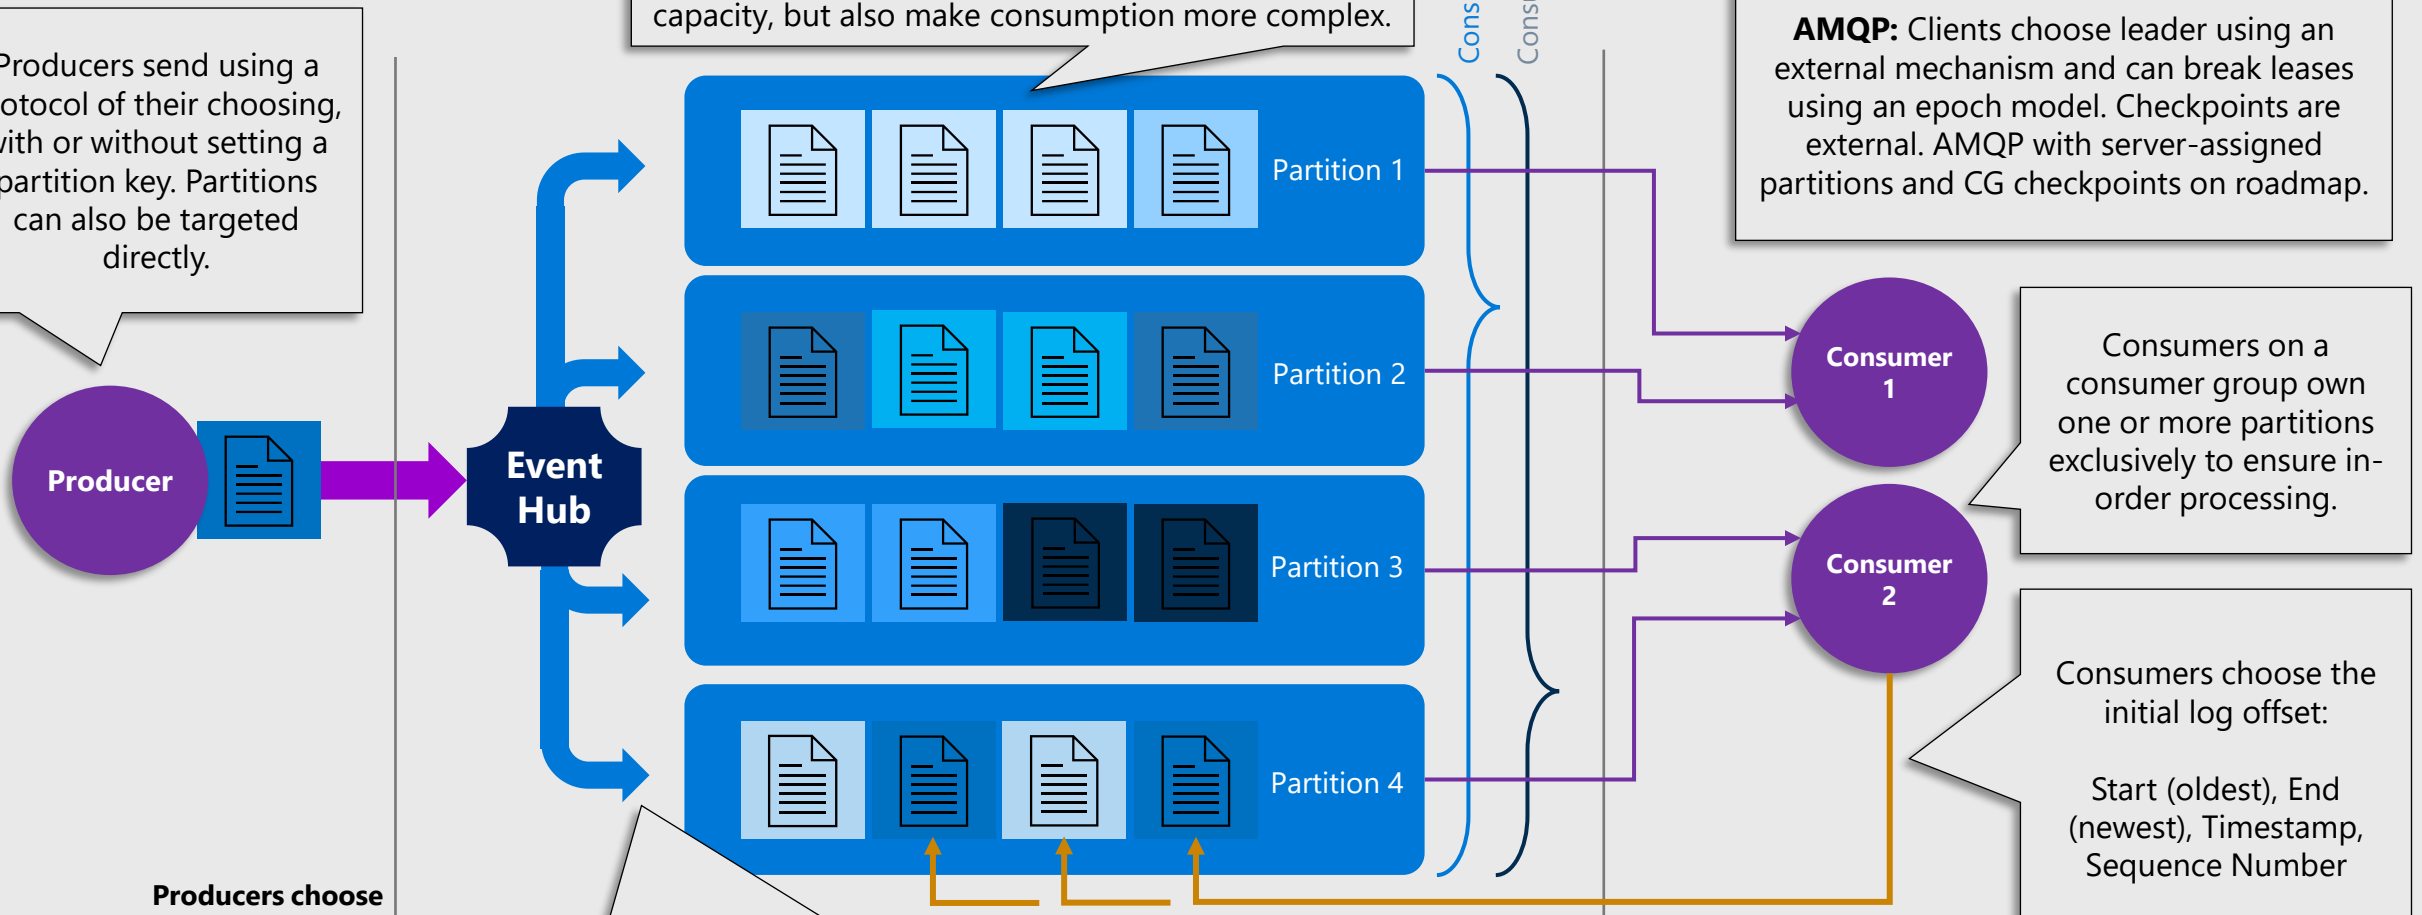

Consumers on a consumer group own one or more partitions exclusively to ensure in-order processing.

Consumers choose the initial log offset:

Start (oldest), End (newest), Timestamp, Sequence Number

Read progress is noted in checkpoints that can be stored anywhere.

Producers choose their preferred protocol

OASIS AMQP 1.0
Apache Kafka® RPC
HTTPS 1.1

An Event Hub is an event stream manager handling one or more independent log partitions. Different from Apache Kafka®, Event Hubs can optionally assign events to partitions on the server-side.

Consumers choose their own preferred protocol

OASIS AMQP 1.0
Apache Kafka® RPC

Event Hubs Capture

{Namespace}/**{EventHub}**/**{PartitionId}**/{Year}/{Month}/{Day}/{Hour}/{Minute}/{Second}

https://mystorageaccount.blob.core.windows.net/mycontainer/mynamespace/myeventhub/0/2017/12/08/03/03/17.avro

Event Hubs Capture allows organizing torrents of telemetry in batches that can be easily consumed by batch-oriented analytics frameworks.

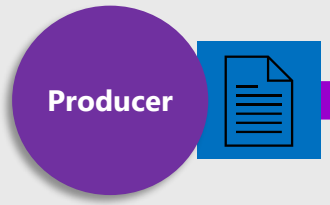

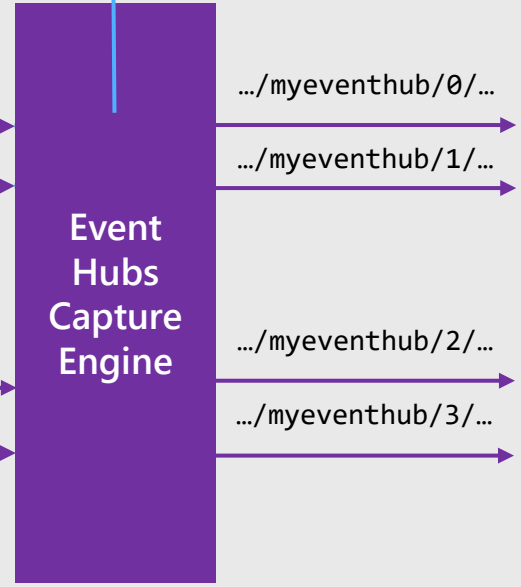

Event is raised when capture file has been written

.../myeventhub/0/...

.../myeventhub/1/...

.../myeventhub/2/...

.../myeventhub/3/...

Azure Blob Storage

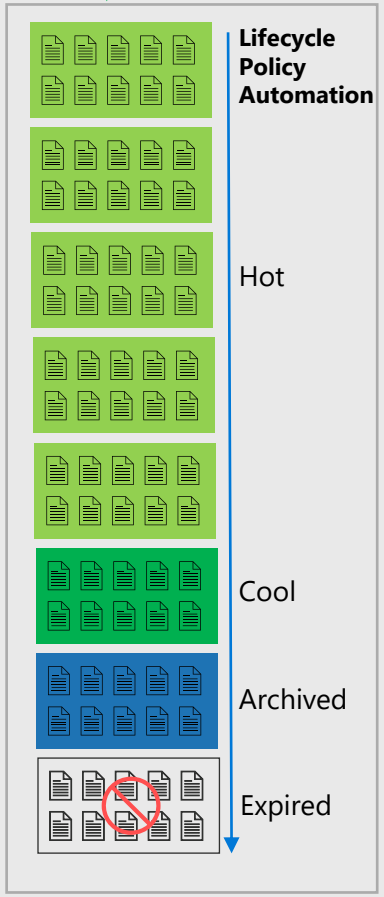

Azure Data Lake Gen2

Producers choose their preferred protocol

OASIS AMQP 1.0
Apache Kafka® RPC
HTTPS 1.1

Capture throughput keeps up with the Event Hubs' ingress throughput capacity and provides automated low latency capture to batch storage.

Capture policy allows balancing latency and batch size.

Event Hubs* is not…

* Apache Kafka isn't any of these things, either.

… a publish/subscribe broker. Partitions are not subscriptions. They are chosen by the producer or the broker on ingress. There is also no server-side filtering.

→ Azure Service Bus, Azure Event Grid

… a queue broker. Read progress over the log is handled by the client and there is no event-level ownership and delivery state handling.

→ Azure Service Bus

… a discrete event distribution engine. Event Hubs does not do push deliveries, and delivery failures need to be tracked individually.

→ Azure Event Grid

… a database or long-term event store. Event Hubs exists to catch, store and provide fast access to event data organized around time axis. As data ages (days, not months), you need better indexing.

→ Azure Cosmos DB, Azure SQL, Azure Synapse, Azure Data Explorer

Event Hubs Namespaces: Virtual Brokers

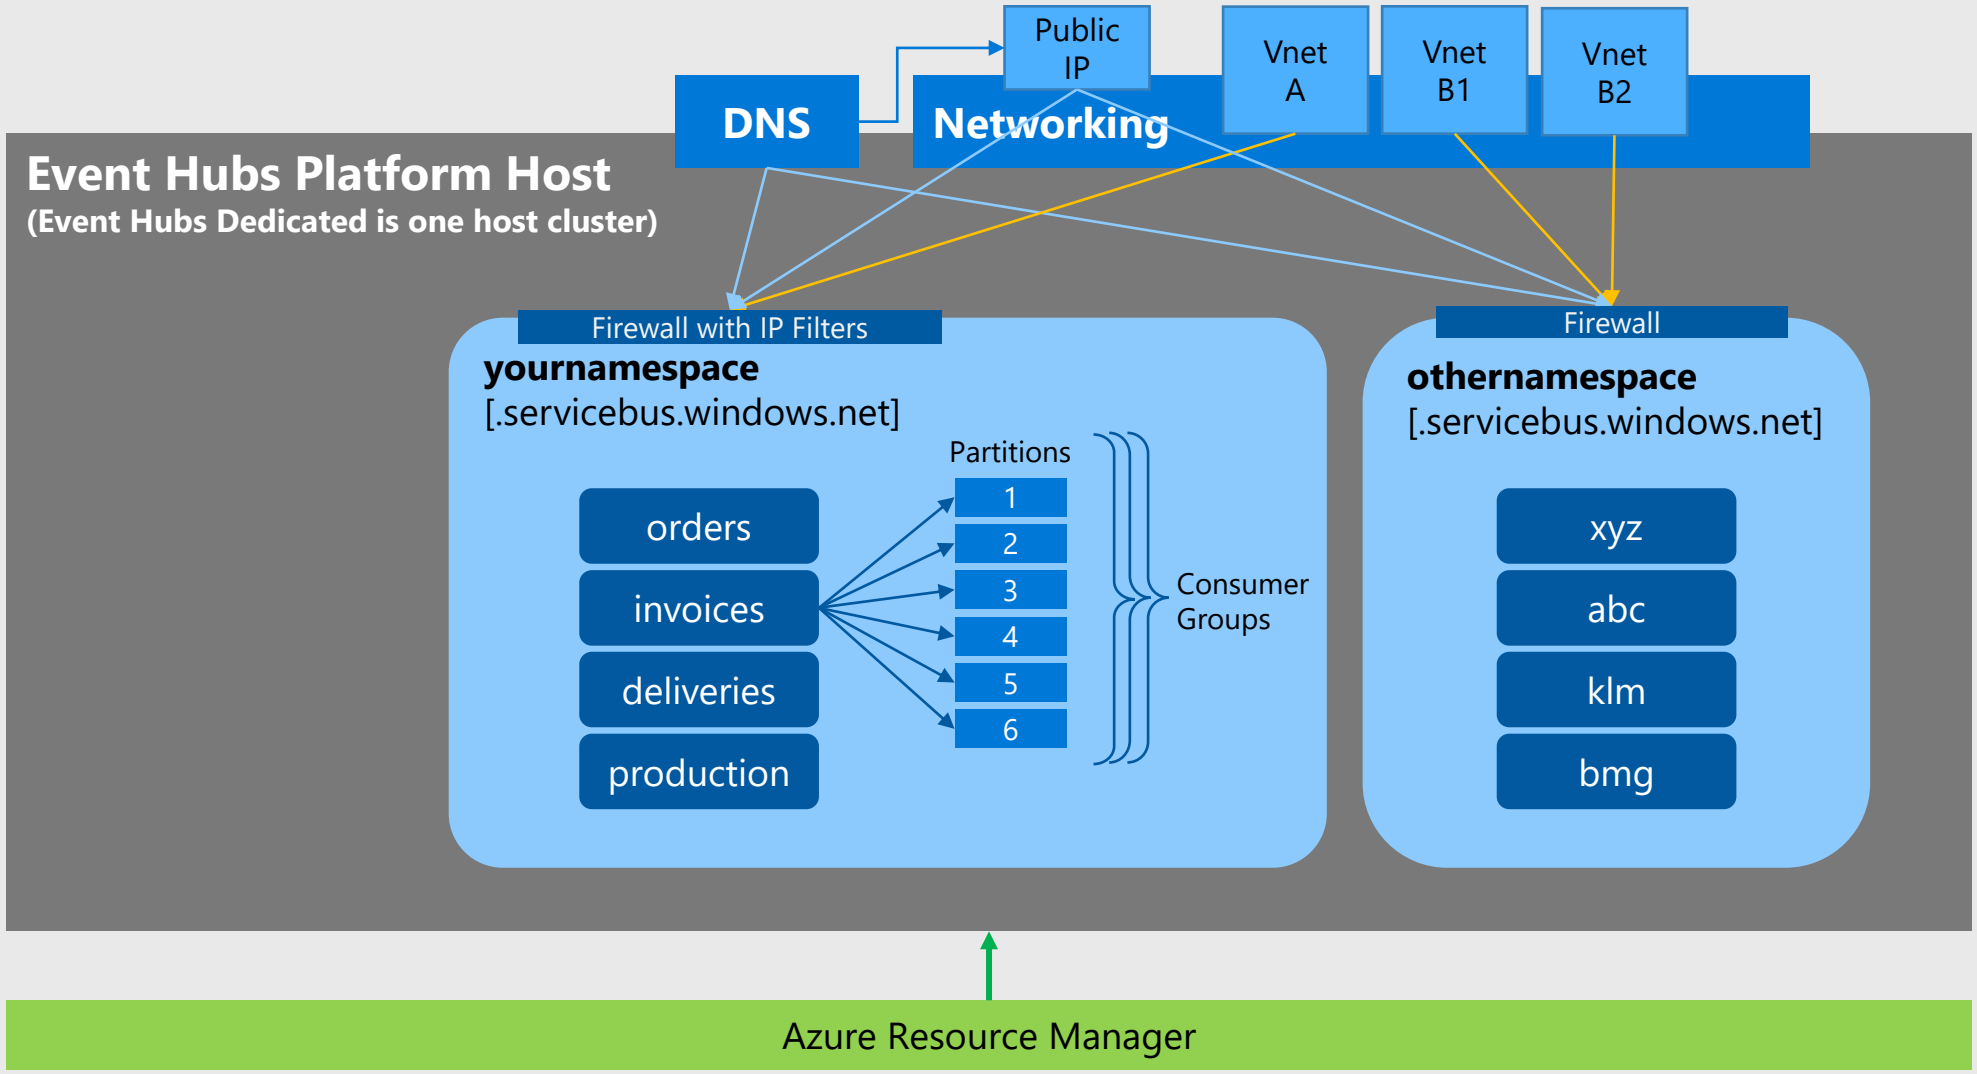

Logical Architecture

3-tier architecture: Networking, Messaging, Storage

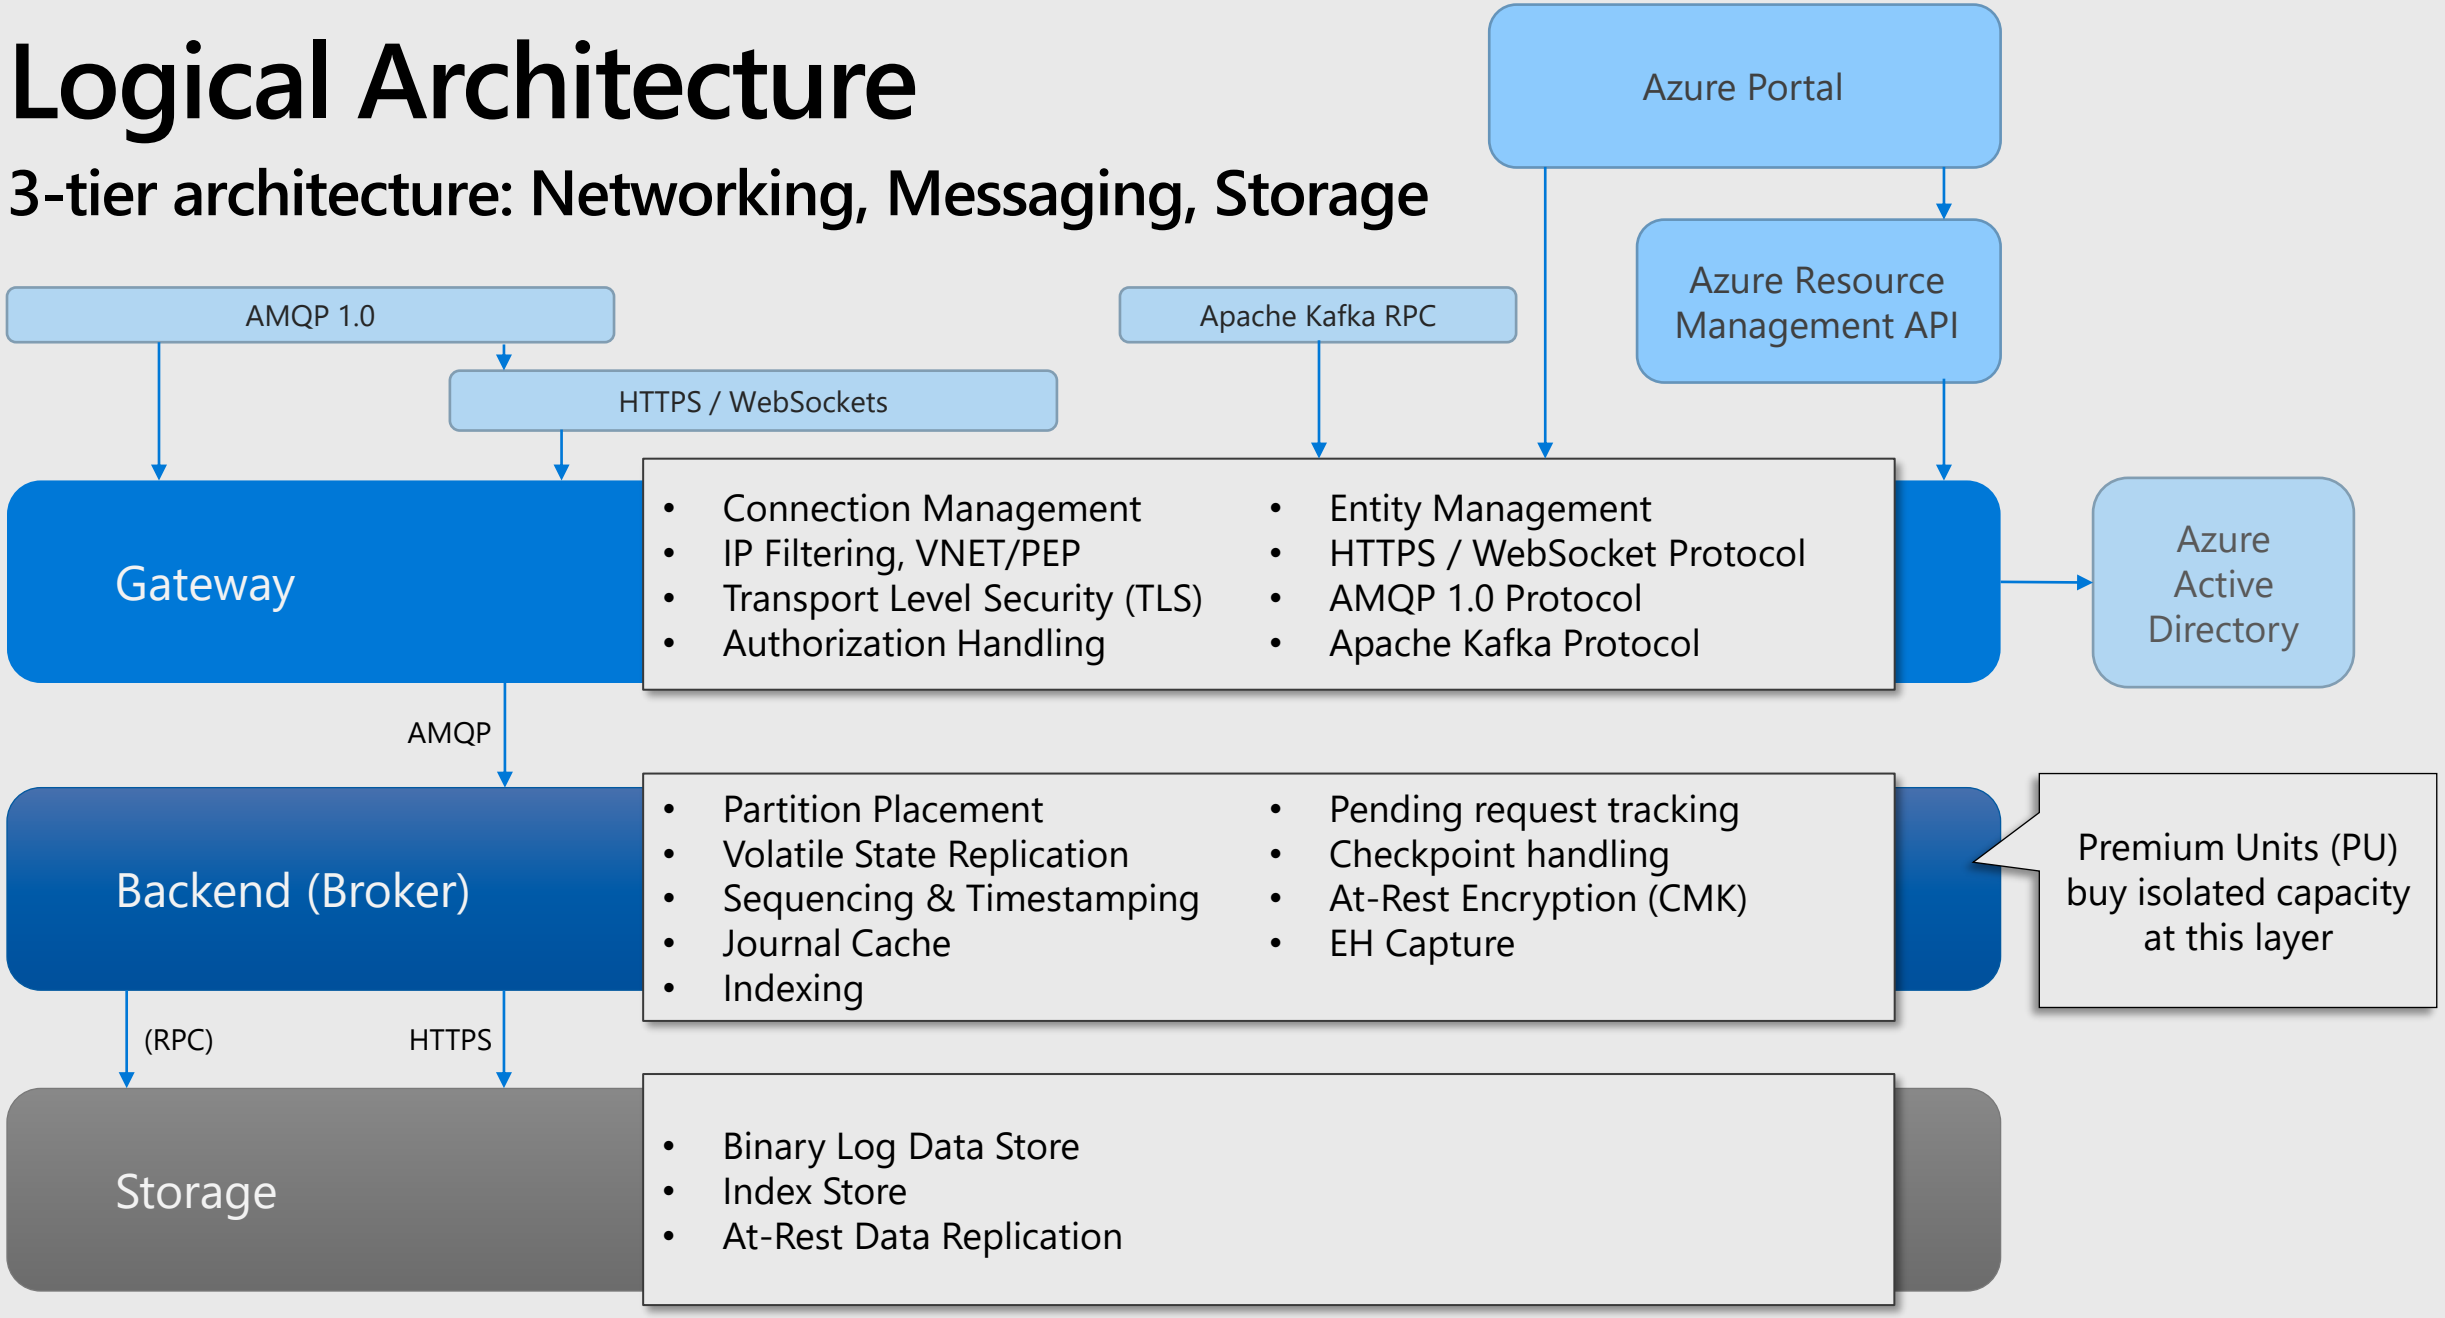

Fault Domains & Availability Zones

High-Availability starts at the physical reality

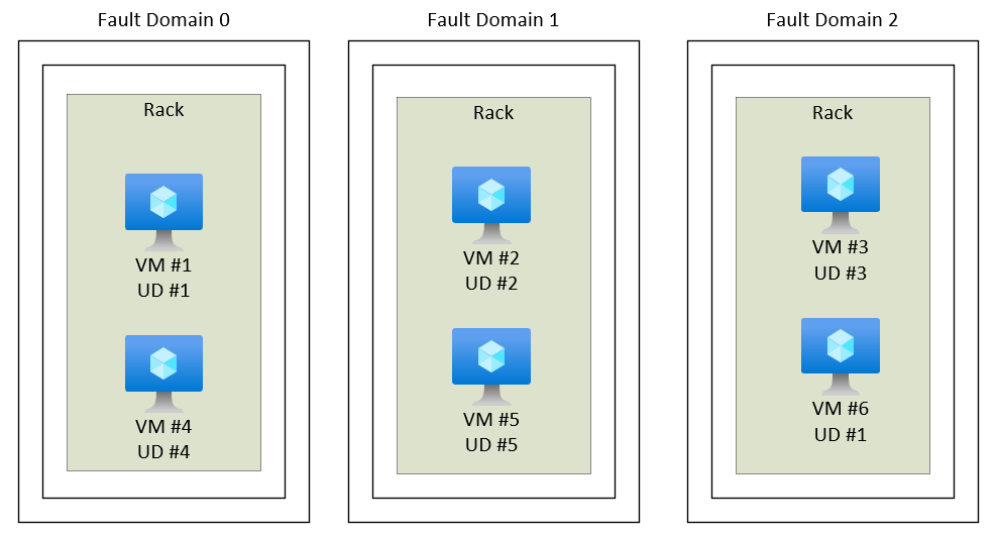

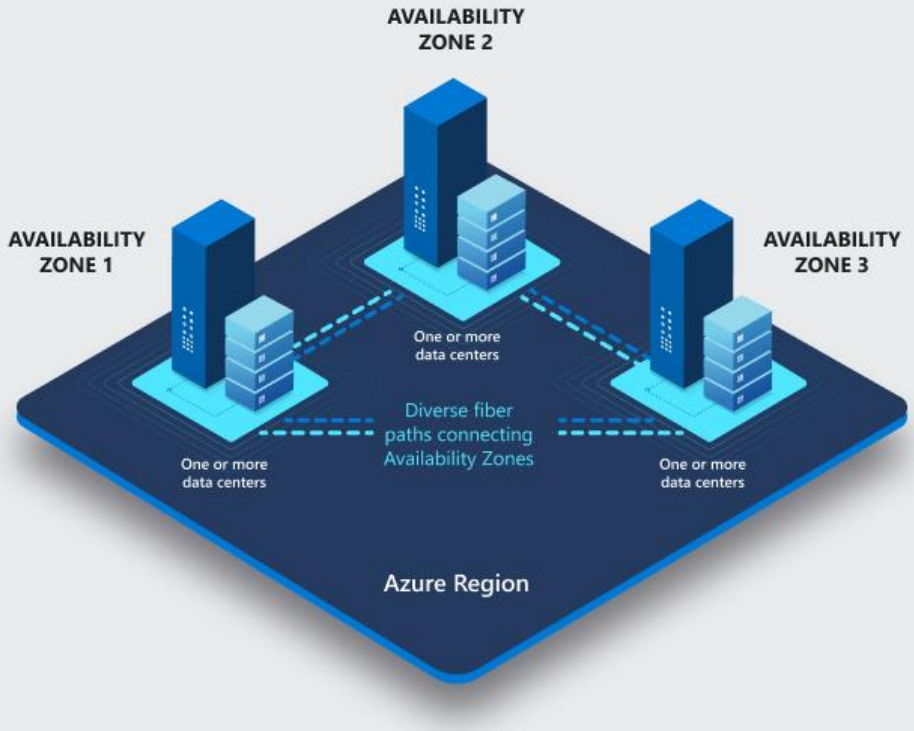

Fault Domain Placement:
Cluster VMs are spread across at least 3 fault domains such that the loss of a rack or network poses no availability risk. Recovery from a fault domain failure is fully automated and the system maintains SLA.

Availability Zones Placement:
Each cluster spans three availability zones and maintains SLA without any tolerance for data loss when one or two zones fail.

Backend and Gateway Clusters

Logical Architecture meets Placement

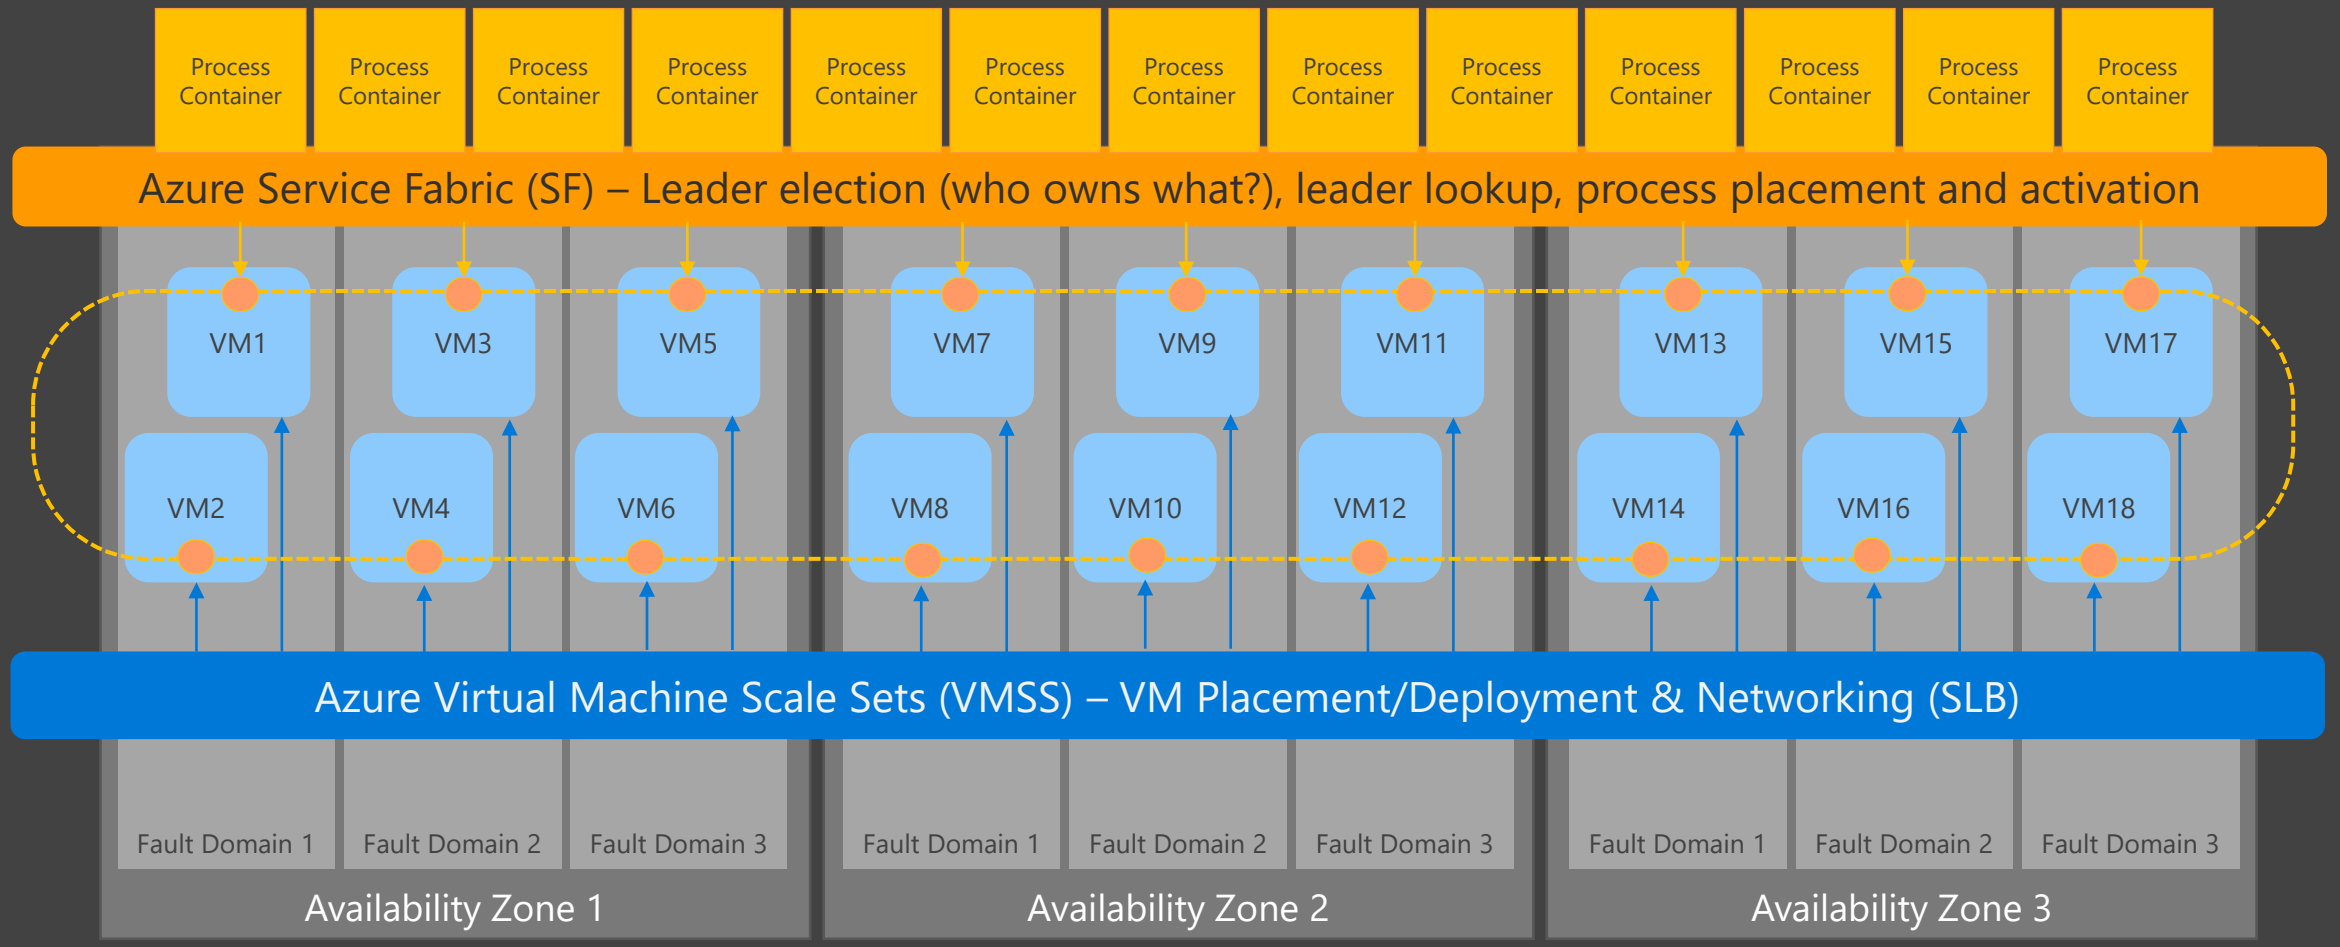

Backend and Gateway Clusters

Principles

Everything is automated. There's no manual intervention in placement of VMs or processes or workloads.

Process placement and servicing/upgrade cycles are controlled by Azure Service Fabric.

Each cluster has a well-known number of virtual machines. We don't use auto-scaling with VMs. We may (rarely) choose to initiate scale up/down to a different level.

Allocation of namespaces to clusters are based on well-established heuristics (Standard) or deterministic capacity limits (Premium). Customer-defined allocation in Dedicated.

Gateway and Backend

Workload placement

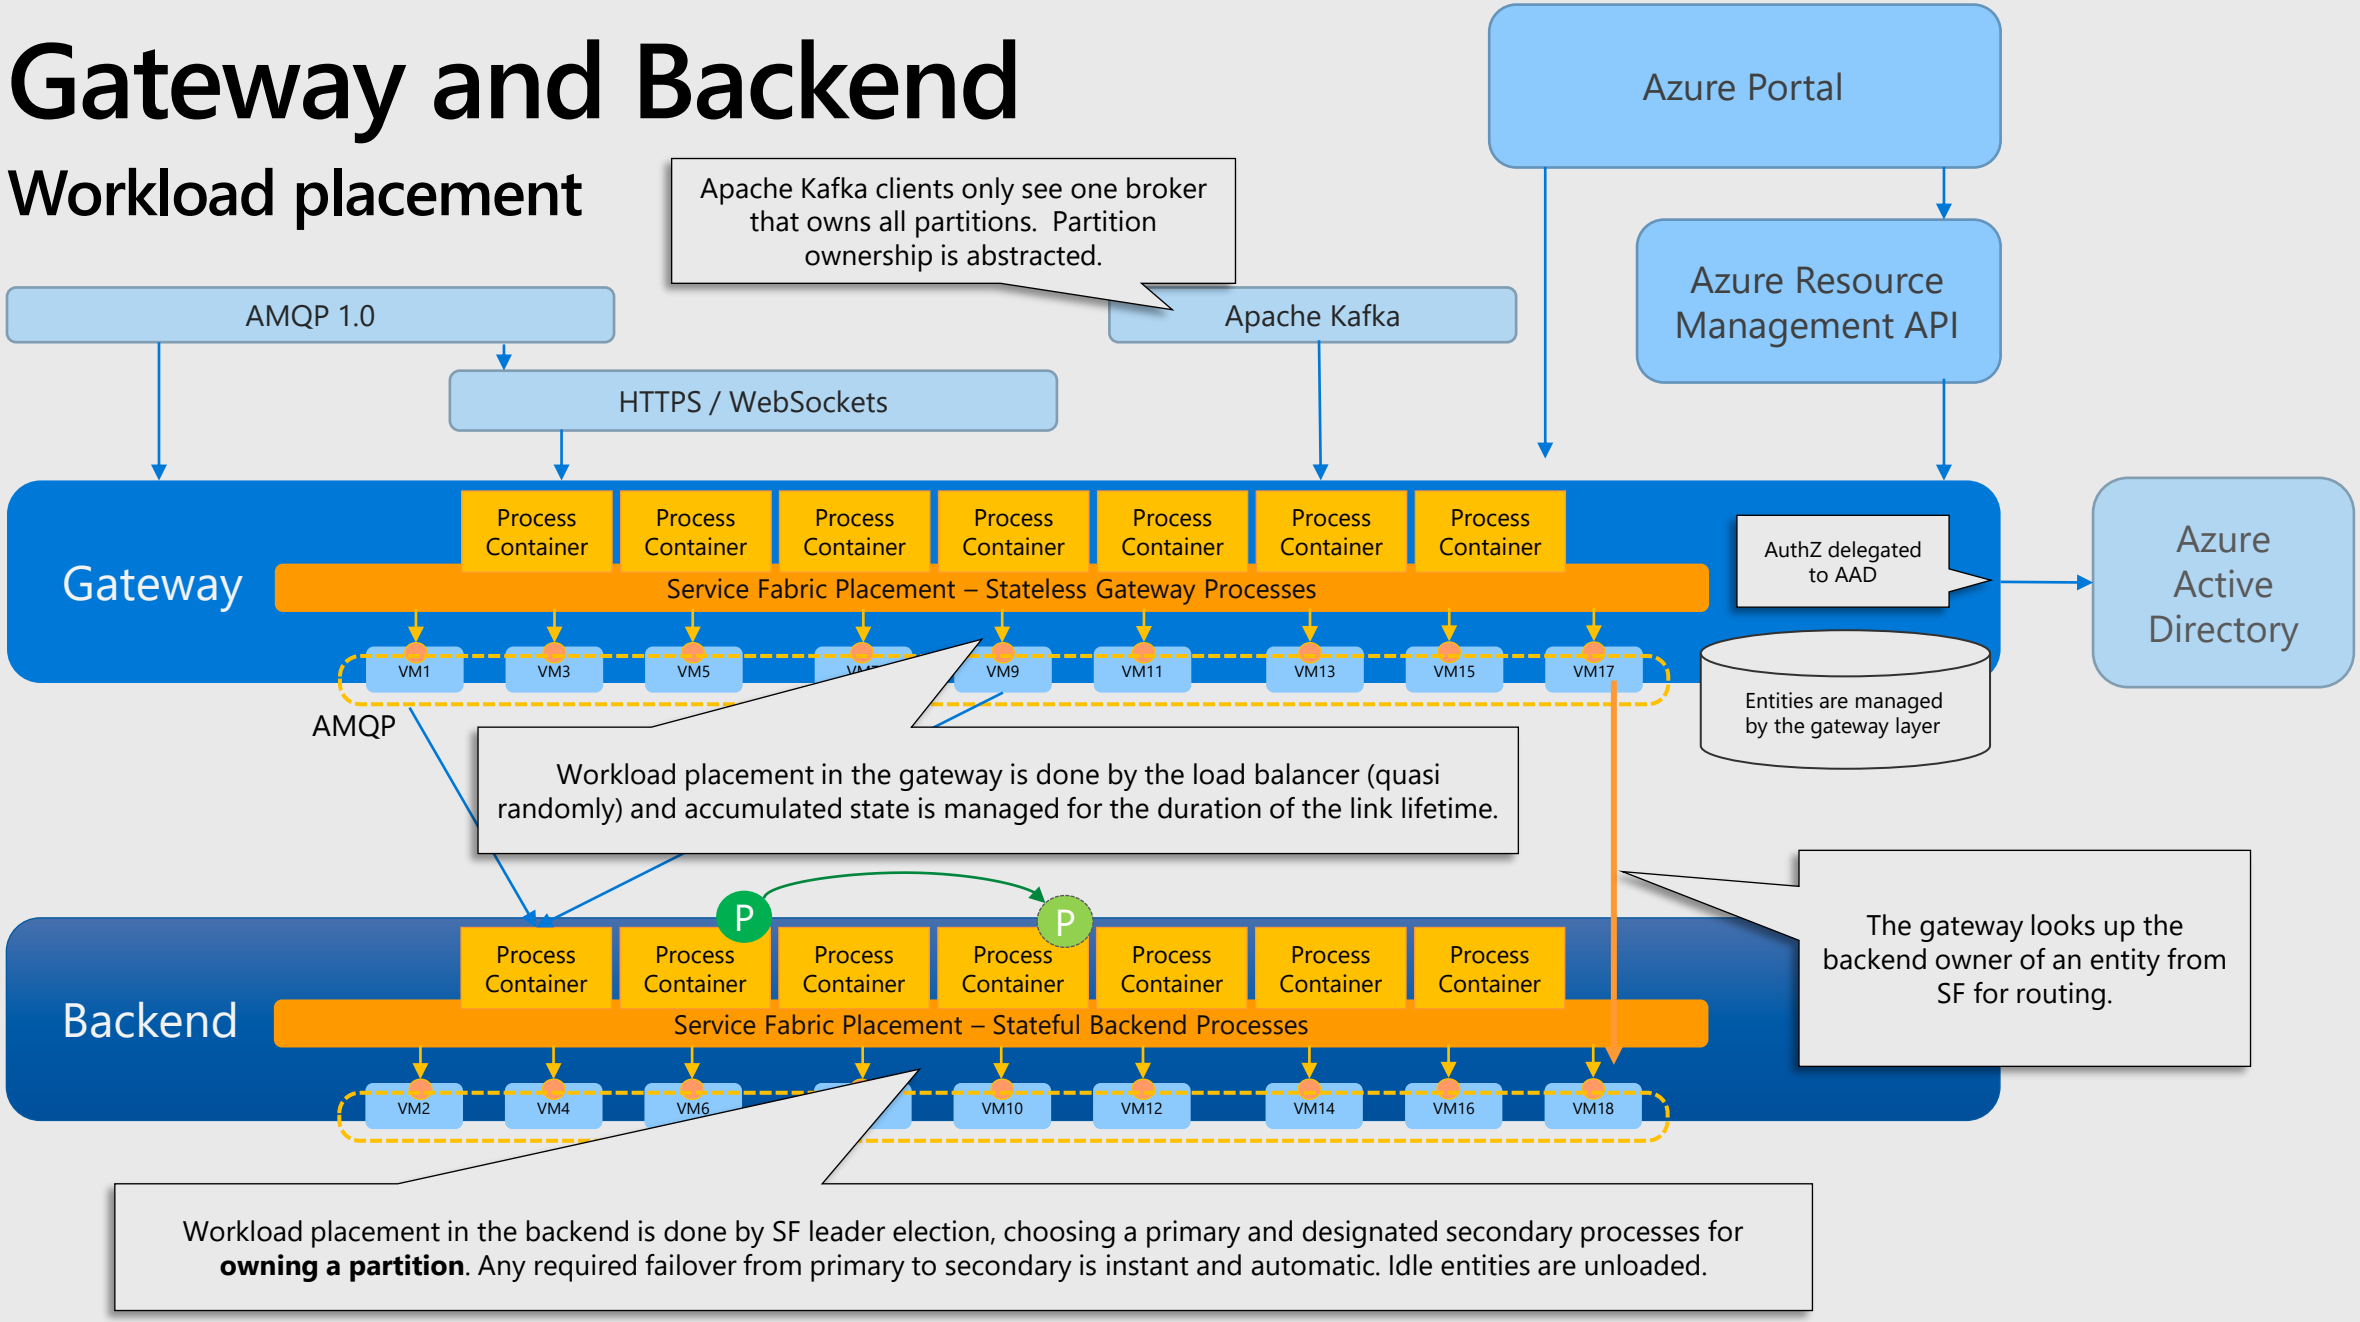

Backend: Event Hubs Premium.

Workload placement

Namespace PU are split across processes:
1 PU = 2 Proc (8GB Mem), 1 Core/Proc (2C)
2 PU = 2 Proc, 1 Core/Proc + 1 Core (3C)
4 PU = 2 Proc, 2 Core/Proc + 1 Core (5C)
8 PU = 4 Proc, 2 Core/Proc + 1 Core (9C)
16 PU = 4 Proc, 4 Core/Proc + 1 Core (17C)

Partition ownership is dynamically mapped to the process(es) associated with a namespace via SF

Host 2

Namespace Broker Process 2 (**4PU**)

P P P P P P P P

Host 1

Namespace Broker Process 1 (**4PU**)

P P P P P P P P

Cores are exclusively mapped to a broker process.

>=2PU: 1 Core extra for utility tasks.

Isolated VMs / L88is

CPU Cores L-Series

1 2 +

Memory Allocations

8GB

VM6

Backend

Process Container | Process Container | Process Container | Process Container | Process Container | Process Container | Process Container

Service Fabric Placement – Stateful Backend Processes

VM2 VM4 VM6 VM8 VM10 VM12 VM14 VM16 VM18

Designated secondaries in different fault domains and AZ reserved for failover

Storage Layer – Event Hubs Standard

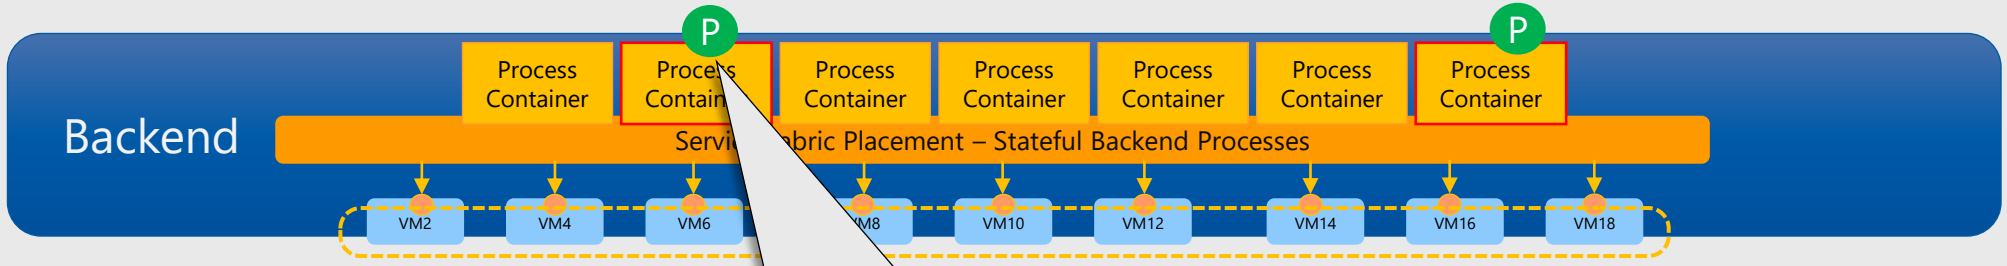

Events are written to append blobs ("extents"). All writes are synchronously triple replicated across availability zones.

Extents are sealed after hitting a size threshold.

Expiry removes extents.

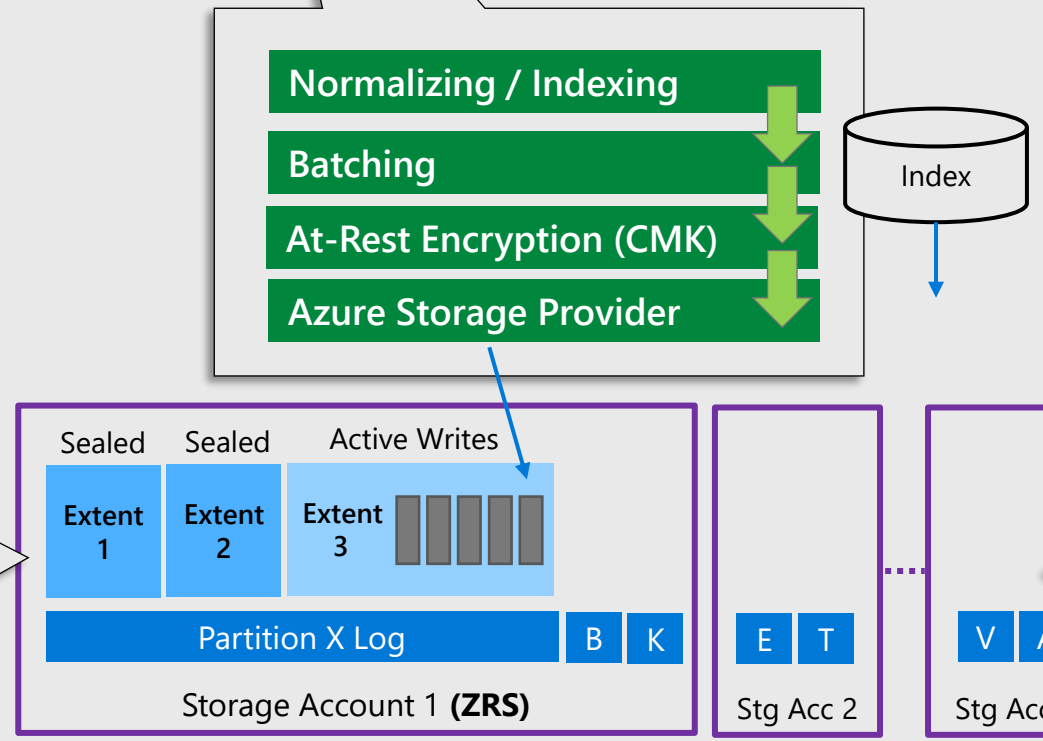

Each Event Hub Cluster uses numerous storage accounts to spread the load and overcome storage quota limitations.

Partition storage is allocated to those storage accounts.

Storage Layer – Event Hubs Premium

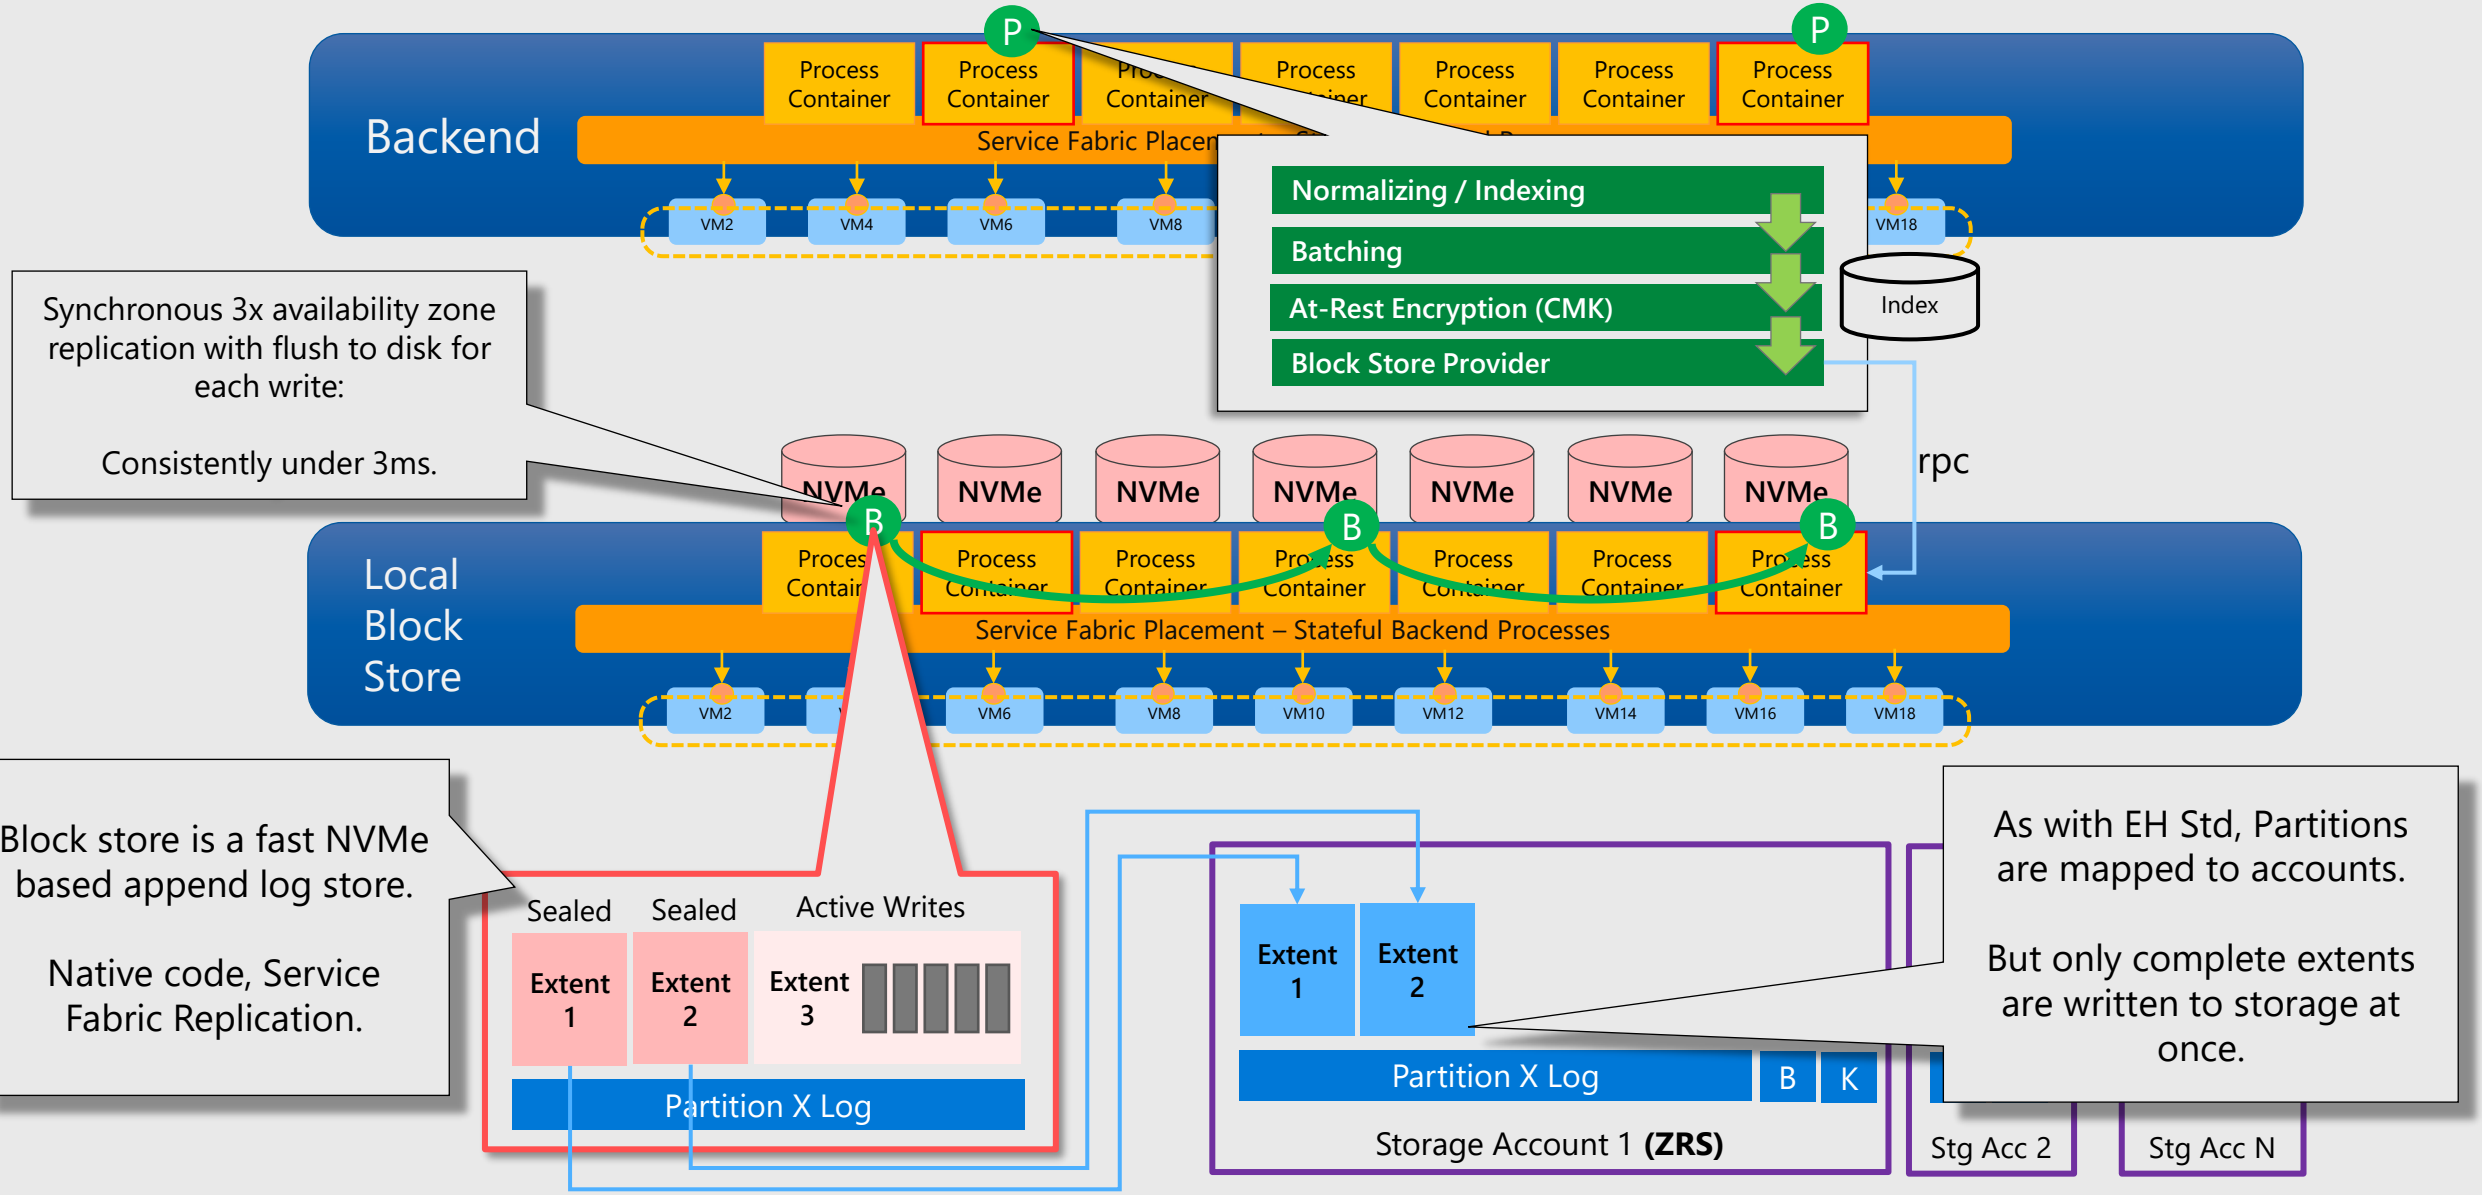

Networking Features

Firewall, Virtual Network Integration with Private Endpoints

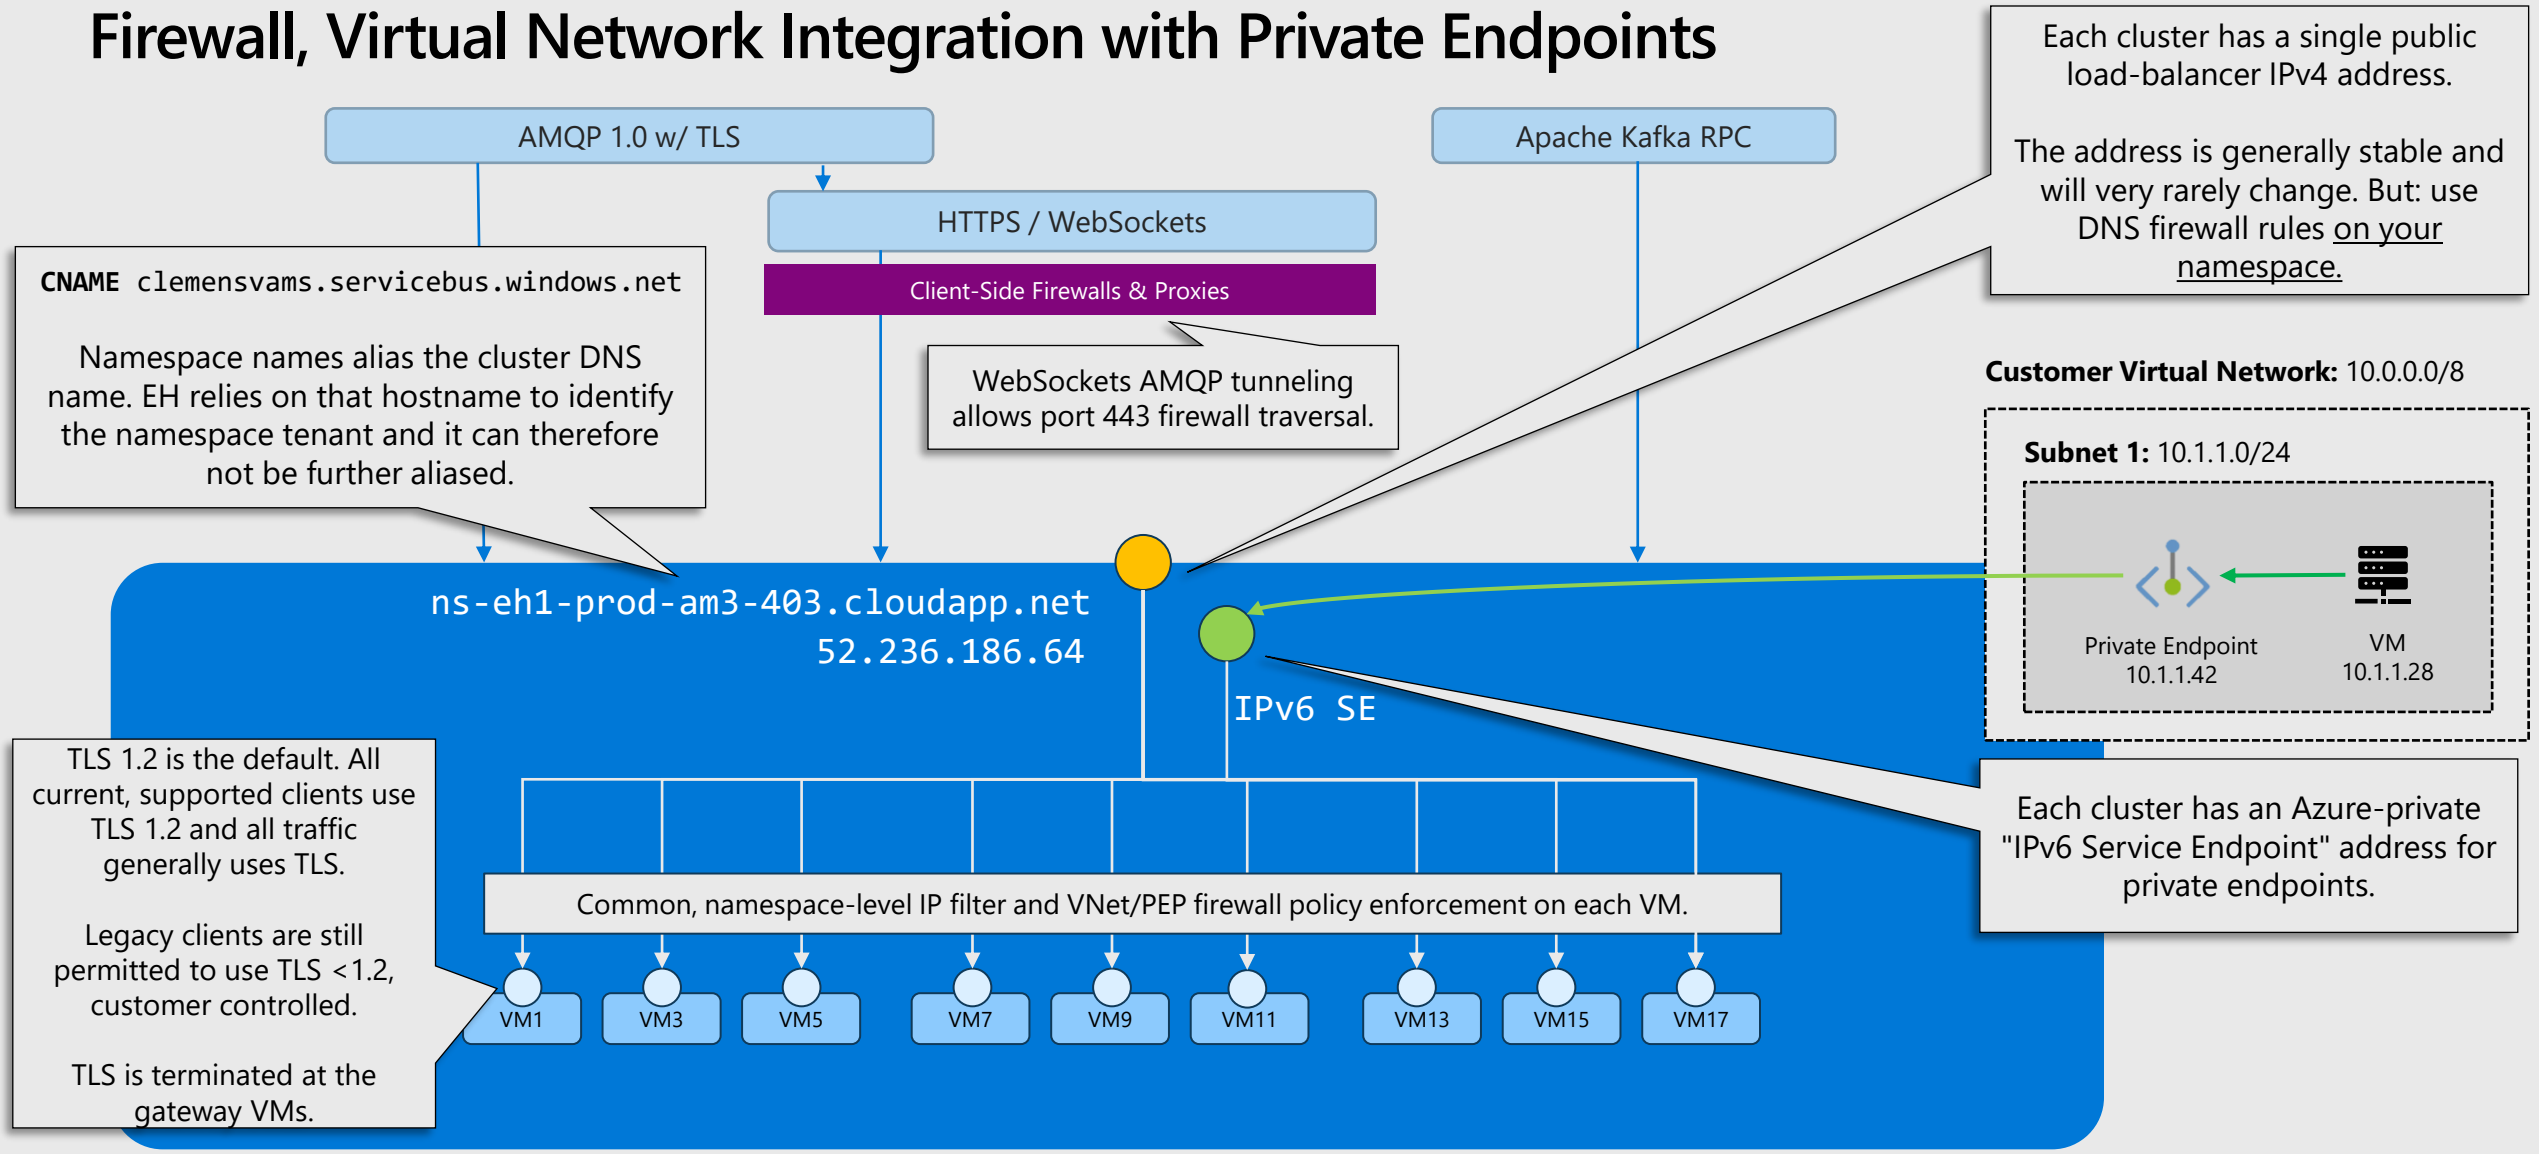

Networking Features

Event Hubs namespaces can be attached to one or more virtual networks and the public IP address space concurrently and act as safe "Layer 7" stream bridges.

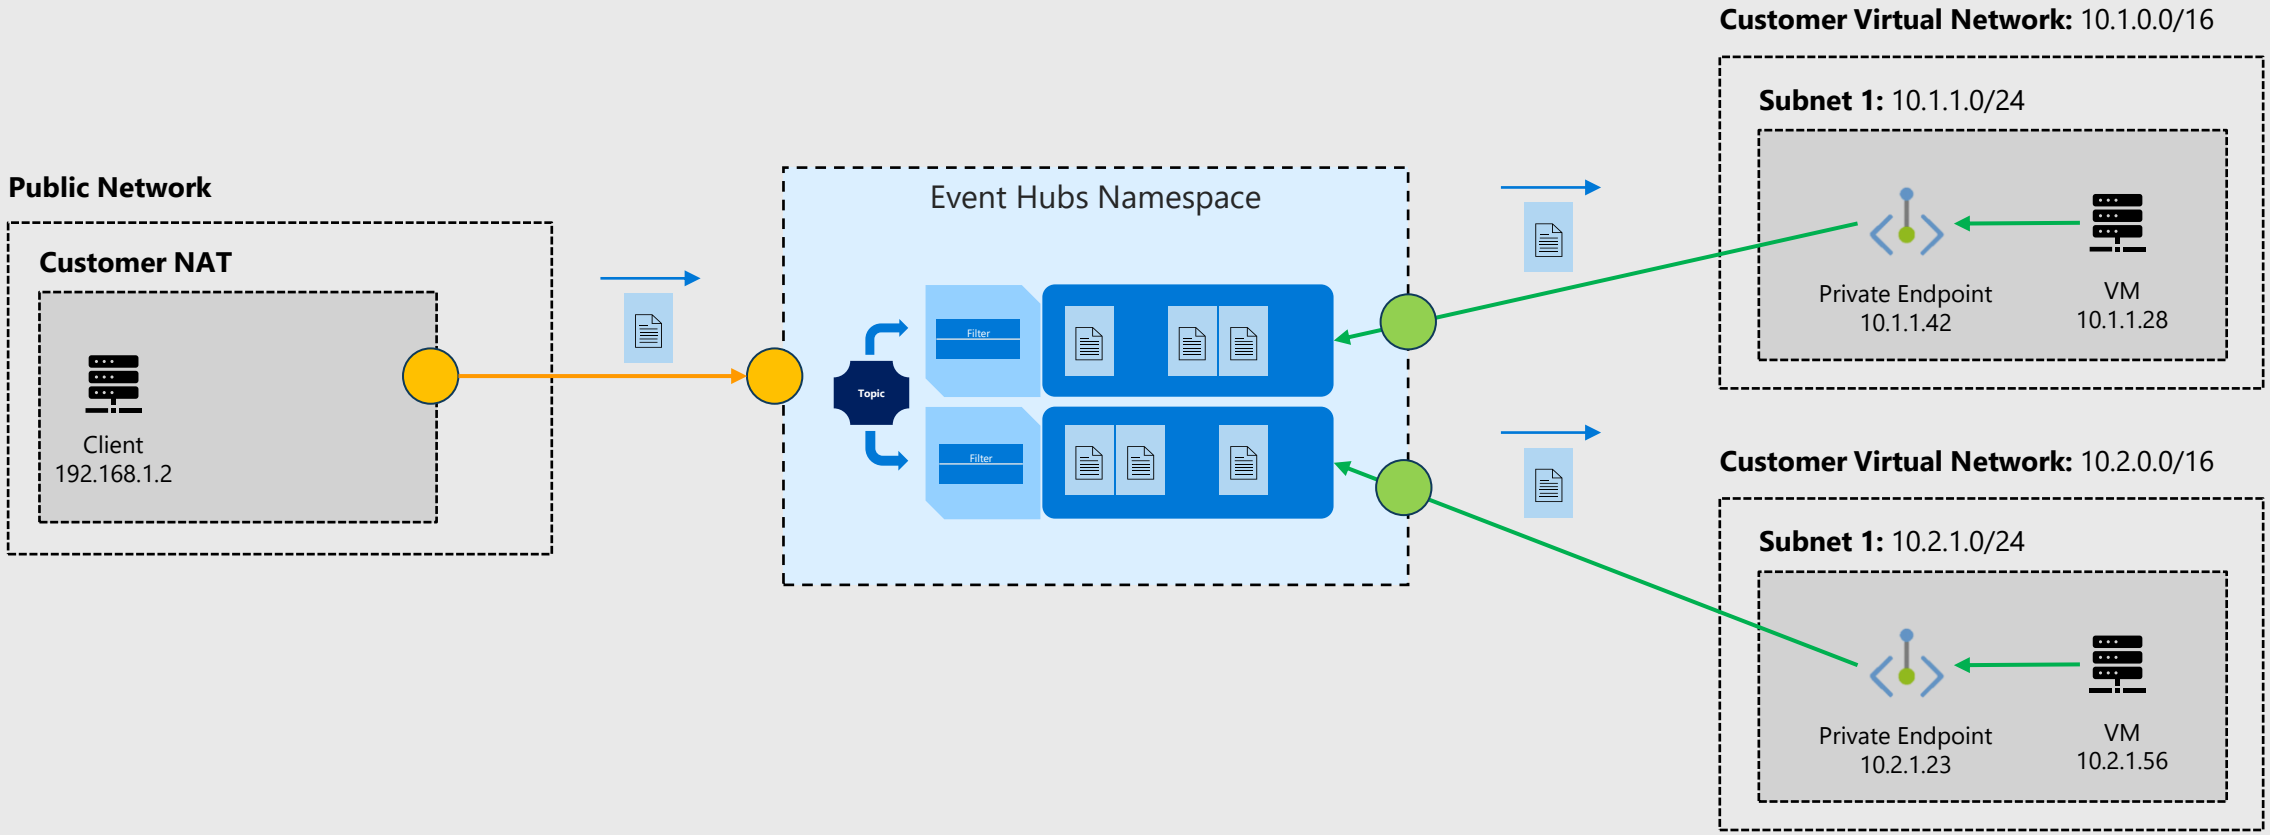

Flows: Automation Tasks

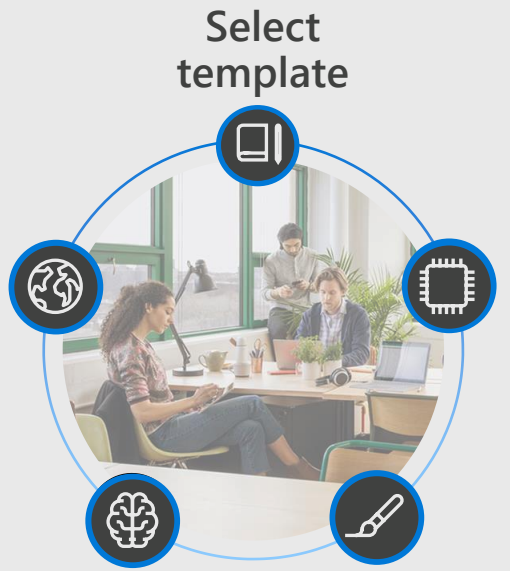

EventHub replication visualization

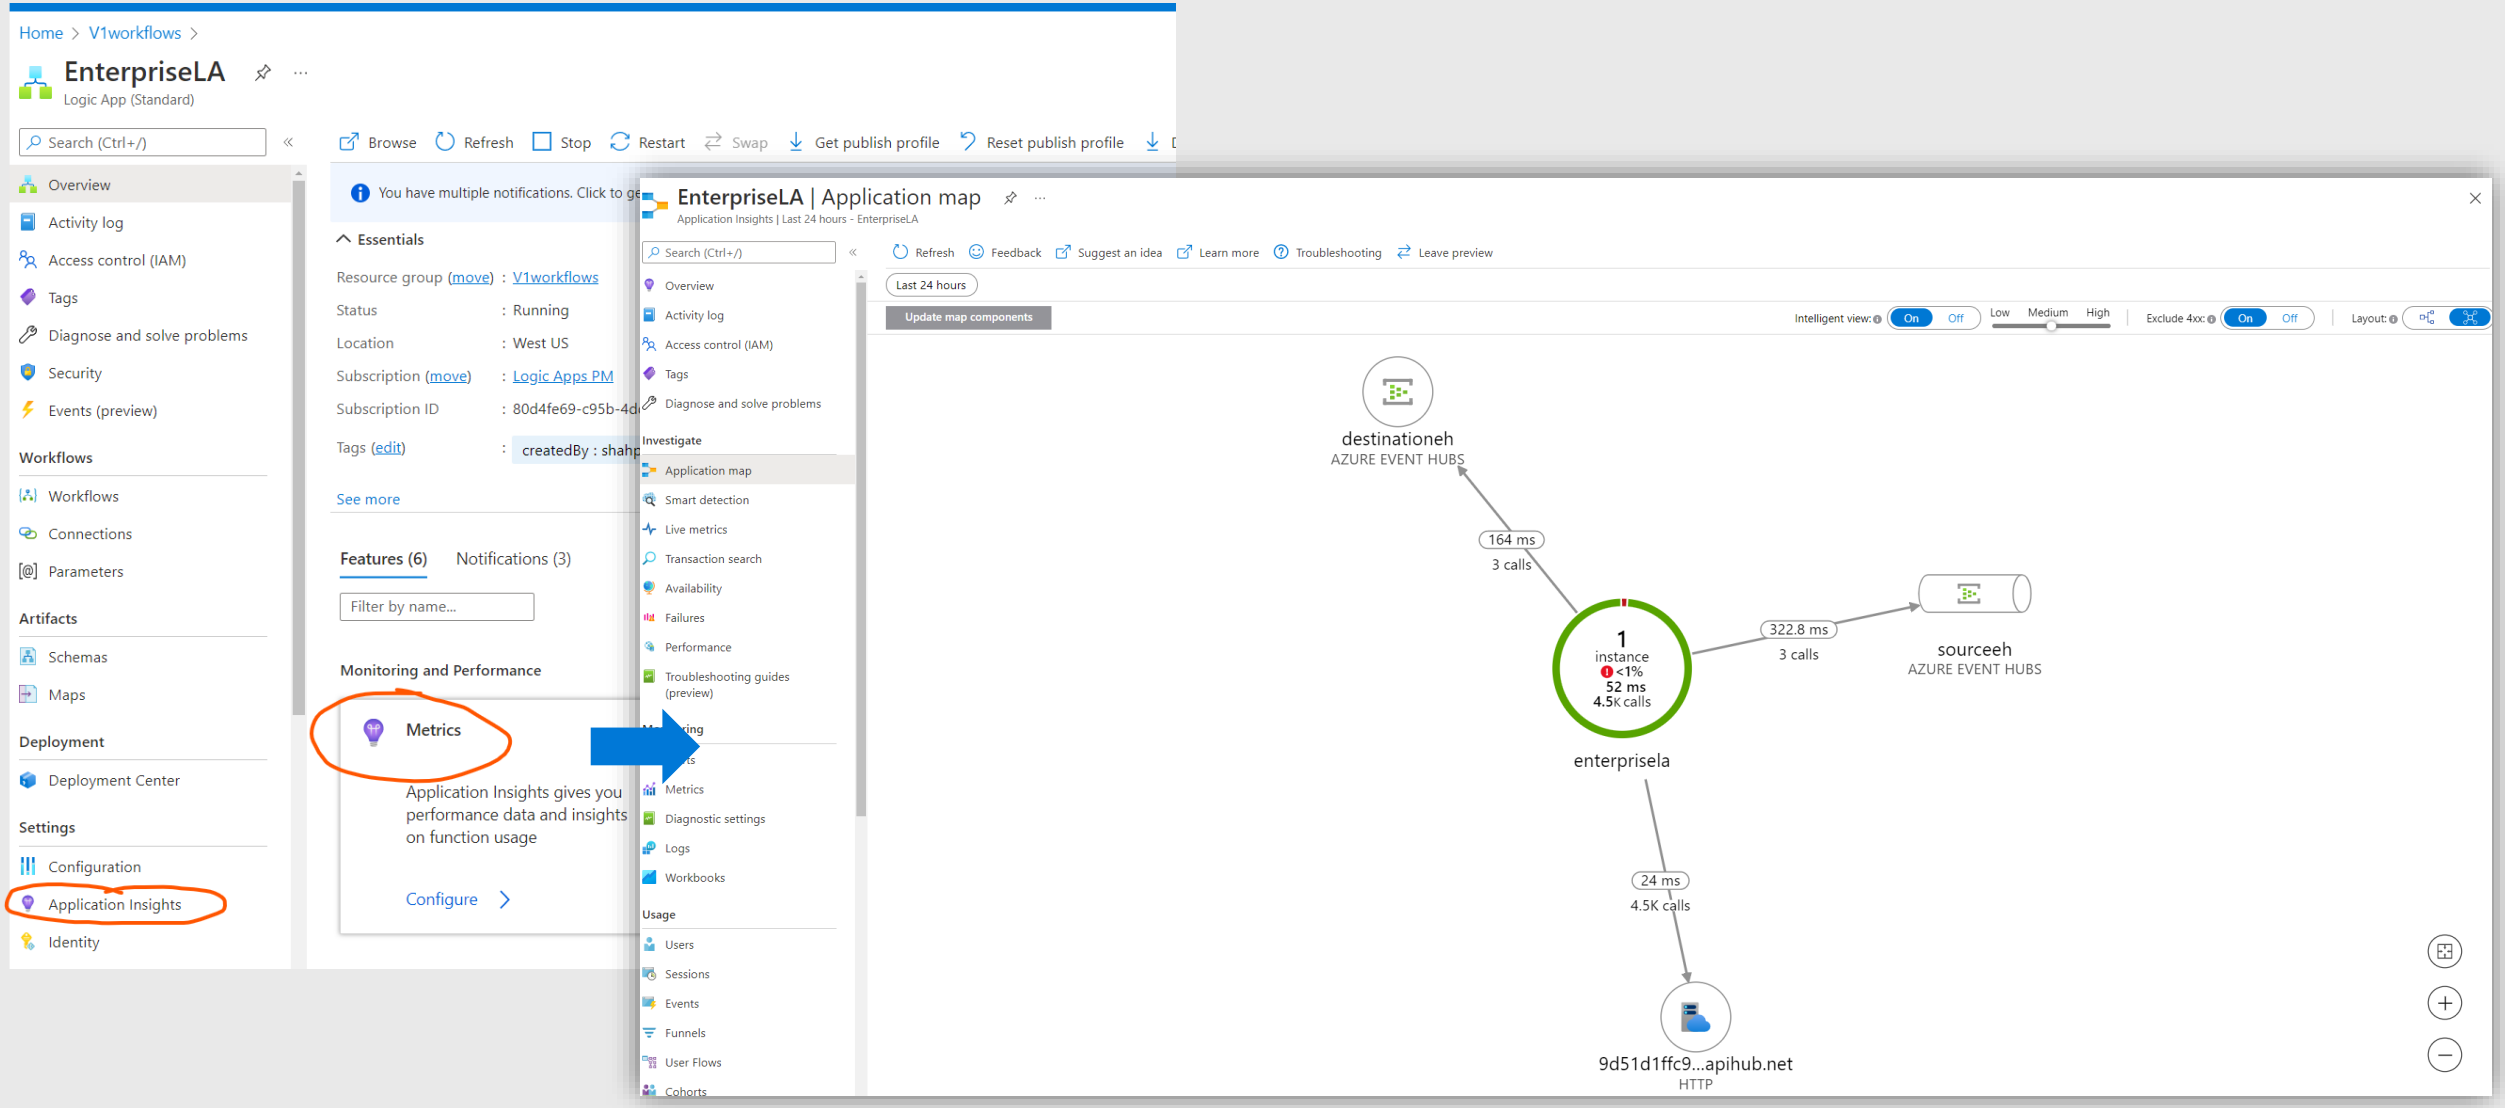

devreleventhub (devreldemoeventhub/devreleventhub)

Event Hubs Instance

Search (Ctrl+/)

Overview
Access control (IAM)
Diagnose and solve problems

Settings
Shared access policies
Properties
Locks

Entities
Consumer groups

Features
Capture
Process data

Automation
Tasks (preview)
Export template

Support + troubleshooting
New Support Request

+ Consumer group | Delete | Refresh

Essentials

JSON View

Resource group (move)	: DBGroup	Status	: Active
Location	: West US	Namespace	: devreldemoeventhub
Subscription (move)	: Visual Studio Enterprise	Created	: Monday, July 18, 2022 at 21:32:42 PDT
Subscription ID	: feaa8ca8-ba47-40fb-b2ce-7ddf1b49258b	Updated	: Monday, July 18, 2022 at 22:13:29 PDT
Partition Count	: 1	Message Retention	: 1 day

Capture events

Use Capture to save your events to persistent storage.

Process data

Process data instantly with Azure Stream Analytics using no-code editor or SQL.

Connect

Authenticate with connection strings and SAS policies.

Checkpoint

Create consumer groups to checkpoint your events.

Event Hub Contents: 1 CONSUMER GROUP
Event Hub Status: ACTIVE
Message Retention: 1 DAY
Partition Count: 1

Show data for the last: 1 hour | 6 hours | 12 hours | 1 day | 7 days | 30 days

Requests

600

Messages

100
90

Throughput

100B
90B

Process Event Hub data with Azure Stream Analytics

Process your Event Hub data using no-code drag and drop experience. (Preview)

Filter and ingest to Synapse SQL

- Preview Event Hub data
- Decide table schema
- Select Synapse SQL table

Start

Capture data to ADLS Gen2 in Parquet format

- Save events to ADLS Gen2 in Parquet
- Specify a time or size interval

Start

Materialize data in Cosmos DB

- Maintain a view of your data in Cosmos DB
- Select the fields to group by
- Define aggregations like count, sum, average
- Set a time period

Start

Filter and ingest to ADLS Gen2

- Preview Event Hub data
- Decide table schema
- Select ADLS Gen2 account

Start

Start with a blank canvas

- View incoming data and define schema
- Define transformations on your input data
- Select output to egress streaming data

Start

Event Hubs Premium is *Fast*!

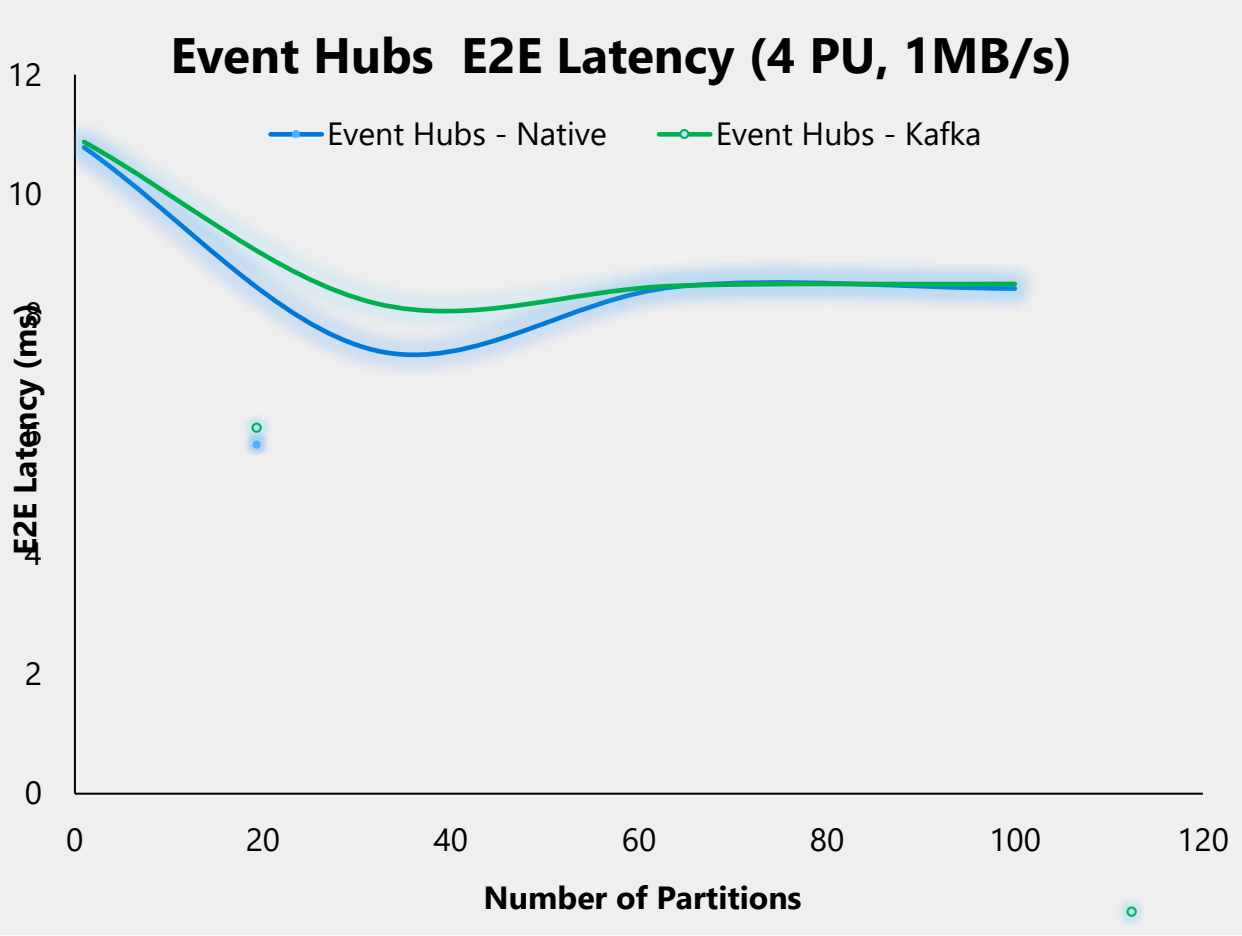

Event Hubs Premium end-to-end latency < 10ms for most Event Streaming workloads.

Performance Metric	Expected Range	Comments
Throughput	~5-10 MB/s per 1 PU	Maximum throughput can be achieved by scaling partitions For both Kafka and AMQP
Latency	<10 ms	With 1 MB/s load, 4 PU namespace For both Kafka and AMQP

Effective Global Reliability

Uptime – Is the cluster available to accept connections?

100%

Reliability – Are individual operations suceeding?

>99,9999% weekly average

Azure Event Hubs

Platform-as-a-Service Event Stream Broker

Use the Apache Kafka® API, but with far lower cost and better performance.

Fully managed: You use the features, Azure deals with everything else

AMQP 1.0 standards compliant, Apache Kafka® wire-compatible

Polyglot Azure SDK and cross-platform client support

Industry-leading reliability and availability

Fast.

Thank You!

CARIBBEAN
DEVELOPERS
CONFERENCE

That was a lot, wasn't it?

Q&A