



ITARC Stockholm 2021

How to architect something really big?

Clemens Vasters
Cloud Platform Architect
Azure Messaging
Microsoft

@clemensv



Do not.

I'm absolutely serious, and here's why ...

Before we go there





Architecture
is
Opinionated





Architecture
is
Engineering



Dear Barry,

“How to architect something really big?” is a trap of a session title and you knew this when you asked me to do this talk.

So here we are. There is no good, universal answer to a question that broad.

Yet we still have some 40 minutes to fill...

Azure Messaging. Arguably Big.



>140 Trillion

Monthly Requests
across Messaging

500 GB

Per minute at peak
Written to storage by
Event Hubs Capture

99.99997%

Weekly success rate
across Messaging

>168 PB

Monthly Data Volume

>60 regions

Running our services
Ring 0

100%

Of the top 500 Azure
customers use
messaging services


>50,000+

VMs run our service


4.5
Trillion

Daily Requests


What did we build?



7 operationally distinct
service types.
Starting from one.



Regions.
Availability zones.
Clusters.



Cluster Tenants.
In-Cluster Isolation.

Big things
split things.

Meet Your Adversaries

- Budget !!
- Business Model !!
- Capacity
- Latency
- Affinity
- Failure
- Thieves
- Idiots



Capacity

- 'Capacity' defines the limits of how much information can be held, transferred, or processed by a system during some period
 - Machine instructions that can be processed within the period to compute some output based on some input
 - Amount of information that can be held in the storage container at any given time
 - Amount of information that can be transferred into and out of a scope of location within the period
- If capacity were 'unlimited', all information would be ubiquitous



Latency

- 'Latency' is the time it takes to accomplish a certain task
 - Compute an output based on some input using an algorithm
 - Store information into or retrieve information from some kind of container
 - Transfer information between different system scopes and/or physical locations
- If latency were 'zero', all information would be instantly available everywhere
- Latency is Software's Friction



Affinity

- 'Affinity' is the degree to which a particular set of information is bound to a scope or location
 - Affinity of original and intermediate state to a particular sequence of processing while producing a desired output
 - Affinity of information to a particular container due to size or transfer latency
 - Affinity of transfer to a particular communication route
 - Affinity to a particular physical device or location (including someone's pocket)
- Affinity is the biggest inhibitor for scaling out



Failure

- 'Failure' is any condition that interrupts or otherwise adversely affects the common, planned flow of information processing
 - Programming errors
 - Physical failure of equipment
 - Exhaustion of system capacity
 - Environment conditions
- 'Permanent Failure': 'Call the Doctor'
- 'Intermittent Failure': 'Wait for it to go away'



Thieves and Idiots

- “Thieves and/or idiots”
 - Security protects against thieves
 - Identify who you’re dealing with (authentication)
 - Only allow them to do what they’re entitled to do (authorization)
 - Safety protects against idiots
 - Make sure information is safe against inadvertent destruction (“Oops. I didn’t mean to do that”)
- Security and Safety are about undesired interactions with people.



Considerations.

Business

- What are you selling?
- How does a new feature impact what you are selling?
- What does the new feature do for your bottom line?
- Might that new feature hurt your bottom line?



Blast radius

- How much impact can a bug have?
- How much impact can a dependency issue have?
- How much impact can an attack have?
- How much impact can an operator have?

A large, stylized purple number '2' is positioned on the left side of the slide, partially overlapping the title area.

Tenancy and Partitions

- All "really big" systems have tenants.
 - Users, Accounts, Namespaces, etc.
- Natural partitioning boundary
- Consider workflows and objects as tenants
 - Short-lived tenants
 - Dynamic allocation

Failures and Disasters

- Failures will happen. Big ones, too. Sorry.
 - (Hi Facebook!)
- How much are you willing to pay for what kind of risk mitigation?
 - Lots of disaster scenarios are the boogeyman.
 - No, you can't run platform better than a hyperscaler
 - Availability zones are GREAT tech
- Systems don't die clean deaths. Your disaster drills are a lie.

Automation

- Radically automate (almost) everything
 - Every automated step is less of an opportunity for someone making a manual mistake
 - You can't scale globally with manual steps
- Do not hand the reliability predictability of your system to a load-adaptive algorithm unless you really understand what it does and why.



Observability

- Instrument deeply: "weird" stuff will happen
- 0.005% reliability gains matter at scale
- The root cause is often elsewhere



Truth and Lies

- You know the system reality. Does your user?
Do they need to?





Thank you